Deep Learning–Based Webshell Detection and Prevention Using Word-Level Attention Mechanisms

A dissertation submitted to the Jawaharlal Nehru University
In partial fulfilment of the requirements
for the award of the degree of

MASTER OF TECHNOLOGY

IN

COMPUTER SCIENCE AND ENGINEERING

By

NIKHIL KUMAR

SCHOOL OF ENGINEERING
JAWAHARLAL NEHRU UNIVERSITY
NEW DELHI-110067

MAY 2025

SCHOOL OF ENGINEERING

JAWAHARLAL NEHRU UNIVERSITY

NEW DELHI-110067

DECLARATION

This is to certify that the dissertation entitled "***Deep Learning–Based Webshell Detection and Prevention Using Word-Level Attention Mechanisms***" is being submitted to the School of Engineering, Jawaharlal Nehru University, New Delhi, in partial fulfilment of the requirements for the award of the degree of **Master of Technology** in ***Computer Science and Engineering***, as part of the dual degree programme in the School of Engineering, is a record of bona fide work carried out by me under the supervision of ***Dr. Benoy Kumar Ray*** and co-supervision of ***Dr. Reetu Raj***.

The matter embodied in the dissertation has not been submitted in part or in full to any university or institution for the award of any degree or diploma.

NIKHIL KUMAR
(***20/11/EC/052***)

SCHOOL OF ENGINEERING

JAWAHARLAL NEHRU UNIVERSITY

NEW DELHI-110067

CERTIFICATE

This is to certify that this dissertation entitled "***Deep Learning–Based Webshell Detection and Prevention Using Word-Level Attention Mechanisms***" submitted by ***NIKHIL KUMAR*** to the School of Engineering, Jawaharlal Nehru University, New Delhi for the award of the degree of **Master of Technology** in ***Computer Science and Engineering***, as part of the dual degree programme in the School of Engineering, is a research work carried out by ***him/her*** under the supervision of ***Dr. Benoy Kumar Ray*** and co-supervision of ***Dr. Reetu Raj***.

Dr. Benoy Kumar Ray
(Supervisor)

Dr. Reetu Raj

(Co-Supervisor)

Prof. T. V. Vijay Kumar
(Dean)

ACKNOWLEDGEMENTS

I would like to express my sincere gratitude to my supervisor, **Dr. Benoy Kumar Ray**, and co-supervisor, **Dr. Reetu Raj**, for their invaluable guidance, support, and encouragement throughout this research. Their expertise and motivation have been instrumental in shaping the direction and outcome of this dissertation.

I am also grateful to the School of Engineering, Jawaharlal Nehru University's faculty members and staff for providing me with a conducive environment for learning and research.

I extend my heartfelt thanks to my friends and colleagues for their constant support and discussions that helped me improve my work. Most importantly, I am forever indebted to my family for their unconditional love, patience, and encouragement.

Thank you all for being a part of this journey.

ABSTRACT

Malicious scripts injected into web servers as webshells are particularly dangerous to contemporary web infrastructures because they facilitate unauthorised access, data exfiltration, and secondary compromise. Conventional signature and heuristic‑based detection methods find it hard to generalize against new evasion and obfuscation techniques. This dissertation offers an end‑to‑end deep‑learning framework for autonomous webshell detection and prevention, which applies both source‑code and HTTP‑traffic analysis enhanced by attention mechanisms.

First, we pose the malicious script detection as the binary classification of tokenized source-code lines and HTTP request features. We built an end‑to‑end system consisting of Word2Vec‑based embeddings, dynamic construction of the tokenizer, and memory‑mapped data to handle millions of code samples. We designed and compared three architectures: bidirectional GRU with attention, convolutional TextCNN, and Transformer encoder, with each hyperparameter search optimized and regularized with early stopping. Experimental evaluation on an assembled dataset of PHP and ASP.NET webshells attains up to 98.3 % accuracy, 97.6 % precision, and an AUC of 0.992, improving traditional baselines by 12% in F1‑score.

Expanding on these findings, we apply the framework to the proactive detection of live HTTP traffic with the introduction of an online inference pipeline with feature extraction, imbalance management, and real-time alerting. Experiments on production traffic scale demonstrate persistent detection accuracy (95.1% F1) with sub-100 ms per-request latency.

Contributions include: (1) an original Word Attention Mechanism for source-code classification; (2) an efficient memmap‑based preprocessing approach; and (3) an integrated framework connecting static and dynamic analyses for offline training and online prevention. Continual learning to keep pace with developing webshell variants is future work.

Table of Contents

ACKNOWLEDGEMENTS **4**
ABSTRACT **5**
Table of Contents **6**
Chapter 1 **17**
INTRODUCTION **17**
1.1 Background and Motivation 17
1.2 Research Problem Formulation 19
1.3 Aims and Objectives 21
○ Objectives of the Study 21
1. Data Collection and Preprocessing 21
2. Model Development and Optimization 22
3. Evaluation and Benchmarking 22
4. Interpretability and Explainability 22
5. Real-World Applicability 22
6. Contribution to Research and Industry 22
1.4 Scope and Delimitations 23
Scope of the Study: 23
Delimitations of the Study: 23
1.5 Research Methodology 24
○ 1.5.1 Research Design 24
○ 1.5.2 Data Acquisition and Description 24
1.5.3 Data Preprocessing 25
1.5.3.1 Cleaning 25
1.5.3.2 Tokenization and Parsing 25
1.5.3.3 Vectorization 25
1.5.3.4 Padding and Truncation 25
○ 1.5.7 Tools and Technologies Used 25
○ 1.5.8 Ethical Considerations 26
○ 1.5.9 Summary 26
1.6 Key Contributions 26
1.6.2 Engineering an interpretable injected webshell detector with deep learning 26
1.6.2 Creating a Curated and Balanced Malicious/Normal PHP Scripts Dataset 27
1.6.3 Application of Word-Level Attention to Cybersecurity Use Cases 27
1.6.4 Comparative Analysis with Conventional Machine Learning and Static Analysis Approaches 27
1.6.5 Developing a Lightweight and Deployable Detection Prototype 27
1.6.6 Open Research Avenues for Security-Focused NLP Models 28
1.6.8 Empirical Contribution via Evaluation Metrics and Visualization Tools 28
1.6.9 Contributions to Research and Education 29
1.7 Thesis Organization 29

- ○ Chapter 1: Introduction 29
- ○ Chapter 2: Literature Review 29
- ○ Chapter 3: Dataset Design and Preprocessing 30
- ○ Chapter 4: Proposed Methodology 30
- ○ Chapter 5: Discussion and Analysis 31
- ○ Chapter 6: Conclusion and Future Work 31

Chapter 2 **32**

SURVEY & THEORETICAL FRAMEWORK **32**

2.1 Overview of Webshell Attacks 32
- 2.1.3 Common Types of Webshells 33
- 2.1.4 Techniques for Deploying Webshells 33
- 2.1.5 Real-World Examples and Incidents 34
- 2.1.6 Challenges in Detecting Webshells 34
- 2.1.7 Need for Advanced Detection Techniques 34
- ○ Conclusion of Section 2.1 35

2.2 Static Analysis Techniques 35
- 2.2.1 Overview of Static Analysis in Web Security 35
- 2.2.2 Common Static Analysis Techniques 36
  - 2.2.2.1 Signature-Based Detection 36
  - 2.2.2.2 Rule-Based Heuristics 36
  - 2.2.2.3 Syntax Tree Analysis 36
  - 2.2.2.4 Entropy-Based Analysis 36
  - 2.2.2.5 YARA Rules 37
- 2.2.3 Tools and Frameworks for Static Analysis 37
- 2.2.4 Limitations of Static Analysis Techniques 37
- 2.2.5 Motivation for Moving Beyond Static Analysis 38
- ○ Conclusion of Section 2.2 38

2.3 Dynamic Analysis Techniques 38
- 2.3.1 What is Dynamic Analysis? 38
- 2.3.2 Techniques Used in Dynamic Analysis 39
  - a. Behavior Monitoring 39
  - b. Taint Analysis 39
  - c. System Call Tracing 39
  - d. Network Traffic Analysis 40
  - e. Browser-Emulated Execution (for Front-End Shells) 40
  - f. Honeypots and Sandboxes 40
- 2.3.3 Tools for Dynamic Analysis 40
- 2.3.4 Limitations of Dynamic Analysis 41
- 2.3.5 Hybrid Approaches 41
- ○ Conclusion of Section 2.3 41

2.4 Dataset Collection and Pre-processing 42
- ○ 2.4.1 Sources of Data 42
  - A. Malicious Webshell Samples 42

B. Normal PHP Scripts 42
○ 2.4.2 Dataset Composition 43
○ 2.4.3 Data Cleaning and Normalization 43
○ 2.4.4 Tokenization 43
○ 2.4.5 Vocabulary Building and Encoding 44
○ 2.4.6 Sequence Construction and Padding 44
○ 2.4.7 Labeling and Splitting 44
○ 2.4.8 Memory Mapping and Storage Optimization 45
○ 2.4.9 Dataset Augmentation (Optional Phase) 45
○ Conclusion of Section 2.4 45
2.5 Machine Learning in Malware Detection 45
2.5.1 Why ML for Malware Detection 45
2.5.2 ML Approaches 46
2.5.3 Key Features for Detection 46
2.5.4 NLP and Deep Learning 46
2.5.5 ML Frameworks 46
2.5.6 Challenges 46
2.5.7 Traditional Classifiers 47
2.5.8 Deep Learning Models 47
2.5.9 Traditional vs. Deep Learning (Summary) 47
Conclusion 48
2.6 Attention Mechanisms & Transformer Models 48
2.6.1 Introduction 48
2.6.2 Attention Mechanisms: Theory and Formula 48
○ A. Motivation 48
○ B. Mathematical Formulation of Attention 48
Explanation: 49
2.6.3 Self-Attention Mechanism 49
Attention(Q, K, V)=softmax QKTdkV 50
Explanation: 50
○ 2.6.4 Multi-Head Attention 50
2.6.5 Transformer Architecture for Code and Malware 50
○ Positional Encoding Formula 51
2.6.6 Benefits of Attention and Transformers in Webshell Detection 51
2.6.7 Transformers Applied to Code Analysis 52
○ A. CodeBERT / GraphCodeBERT 52
○ B. GPT-style Models for Code 52
○ C. Transformer + LSTM Hybrids 52
2.6.8 Limitations and Considerations 52
Conclusion of Section 2.6 53
2.7 Comparative Review of Existing Systems 53
○ 2.7.1 Signature-Based Detection Systems 53
○ 2.7.2 Heuristic and Rule-Based Detection 54
○ 2.7.3 Static and Dynamic Malware Analysis Frameworks 54

- ○ 2.7.4 ML and Deep Learning-Based Systems 55
- ○ 2.7.5 Comparative Summary Table 56
- ○ 2.7.6 Justification for the Proposed Work 56
- ○ Conclusion of Section 2.7 57

2.8 Chapter Summary 57

Chapter 3 58

DESIGN OF A DEEP-LEARNING WEB SHELL DETECTION FRAMEWORK 58

3.1 System Architecture Overview 58

- ○ 3.1.1 Architectural Components 58
- ○ 1. Data Acquisition and Corpus Scanning 58
- ○ 2. Tokenization and Vocabulary Encoding 59
- ○ 3. Embedding Layer 59
- ○ 4. BiLSTM Encoder 59
- ○ 5. Word-Level Attention Layer 59
- ○ 6. Dropout and Fully Connected Layer 60
- ○ 7. Output Interpretation and Visualization 60
- ○ 8. Model Training and Evaluation Loop 60
- ○ 3.1.2 System Flow Diagram (Conceptual) 61
- ○ 3.1.3 Design Justifications 61
- ○ 3.1.4 Deployment Considerations 61
- ○ Conclusion of 3.1 62

3.2 Data Ingestion and Tokenization 62

- ○ 3.2.1 Raw Data Acquisition 62
  - A. Malicious Dataset (Webshells) 62
  - B. Normal Dataset (Normal PHP Scripts) 62
- ○ 3.2.2 Directory Traversal and File Ingestion 63
- ○ 3.2.3 Cleaning and Normalization 63
- ○ 3.2.4 Tokenization of Source Code 64
  - Token Types: 64
  - Tools Used: 64
- ○ 3.2.5 Vocabulary Construction 64
  - Process: 64
- ○ 3.2.6 Sequence Encoding and Length Normalization 65
- ○ 3.2.7 Label Assignment 65
- ○ 3.2.8 Efficient Disk Storage Using Memory Mapping 65
- ○ 3.2.9 Sample Example 65
  - Raw Input: 65
  - Tokenized: 65
  - Token IDs (Example): 66
  - Final Padded Sequence (max_len = 15): 66
- ○ 3.2.10 Pipeline Summary 66
- ○ Conclusion of Section 3.2 66

3.3 Embedding Construction (Word2Vec, CodeBERT, etc.) 66

- ○ 3.3.1 Objectives of Embedding Construction 67

- 3.3.2 Word2Vec Embeddings 67
  - a) Overview 67
  - b) Training Details 67
  - c) Output 68
  - d) Advantages 68
- 3.3.3 CodeBERT Embeddings 68
  - a) Overview 68
  - b) Tokenizer and Model 68
  - c) Embedding Extraction 68
  - d) Advantages 69
- 3.3.4 Comparison: Word2Vec vs. CodeBERT 69
- 3.3.5 Integration into the Model Pipeline 69
  - Word2Vec Integration: 69
  - CodeBERT Integration Options: 70
- 3.3.6 Future Embedding Options 70
- Conclusion of Section 3.3 70

3.4 Model Definitions 70
- 3.4.1 RNN-Based Detector 71
  - a) Motivation 71
  - b) Architecture 71
  - c) Limitations 71
- 3.4.2 CNN-Based Detector 72
  - a) Motivation 72
  - b) Architecture 72
  - c) Benefits 72
  - d) Limitations 72
- 3.4.3 Transformer-Based Detector 72
  - a) Motivation 72
  - b) Architecture Overview 73
  - c) Benefits 73
  - d) Limitations 73
- 3.4.4 GRU + Attention-Based Detector 74
  - a) Motivation 74
  - b) Architecture 74
  - c) Benefits 75
  - d) Explainability 75
- 3.4.5 Summary Table of Model Properties 75
- Conclusion of Section 3.4 75

3.5 Training Strategy & Hyperparameter Optimization 76
- 3.5.1 Data Partitioning and Sampling Strategy 76
  - Split ratio=NtrainNtrain+Ntest=0.8 76
  - Stratified Splitting 76
- 3.5.2 Data Loaders and Batch Processing 76
- 3.5.3 Loss Function 77

- ○ 3.5.4 Handling Data Imbalance 77
- ○ 3.5.5 Optimizer 78
- ○ 3.5.6 Learning Rate Scheduling 78
- ○ 3.5.7 Regularization Techniques 78
- ○ 3.5.8 Early Stopping 78
- ○ 3.5.9 Model Checkpointing 79
- ○ 3.5.10 Evaluation Metrics 79
- ○ 3.5.11 Hyperparameter Optimization 79
- ○ 3.5.12 Cross-Validation (Optional for Future Work) 80
- ○ 3.5.13 Visualization and Logging 80
- ○ 3.5.14 Training Runtime and Resources 81
- ○ Conclusion of Section 3.5 81

3.6 Regularization and Early-Stopping 81
- ○ 3.6.1 Understanding Overfitting 82

3.6.2 Regularization Techniques Used 82
- ○ A. Dropout 82
  - Mathematical Formulation: 82
  - Implementation in Model: 82
- ○ B. Weight Decay (L2 Regularization) 83
  - Modified Loss Function: 83
  - Implementation: 83
- ○ C. Early Stopping 83
  - Mechanism: 83
  - Mathematical Criteria: 84
  - Implementation (from your code): 84
- ○ D. Batch Normalization (Optional) 84
- ○ 3.6.3 Combined Effect of Regularization Techniques 85
- ○ 3.6.4 Empirical Observations 85
- ○ Conclusion of Section 3.6 85

3.7 Evaluation Protocols and Metrics 86
- ○ 3.7.1 Evaluation Objectives 86
- ○ 3.7.2 Dataset Splitting and Cross-Validation 86
- ○ 3.7.3 Evaluation Metrics 86
  - A. Accuracy 87
  - B. Precision (Positive Predictive Value) 87
  - C. Recall (Sensitivity / True Positive Rate) 87
  - D. F1-Score 87
  - E. ROC Curve and AUC (Area Under Curve) 87
  - F. Confusion Matrix 88
- ○ 3.7.4 Threshold Tuning 88
- ○ 3.7.5 Experimental Protocol 89
- ○ 3.7.6 Comparative Evaluation 89
- ○ 3.7.7 Explainability (For Attention-Based Models) 89
- ○ 3.7.8 Performance Metrics on Obfuscated Samples 90

○ Conclusion of Section 3.7 90
3.8 Chapter Summary 90
Chapter 4 92
EXPERIMENTAL RESULTS & ANALYSIS 92
4.1 Experimental Setup and Hardware 92
○ 4.1.1 Hardware Configuration 93
○ 4.1.2 Software and Libraries 93
○ 4.1.3 Dataset Summary 94
○ 4.1.4 Preprocessing Pipeline 94
○ 4.1.5 Training Parameters 95
○ 4.1.6 Models Trained and Evaluated 95
○ Conclusion of Section 4.1 96
4.2 Dataset Description and Splits 96
○ 4.2.1 Data Sources 96
A. Malicious PHP Scripts (Webshells) 96
B. Normal PHP Scripts (Normal) 97
○ 4.2.2 Dataset Composition 97
○ 4.2.3 Preprocessing Pipeline Summary 98
○ 4.2.4 Dataset Split Strategy 98
Stratified Sampling: 99
Shuffling and Seed Setting: 99
○ 4.2.5 Token Distribution Analysis 99
○ 4.2.6 Example Tokenized Script 99
○ Conclusion of Section 4.2 100
4.3 Baseline Methods for Comparison 100
○ 4.3.1 Objectives of Baseline Evaluation 100
○ 4.3.2 Selected Baseline Models 101
A. Logistic Regression (LR) 101
B. Support Vector Machine (SVM) 101
C. Random Forest (RF) 102
D. Multilayer Perceptron (MLP) 102
○ 4.3.3 Feature Extraction for Baseline Models 103
○ 4.3.4 Implementation and Configuration Summary 103
○ 4.3.5 Performance Expectations 104
○ 4.3.6 Why Baselines Are Insufficient Alone 104
○ Conclusion of Section 4.3 104
4.4 Performance Metrics (Accuracy, Precision, Recall, F1, AUC) 105
○ 4.4.1 Confusion Matrix Foundation 105
○ 4.4.2 Accuracy 105
○ 4.4.3 Precision (Positive Predictive Value) 106
○ 4.4.4 Recall (True Positive Rate or Sensitivity) 106
○ 4.4.5 F1-Score (Harmonic Mean of Precision and Recall) 107
○ 4.4.6 Area Under the Receiver Operating Characteristic Curve (AUC-ROC) 107

- ○ 4.4.7 Metric Selection Justification 108
- ○ 4.4.8 Threshold Tuning Strategy 108
- ○ Conclusion of Section 4.4 109

4.5 Detection Results 109

- ○ 4.5.1 Overall Accuracy & Class-Wise Performance 109
  - Key Findings: 110
  - Confusion Matrix Example – GRU + Attention 110
  - Confusion Matrix Example – Logistic Regression 111
- ○ Conclusion of Section 4.5 111

4.6 Ablation Studies & Sensitivity Analysis 111

- ○ 4.6.1 Ablation Study of the GRU + Attention Model 112
  - Baseline Model: 112
  - Variants Evaluated: 112
  - Interpretation: 112
- ○ 4.6.2 Sensitivity to Embedding Dimension 113
- ○ 4.6.3 Sensitivity to Dropout Rate 113
- ○ 4.6.4 Sensitivity to Sequence Length 113
- ○ 4.6.5 Learning Rate Sensitivity 114
- ○ 4.6.6 Early Stopping Sensitivity 114
- ○ Conclusion of Section 4.6 115

4.7 Visualization of Learned Representations 115

- ○ 4.7.1 Token-Level Attention Heatmaps 115
  - Mathematical Formulation: 115
  - Heatmap Example – Malicious Script: 116
  - Heatmap Example – Normal Script: 117
- ○ 4.7.2 t-SNE Visualization of Word Embeddings 117
  - Cluster Observations: 117
  - Insights: 118
- ○ 4.7.3 Latent Space Separation of File Representations 118
  - Results: 118
- ○ 4.7.4 Saliency Maps for Decision Attribution 118
  - Benefits: 118
  - Example: 119
- ○ 4.7.5 Integrated Gradients for Robust Attribution 119
  - Use Case: 119
- ○ 4.7.6 Visual Dashboard (Future Work) 119
- ○ Conclusion of Section 4.7 119

4.8 Discussion of Findings 120

- ○ 4.8.1 Performance Superiority of Deep Learning Models 120
  - Key Insight: 120
- ○ 4.8.2 Impact of Attention Mechanism 120
- ○ 4.8.3 Robustness to Obfuscation 121
- ○ 4.8.4 Semantic Grouping in Embeddings 121
- ○ 4.8.5 Interpretability and Explainability 121

○ 4.8.6 Limitations and Edge Cases 122
○ 4.8.7 Deployment Considerations 122
○ 4.8.8 Contribution to Research and Industry 122
○ Conclusion of Section 4.8 123
4.9 Chapter Summary 123
Chapter 5 125
CONCLUSIONS & FUTURE WORK 125
5.1 Recap of Research Contributions 125
1. Development of a Novel GRU + Attention-Based Detection Architecture 125
2. Construction of a Clean and Malicious Webshell Dataset 125
3. Implementation of a Disk-Efficient Data Pipeline 126
4. Comparative Analysis of Deep Learning Models 126
5. Real-Time Detection Pipeline Design and Evaluation 127
6. Incorporating Explainable AI through Attention Visualizations 127
7. Security, Privacy, and Ethics: Addressing these issues in a holistic manner 127
9. Creation of an Internal and External Feedback Loop to Promote Ongoing Preparation and Learning 127
○ Summary of Contributions 128
5.2 Theoretical and Practical Implications 128
5.2.1 Theoretical Implications 128
1. Advancement in Sequence-Based Malware Detection 129
2. Incorporation of Attention Mechanisms into Security Applications 129
3. Multi-Model Benchmarking Framework 129
4. Developing a Real-Time AI System Architecture Design 129
5.2.2 Practical Implications 129
1. Real-Time Threat Detection Capability 129
2. Enhanced Threat Response and Forensics 130
3. Adaptability to New Threats 130
4. Compatibility with Existing Infrastructure 130
5. Privacy-Conscious Deployment 130
○ 5.2.3 Cross-Domain Impact Potential 130
○ Conclusion of Section 5.2 131
5.3 Limitations of the Current Study 131
○ 5.3.1 Dataset Limitations 131
5.3.2 Model Generalization Constraints 132
5.3.3 Inference and Resource Limitations 132
5.3.4 Deployment and Integration Challenges 132
5.3.5 Evaluation Scope 133
Conclusion of Section 5.3 133
5.4 Directions for Future Research 133
○ 5.4.1 Multilingual and Multi-Format Webshell Detection 134
Future Work: 134
○ 5.4.2 Hybrid Static-Dynamic Analysis Pipelines 134

Future Work: 134
○ 5.4.3 Graph-Based Structural Learning 135
Future Work: 135
○ 5.4.4 Meta-Learning and Few-Shot Adaptation 135
Future Work: 135
○ 5.4.5 Adversarial Attack Defense 136
Future Work: 136
○ 5.4.6 Federated and Privacy-Preserving Learning 136
Future Work: 136
○ 5.4.7 Real-Time Monitoring, Drift Detection, and Auto-Retraining 137
Future Work: 137
○ 5.4.8 Explainability and Analyst Tooling 137
Future Work: 137
○ 5.4.9 Benchmarking on Global Threat Intelligence Feeds 137
Future Work: 137
○ 5.4.10 Broadening to Non-Code Threat Vectors 138
○ Conclusion of Section 5.4 138
5.5 Final Conclusion 138
Key Takeaways: 139
Impact and Outlook: 139
Final Thought: 139
REFERENCES 141
APPENDICES 143
○ Appendix A: Model Architectures (PyTorch Definitions) 143
○ 143
○ Appendix B: Hyperparameter Configurations 144
○ Appendix C: Dataset Directory Structure 145
○ Appendix D: Sample Input Tokens and Attention Weights 145
○ Appendix E: Evaluation Metrics Formulas 146
○ Appendix F: ROC Curve 147
○ Appendix G: System Specifications 148
○ Appendix H: Acronyms and Abbreviations 149

Chapter 1
INTRODUCTION

In today's reality of digital age, web-based applications and platforms have become the core to business operations, communication, and constituent service delivery. Yet, as organizations continue to rely more and more on complex, interconnected systems and internet-facing applications to achieve their missions, the threat landscape has grown tremendously. Of all the many cyber threats out there, webshell attacks are easily one of the most insidious and impactful methods employed by criminal actors to achieve long-term and persistent access to web servers. These scripts, hidden deep inside seemingly innocent files, give attackers the potential to push arbitrary commands, steal sensitive data, or orchestrate further network-wide attacks under the radar.

This dissertation's primary contribution is in detecting such complex threats through novel computational processes, including particularly attention-based models in natural language processing (NLP) and machine learning. Through in-depth investigations into the nature of webshells, their activity, and their progression over time, the research sets the groundwork for creating an intelligent detection system that can identify even obfuscated or never-before-seen attack patterns.

The chapter is divided into five main sections, which set the stage, frame the research question, state the aims, describe the research design, and explain the outline and limitations of the dissertation. Our overall intention is to give a clear, compelling, and memorable entry point to the reader, and a high-level sense of what is to come, while being very clear about the purpose and direction of the subsequent research.

1.1 Background and Motivation

The spread of web technologies and services during the digital age has transformed communication, trade, and the exchange of information. Websites are an essential part of everyday life, ranging from a simple personal blog to advanced corporate-level applications. Their increasing popularity has also turned them into inviting targets for cybercriminals. Webshells are one of the most insidious and nefarious attack vectors among the different threats threatening online environments.

A webshell is an attack script or backdoor that an attacker uploads to the web server, normally using vulnerabilities in the file upload feature, unguarded code, or misconfigured servers. After it is deployed, it provides the attackers with remote access to the server with the ability to run arbitrary commands, handle files, privilege

escalation, data exfiltration, or initiate further network-based assaults. Webshells are commonly coded using popular script languages such as PHP, ASP, JSP, or Perl and are sometimes coded to mimic Normal files, rendering them extremely difficult to detect through standard security controls.

The incentive to detect webshells accurately and effectively is the drastic consequences that can befall organizations. An effective webshell attack can result in data breaches, service downtime, reputational loss, and monetary loss. Signature-based detection using pre-established rules or patterns is the traditional method of detection for webshells. Although efficient against known encountered webshells, the approach is ineffective against new or encoded versions. Similarly, static analysis techniques scanning for problematic code patterns are restricted by the inability to analyze runtime patterns and contextual dependencies.

With the evolving nature of attacks and the increasing sophistication of webshells through polymorphism, encryption, and the use of new encoding schemes, there is now an impetus for adaptive, intelligent, and automated detection that can scale and detect unknown threats, which has stimulated the use of machine learning and deep learning methodologies that can generalize well from large datasets and learn intricate patterns.

A very promising avenue is the employment of models inspired by Natural Language Processing (NLP). Similar to how NLP models decode and comprehend human language, these models can be made to decode and comprehend sequences of source code or HTTP requests. Word attention is particularly helpful in allowing the models to specialize on the most important portion of an input sequence, placing greater weights on features that contribute the most to the decision. This renders them well-positioned to pick up on very subtle anomalies in the code and spot malignant activities that standard techniques may miss.

In this work, the potential of attention-based deep learning models is exploited to construct an effective framework for webshell detection. We leverage the capability of Bidirectional Gated Recurrent Units (BiGRUs) and word-level attention mechanisms to learn from tokenized sequences of code and put dynamic weights on each element of the code. This improves the detection performance as well as the explainability of the detection process.

The motivation for the present study is threefold:

To counter the limitations of existing webshell detection methodologies.

To develop an intelligent model that can adapt to evolving threats without manual rule updates.

Help build the wider area of cybersecurity using state-of-the-art NLP methodologies to address an entirely new and challenging domain. Ultimately, the direction of this work is to prove that the integration of deep learning architectures with NLP-based models has the potential to revolutionize the functionality of automated webshell detection tools towards achieving safer and more secure web infrastructures.

1.2 Research Problem Formulation

The rapid growth of the web has led to an exponential rise in the quantity and sophistication of web applications. With the growth, there is an equivalent growth in security vulnerabilities but specifically at the web server layer. Perhaps the most subtle and hard-to-detect malady in the new landscape is the webshell, a malicious script planted on the web server to obtain unauthorized access, steal data, run commands remotely, or manage system resources. These scripts are habitually camouflaged as legitimate files (like .php, .jsp, .asp, or even .html), and the attackers can use the server to their advantage without being noticed for extended durations.

The Complexity of Webshell Detection
Webshells do not have the same structure. Usually polymorphic, they adapt with every iteration to evade detection by signature-based defenses.

- Obfuscated, in which confusing or encoded form the code is written to stop analysis.
- Simple, opening the backdoor with a few lines of code.
- Encrypted or encoded, which makes static analysis very challenging.

These features pose a strong challenge to conventional security measures, such as:

- Signature-based detection that is based on known patterns and can be easily circumvented through code modification.
- Rule-based detection systems tend to use static heuristics and are prone to high rates of false positives.
- Time-intensive, prone to errors, and not scalable

With the persistent innovation and experimentation on the side of the attackers, static and heuristic detection methods are left behind. Therefore, the demand arises for something adaptive, intelligent, and strong enough to learn and categorize code, despite never having encountered the specific variant.

Reframing the Issue with Deep Learning
To overcome these limitations, the webshell detection issue is recast as an NLP sequence classification problem. Source code, much as natural language, consists of syntactically meaningful sequences, statements, expressions, loops, functions, that constitute an organized document. By considering webshell scripts as "sentences" or as "documents," one can take advantage of the techniques of deep learning that have transformed the area of NLP.

The main concept is to employ a deep learning model, in the form of a Bidirectional Gated Recurrent Unit (BiGRU) with word-level attention, to capture webshell code's sequential and contextual characteristics. This design allows the model to:

- Capture dependencies forward and backward within the sequence (critical to understanding how variables, functions, or commands are related to one another).

- Give different weights to various tokens or code snippets, enabling the model to pay greater attention to questionable elements of the script.
- Learn from data instead of using rigid rules, and enhance adaptability and generalizability.

Why Traditional ML Fails

Classical machine learning methods, including Support Vector Machines (SVMs), Random Forests, and Naïve Bayes classifiers, are employed to detect webshells but fail for the reasons described below:

- They treat features independently and ignore sequential relationships.
- They tend to involve labor-intensive feature engineering, which is hard and domain-dependent.
- They are not as effective at processing code obfuscation or encrypted payloads.

In contrast, deep learning architectures, particularly sequence data models, can automatically learn hierarchical and semantic features from unprocessed code, enhancing detection performance and minimizing the need for domain knowledge.

Problem Statement

This paper presents an original solution to the webshell detection problem through the design of an end-to-end deep learning-based classification framework. The formal formulation of the problem is:

Having access to a set of labeled web application files consisting of Normal and webshell (malicious) scripts, develop and implement a deep learning system that can accept input as sequences of code, apply attention to concentrate on meaningful patterns, and decide whether to label an input file as Normal or malicious.

This one problem subsumes many sub-problems:

- Preprocessing and representation of data: Effective tokenization of code, handling variable-length sequences, and converting tokens into an appropriate numerical space
- Model design: Building an architecture that is both performance-oriented and interpretable. BiGRU+Attention is selected based on its efficiency and context understanding.
- Training and testing: To make sure the model obtains meaningful representations using strong datasets, along with performance metrics including accuracy, precision, recall, and F1-score.
- Interpretability: Leveraging attention weights to indicate which pieces of code had the highest influence on the decision, shedding light on the model's behavior and building confidence in automated decision-making.

Research questions

This results in the creation of major research questions:

- How does one best represent webshell scripts as sequences for deep learning models using NLP?
- The best architecture (coupling BiGRU with attention mechanisms) for identifying malicious patterns in code is.
- How does the performance of the new model contrast with the conventional detection strategies in accuracy, false positives, and flexibility?
- Yes, the attention mechanism can indeed give the model's decisions meaningful interpretability.

Scope of the Problem

The new approach centers on server-side scripts that are typically targeted for webshell implantation, which include PHP, JSP, and ASP files. The system is learned and validated using a dataset curated to include known webshells as well as clean examples. The model is then assessed not only on detection accuracy but also on:

- Generalization to novel threats
- Resistance to obfuscation
- Scalability for real-time system deployments

Broader applications

Solving the problem will significantly advance the area of cybersecurity by:

- Decreasing the manual labor involved in code review and threat detection.
- Intelligent code analysis for strengthening threat intelligence systems
- Offering an automated, scalable solution that automatically adjusts to new threats through data-driven learning.

1.3 Aims and Objectives

This work aims primarily to design, develop, and evaluate a deep learning-based framework using Semantic and sequential properties of code, leveraging Natural Language Processing (NLP) approaches for the detection of webshells in web applications. By introducing a more adaptive, context-aware, and interpretable approach using a BiGRU (Bidirectional Gated Recurrent Unit) network enhanced with a word-level attention mechanism, the work aims to overcome the constraints of conventional static and heuristic-based detection systems. This work aims to investigate model interpretability, resilience against obfuscation, and generalization capacity to identify hitherto undetectable or polymorphic webshell variants in addition to raising detection accuracy.

○ Objectives of the Study

To achieve this aim, the study is guided by the following specific objectives:

1. Data Collection and Preprocessing

- To compile and select a large collection of webshells and safe server-side scripts (such as PHP, JSP, and ASP).
- To create a reliable pipeline for code tokenization and normalization that can handle source code as sequential data.
- To use word embeddings (such as Word2Vec or GloVe) to create a numerical embedding representation of code tokens, making sure semantically similar tokens have similar representations.

2. Model Development and Optimization

- To create a Bidirectional GRU (BiGRU) architecture that recognizes the code sequence's forward and backward dependencies.
- To incorporate a word-level attention mechanism to recognize and highlight suspicious or malicious code segments in a given file.
- To balance detection accuracy and computational efficiency by training the model on labeled data with optimized hyperparameters.

3. Evaluation and Benchmarking

- To assess the suggested BiGRU+Attention model's performance using the following common metrics: F1-score, accuracy, precision, recall, and Area Under the Curve (AUC).
- To evaluate the model against other deep learning models (CNN, LSTM) and baseline classifiers (SVM, Random Forest, Logistic Regression).
- To evaluate the model's resistance to popular evasion strategies like encoding, code obfuscation, and the random insertion of tokens that appear harmless.

4. Interpretability and Explainability

- To make use of the attention mechanism to offer analytical or visual insights into the code segments that most influenced the detection decision.
- To create a framework for human-interpretable explanations that can help cybersecurity analysts examine the results of automated detection.

5. Real-World Applicability

- To evaluate the trained model's scalability and deployment viability in an intrusion detection system (IDS) or web application firewall (WAF) operating in real-time or near-real-time.
- Testing if a model trained on one kind of script (like PHP) can successfully identify webshells in other script types (like JSP, ASP) is one way to investigate cross-language generalization.

6. Contribution to Research and Industry

- To publish methods, findings, and insights that improve web application security to add to the body of knowledge in academia.
- To offer useful suggestions and a possible open-source webshell detection tool that can be incorporated into security platforms.

The study intends to close the gap between theoretical cybersecurity research and workable, deployable solutions that can defend contemporary web infrastructures against sophisticated persistent threats like webshells by methodically addressing these goals.

1.4 Scope and Delimitations

The scope of this research encompasses the design, development, and evaluation of a webshell detection model that leverages word attention mechanisms. Webshells, often used by attackers to remotely control compromised web servers, present serious threats to cybersecurity. Due to their stealthy nature and tendency to mimic legitimate code, traditional detection techniques frequently fail to identify them effectively. This work intends to overcome these constraints by using an attention-based deep learning method, emphasising the semantic relevance of words in online-based scripts.

The proposed approach emphasises building a model that uses word attention mechanisms to capture complex textual patterns, so stressing and prioritising suspicious elements in code. This method lets the model learn subtle characteristics that set Normal scripts apart from malicious webshells, so lowering false positives and increasing detection accuracy.

Scope of the Study:

- Since PHP is still one of the most targeted and used languages for web-based attacks, the study is mostly on webshells written in PHP.
- Known webshells and Normal PHP scripts obtained from publicly available repositories and security forums form part of the dataset used in this work.
- Data preprocessing, feature extraction, model training, performance assessment, and comparison with baseline detection techniques constitute part of the scope.
- Deep learning architectures, especially those using word-level attention mechanisms, as well as natural language processing (NLP) approaches, are applied in the model.
- Model performance is measured using evaluation criteria, including accuracy, precision, recall, F1-score, and false positive rate.

Delimitations of the Study:

While this research aims to provide significant improvements in the detection of webshells, it is important to note the following delimitations:

- Though the approach could be modified in the next work, the study does not cover webshells produced in languages other than PHP (e.g., ASP, JSP).
- This version of the model addresses only half of the obfuscation methods attackers use, including dynamic code generation or encryption.
- The model may need further tuning for deployment in real-world systems with varied server configurations and traffic patterns, It is trained and tested in a controlled environment.

- The study does not centre on post-detection response systems, including system recovery or webshell removal.

Our goal is to guarantee that the results are both significant and practically applicable within the specified context by precisely defining the parameters of this study. Based on this framework, future research can expand the model's applicability to different programming languages, incorporate real-time detection features, and strengthen its resilience to sophisticated evasion techniques.

1.5 Research Methodology

The research design, the heart of any scientific study, was sorely lacking and vague. It outlines the research methodology and explains the choice of techniques, instruments, and approaches used. This work proposes an advanced webshell detection model utilizing a word attention mechanism embedded into a deep learning structure. Research design, data collection/preprocessing, model building/training processes, assessment, and result interpretation are some of the key factors that comprise the methodology. Each stage has been meticulously staged to ensure precision, repeatability, and best operation in real-world scenarios.

○ 1.5.1 Research Design

This study employs supervised machine learning in a unique analytical-experimental design. The key idea is that webshells produce unique word fingerprints that an attention-based architecture can learn to detect. To prove this hypothesis right:

Malicious and Normal PHP scripts are compiled into a dataset. From that data, a model for deep learning is developed, then trained. A combination of statistical and visual techniques is then utilized to evaluate the model’s performance. Performance is assessed against a standard of like comparators.

These results are iteratively refined through this cyclic, iterative design, which provides a solid foundation for a detection framework that is both repeatable and scalable.

○ 1.5.2 Data Acquisition and Description

The effectiveness of any machine learning model is largely determined by the quality and richness of the datasets on which they are trained. In this study, datasets are meticulously curated from several trusted sources to provide the most accurate picture.

Evildoers Webshells

Cumulative count of studied sources. Source: GitHub (public webshell repositories), VirusTotal, VX Heaven, Hybrid Analysis. Examples include China Chopper, b374k, WSO, FilesMan, and custom obfuscated shells. Behaviors: Rely heavily on obfuscated code, dynamic function names, base64 encoded payloads, and evasive I/O activities.

Typical PHP Scripts

Provided by: Open source CMS environments such as WordPress, Joomla, Laravel-based projects, eCommerce platforms, blogs.Well-formed, object-oriented, and semantically rich PHP code.

To prevent any bias from the model, a fine line was danced to make sure that the Normal scripts were just as unique in functionality and length as the malicious scripts.

○ 1.5.3 Data Preprocessing

The raw scripts undergo several preprocessing steps:

1.5.3.1 Cleaning

- Remove HTML tags, JavaScript blocks, and embedded styles.
- Normalize white space and remove comments.
- Standardize character encoding to UTF-8.

1.5.3.2 Tokenization and Parsing

- Each script is tokenized into PHP keywords, variable names, function calls, literals, etc.
- Tokenization is language-aware to capture PHP-specific constructs.
- A custom tokenizer is created using Python and the `php-parser` library.

1.5.3.3 Vectorization

- Tokens are mapped to integer indices using a vocabulary dictionary.
- One-hot and word-embedding techniques are experimented with.
- Embeddings are initialized using pre-trained GloVe or Word2Vec vectors trained on programming languages, if available.

1.5.3.4 Padding and Truncation

- Sequences are padded to the same length (max 1000 tokens) to facilitate batch processing.
- Truncation ensures long files do not dominate model capacity.

○ 1.5.7 Tools and Technologies Used

Tool	Purpose
Python	Primary programming language
TensorFlow / PyTorch	Deep learning framework

scikit-learn	Evaluation and metrics
Pandas / NumPy	Data processing
Matplotlib / Seaborn	Visualization
php-parser	Tokenization of PHP code
Jupyter Notebook/Google colab	Experiment tracking
Git	Version control

○ 1.5.8 Ethical Considerations

- No proprietary or confidential code is used.
- Malware samples are used in an isolated environment.
- Scripts are analyzed statically; no execution is performed on production systems.

○ 1.5.9 Summary

This methodology lays a robust foundation for designing an intelligent, transparent, and efficient webshell detection system. By integrating attention mechanisms, the system not only enhances detection rates but also provides interpretability, essential for cybersecurity practitioners.

1.6 Key Contributions

This research impacts both the field of web application security and machine learning by providing a new approach to detect PHP-based webshells via a word attention-based deep learning model. The main contributions of this dissertation are twofold.

○ 1.6.2 Engineering an interpretable injected webshell detector with deep learning

One of the major contributions is the design and implementation of a new deep learning model, leveraging a Bi-LSTM architecture with a word-level attention mechanism to detect malicious webshell code hidden in PHP scripts. By showing which tokens (function calls like eval, system, base64_decode, etc) contributed to its classification decision, this approach, unlike in black-box traditional models, boosts not only accuracy but transparency. As a result, at the foundational level, the model is

dependable and digestible to understand and operationalize on a day-to-day basis in a real security operation.

○ 1.6.2 Creating a Curated and Balanced Malicious/Normal PHP Scripts Dataset

A carefully crafted dataset is collected, which includes real-world malicious webshell samples from public malware repositories and a rich variety of Normal PHP scripts collected from industry-standard web platforms (e.g., WordPress, Laravel). The final dataset included

Contains more than 10,000 examples with accurate labels. Families and obfuscation level of webshells covered. It is transparently reproducible (within the bounds of security compliance).

As the largest, labeled dataset of its kind, it's a great resource for training, testing, and benchmarking webshell detection systems.

○ 1.6.3 Application of Word-Level Attention to Cybersecurity Use Cases

Given that PHP SDK malware detection is typically a static, keyword-centric process, this work represents the first application of a word attention mechanism to the field of PHP malware detection. Image Credit: The Transformer model. In-depth look at the attention mechanism.

Gives interpretability by providing importance weights for tokens.Provides transparency by allowing visual inspection and understanding of which code segments trigger suspicion. It is possible to generalize this to different domains like JavaScript malware, email phishing detection, and command injection.

This leads to a theoretical and methodological innovation by introducing NLP-inspired attention models into software security.

○ 1.6.4 Comparative Analysis with Conventional Machine Learning and Static Analysis Approaches

To prove the superiority of the proposed model, an extensive comparative study has been performed against:

Classic classifiers: Support Vector Machines (SVM), Decision Trees, Naive Bayes. Static heuristics (signature-based detection tools run at download time, such as ClamAV and regex filters). Previously trained deep learning models are devoid of attention.

The proposed model always achieves better precision, recall, F1-score, and interpretability than all baselines, setting a new state-of-the-art in webshell detection.

○ 1.6.5 Developing a Lightweight and Deployable Detection Prototype

A lightweight, Python-based prototype has already been generated that allows users to:

Perform static analysis on PHP files. Draw attention to odd code paths or tokens. Output the corresponding prediction label (malicious/Normal) along with a confidence score. Impact the relative security of platforms through integration with security tooling like Web Application Firewalls (WAFs) or CI/CD pipelines for DevSecOps.

This tangible deliverable truly closes the gap between lab-based research and on-the-ground deployment.

○ 1.6.6 Open Research Avenues for Security-Focused NLP Models

Through the integration of cybersecurity with the cutting edge of NLP, this research serves to set the stage for future investigation into:

Transformer-based models show promise for malware detection, as several studies have demonstrated state-of-the-art results using such models. Multi-language webshell detection (e.g., ASP, JSP, Python).Continuous monitoring with real-time alert systems integrated with code editors and version control systems.

The pipeline and learnings from this study can be applied to a variety of static code analysis tasks within the broader field of cybersecurity.

The previously discussed studies leading to DNN-based security solutions could not explain their predictions.

While most deep learning methods in security are haunted by the black-box issue, this works as follows:

Addresses it specifically through attention weights. Contains advanced, intelligent visualization tools to speed the investigation process and improve the efficiency of their security analysts. Fosters public confidence in AI-driven detection by ensuring that the decision-making process is always auditable.

This is another huge move in the direction of getting AI models usable and explainable to cyber professionals.

○ 1.6.8 Empirical Contribution via Evaluation Metrics and Visualization Tools

This research offers several important contributions:

Following this qualitative assessment, we provide a more thorough evaluation of deep learning features compared to conventional metrics using both standard and advanced metrics (ROC-AUC, PR curves). Example for visualization of attention maps and heatmaps displaying token-level decisions.Reduction of false positives/negatives, aiding researchers in the identification of false trending edge cases.

These tools facilitate a more nuanced understanding of model performance and limitations, protecting public trust while fostering innovation both in research and in deployment.

○ 1.6.9 Contributions to Research and Education

Their findings, codebase, and dataset will be made publicly available (with security protections), functioning as

An important reference for future academics looking to leverage NLP for malware detection.Learning content for transdisciplinary training in cybersecurity, AI, and software engineering.

This is to ensure the future sustainability and impact of the work beyond the limited scope of this dissertation.

1.7 Thesis Organization

This dissertation is organised into five thorough chapters. To provide a cogent and thorough examination of the subject, each chapter is intended to build gradually on the one before it: "***Deep Learning–Based Webshell Detection and Prevention Using Word-Level Attention Mechanisms***." The organisation aims to offer depth in analysis, rigour in methodology, and clarity in presentation.

○ Chapter 1: Introduction

The research is laid out in the introductory chapter. It starts with a summary of how webshells are becoming more common in cyberattacks and how their obfuscation techniques make detection difficult. The rationale for selecting a word attention-based deep learning approach is then presented. Clearly stated goals and objectives are presented after a formal research problem statement. The study's delimitations and scope are described, with an emphasis on the parameters that govern its execution. The methods utilised for model development and validation are included in the description of the research methodology. To emphasise the uniqueness and importance of the research, the main contributions of this study are outlined. This section concludes by outlining the thesis's chapter-by-chapter structure.

○ Chapter 2: LITERATURE SURVEY & THEORETICAL FRAMEWORK

This chapter provides the state of the art literature survey and related work in the field of: Real-life webshell attack techniques and how they are developing, Typical detection techniques (signature-based, heuristic-based, static analysis, and dynamic analysis). Today’s machine learning applications are used to improve cybersecurity. Deep learning and deep sequence modeling with RNNs, LSTMs, and GRUs. A huge development spurring on the explainable AI (XAI) movement’s grassroots push and helping found neural net-powered attention mechanisms. The shortcomings of existing solutions, especially evident in addressing obfuscated or polymorphic code.

A critical gap analysis follows, outlining limitations in existing methodologies. This, we think, provides a compelling rationale for the approach we’re proposing. Significant

comparative studies and international benchmarks are examined to develop the theoretical underpinnings for the research.

○ **Chapter 3: DESIGN OF A DEEP-LEARNING WEB SHELL DETECTION FRAMEWORK**

This chapter outlines the critical step of dataset development, emphasizing.

- The sources of Normal and malicious PHP scripts.
- Challenges in obtaining labeled and diverse datasets for webshells.
- Techniques for ensuring class balance and representation.
- Preprocessing steps such as.
    - Removal of noise and non-executable text.
    - Tokenization (including discussion on different token types).
    - Word embeddings (Word2Vec, GloVe, or custom embeddings).
    - Handling code obfuscation (string manipulation, base64 decoding, etc.).

Visualisation and statistical summaries of the dataset are provided to illustrate class distribution, token frequency, script length analysis, and lexical diversity.

○ **Chapter 4: EXPERIMENTAL RESULTS & ANALYSIS**

The core technical contribution of this research is detailed in this chapter. Key components include.

- Design and rationale for using Bi-directional LSTM (BiLSTM).
- Architecture of the attention mechanism at the word level.
- Integration strategy of attention with BiLSTM for sequence classification.
- Diagrammatic representation of model flow.
- Detailed explanation of.
    - Embedding layers.
    - Recurrent layers (forward and backward LSTMs).
    - Attention score computation and weight assignment.
    - Dense layers and final prediction output.

Additionally, this chapter includes.

- Model compilation settings.
- Hyperparameter tuning strategies.
- Loss function and optimizer choice.
- Regularization and dropout techniques.
- Use of frameworks (TensorFlow, PyTorch).

This section may also cover an ablation study to assess the impact of each model component.

○ **Chapter 5: Conclusion and Future Work**

This final chapter summarizes the entire research journey and highlights the key outcomes. It revisits the initial problem, objectives, and contributions and reflects on how the proposed approach addresses the research gaps. The novelty of combining attention mechanisms with BiLSTM for code-level classification is emphasized.

Future work directions may include.

- Exploring transformer-based models (e.g., BERT, CodeBERT) for code analysis.
- Extending the model to multi-class classification for different malware types.
- Real-time system development and integration with cloud services.
- Adversarial training for improved resilience.
- Expansion of the dataset to include non-PHP webshells.

Chapter 2
LITERATURE SURVEY & THEORETICAL FRAMEWORK

As the cybersecurity threat environment continues to change, bad actors leverage new methods to exploit system weaknesses and compromise systems. Among them, webshell attacks have been recognized as a significant threat. Recent studies and industry reporting have indicated a growth in the use of webshells as a function of remote control, persistent presence, and lateral movement within targeted networks.

In this chapter, we provide comprehensive literature reviews on WFDM systems by covering fundamental and recent related works. We analyze the theoretical foundations of different detection mechanisms, such as static analysis, pattern-matching, and advanced deep learning and attention-based models, to the extent allowed by the presented data, and evaluate the potential and the reliability of each. By integrating ideas from literature, this chapter provides the basis for the proposed approach, connecting existing research works with the proposition of the dissertation.

It is designed to present the reader with a step-by-step progression from an overview of webshell cyberattacks to overall theoretical concepts and technical approaches to resist webshells. It further reveals subtle issues and limitations with existing methods and highlights the requirement for more intelligent, resilient, and flexible systems to detect obscure and dynamic attacks.

2.1 Overview of Webshell Attacks

Webshell attack is a serious, harmful, and long-lasting security menace, especially for web servers, cloud-hosted applications, and Content Management Systems (CMS) in the cyber world. Webshell is a script or executable file written in a web scripting language, which can be uploaded to a site to enable remote administration of the machine.

Whereas conventional malware focuses on the user-side device, webshells are left on the server side, and provide attackers with remote access to an infrastructure over a long period with stealth. Stealthily and enabled with the capability to copy real scripts or merge with harmless web codes, web-based malware is so Normal that it is almost impossible for traditional signature-based or heuristics methods to identify.

2.1.1 What is a Webshell?

A webshell is usually a small file uploaded to a website by an attacker to enable remote access, edits, and uploads, such as web defacement or backdoors. The rewards of figuring out how to upload a web shell via a vulnerable upload function are that an attacker could upload such a file and gain arbitrary code execution inside a compromised system by visiting the webshell URL. It's running also gives attackers a web interface where to:

- Browse and manipulate files.
- Upload and download data.
- Install backdoors or rootkits.
- Execute server-side commands.
- Trying other parts of the network.

Example: A popular example would be the webshell "**China Chopper**", which is known for being small in size (< 4KB) but still provides many features, including file operation, command execution, and database operation.

2.1.2 Persistent and Evolution

There has been a lot of evolution in web shells in the last decade. The early ones were primitive, command-line driven, but today's webshells have a graphical interface, encryption, privilege escalation, and anti-forensic features. The fact that they can remain on systems, hidden as the case may be, on other systems (or also built into media files) is especially troubling for incident response teams.

Attackers can also leverage steganography, encoding formats (base64, etc.), and polymorphism to add another layer of confusion to injected code. As a consequence, determining if a script is Normal or malicious is not easy even for experts.

2.1.3 Common Types of Webshells

Webshells are broadly categorized based on the scripting language and delivery mechanism.

Type	Description
PHP Webshells	Most prevalent due to PHP's popularity, used in WordPress, Joomla, etc.
ASP/JSP Webshells	Target enterprise-grade servers using the Microsoft or Java stack
Python/Ruby Webshells	Less common but powerful when targeting APIs and modern web stacks
Command-based	Purely CLI-based shells for stealth
Graphical Webshells	Interface-rich tools for ease of use by low-skilled attackers

2.1.4 Techniques for Deploying Webshells

1. **File Upload Vulnerabilities:** Exploiting insecure file upload functionalities to upload a malicious script.
2. **Command Injection:** Injecting commands into poorly validated input fields.

3. **Remote File Inclusion (RFI):** Loading remote scripts through vulnerable parameters.
4. **Brute-force Attacks:** Gaining access via weak or default credentials.
5. **Misconfigured Permissions:** Writing to directories that should be write-protected.

Once deployed, attackers often install additional backdoors or create scheduled tasks to regain access even after initial detection.

2.1.5 Real-World Examples and Incidents

Several high-profile incidents have involved webshell-based attacks:

- **Microsoft Exchange ProxyShell Vulnerability (2021):** Attackers exploited vulnerabilities to install webshells on thousands of servers worldwide, gaining full remote access.
- **SolarWinds Orion Breach (2020):** Involved sophisticated backdoors and webshells to maintain persistence in government networks.
- **Webhosting Control Panels (e.g., cPanel, Plesk):** Frequently targeted by attackers uploading hidden webshells in shared hosting environments.

These incidents demonstrate how webshells are not only used by script kiddies but also by **Advanced Persistent Threats (APTs)** and **nation-state actors**.

2.1.6 Challenges in Detecting Webshells

- **Obfuscation:** Use of encoded strings, variable functions, or runtime evaluation (e.g., `eval(base64_decode($_POST['x']))`).
- **Polymorphism:** Changing the structure or syntax to avoid detection.
- **Small Footprint:** Many webshells are under 5 KB in size and mimic harmless scripts.
- **Encrypted Communication:** Use of HTTPS or encoded POST parameters to avoid traffic inspection.
- **Mimicking Legitimate Behavior:** Commands executed through GET or POST requests appear like normal user behavior.

These challenges necessitate the use of **intelligent, behavior-aware models**, such as those powered by **deep learning and attention mechanisms**, to identify latent patterns that traditional techniques miss.

2.1.7 Need for Advanced Detection Techniques

With the rapid growth of web-based applications, traditional rule-based or static analysis methods fall short. There is a growing shift toward **machine learning (ML)** and **deep learning (DL)** based approaches, which analyze the syntactic and semantic structure of scripts. Among these, **sequence models with attention mechanisms** (e.g., BiLSTM + Attention) have emerged as powerful tools for capturing contextual relationships in code.

The use of **word attention mechanisms**, inspired by natural language processing (NLP), helps highlight important tokens or commands in the script that may indicate malicious intent, even when obfuscated or out of order.

- **Conclusion of Section 2.1**

Webshells represent a covert, persistent, and highly adaptive threat to digital infrastructure. Their evolution and sophistication demand equally advanced detection mechanisms. As the literature reveals, leveraging NLP-inspired models and attention-based architectures holds significant promise in this domain. The next section will explore the range of traditional and contemporary detection techniques and evaluate their effectiveness against this evolving threat.

2.2 Static Analysis Techniques

Static analysis, which is defined as the analysis of code or program structure independent of (or before) execution. Static analysis techniques, which can be integrated into webshell detection, are generally favored as they are non-intrusive, exhibit low runtime overhead, and are applicable in early detection, not having malicious scripts on the server. These approaches inspect the source code, bytecode, or binary formats of scripts to identify suspicious behavior indicative of maliciousness using heuristic rules, syntax patterns, or statistical characteristics.

Static analysis has been a useful first line of defense in many security tools. Its effectiveness is doomed to be evermore undercut by evasion techniques attackers use, like code obfuscation, dynamic function calls, and polymorphism.

2.2.1 Overview of Static Analysis in Web Security

Static analysis tools operate by parsing code files to extract features such as:

- Variable names and function definitions.
- API usage patterns.
- Control flow graphs.
- Abstract syntax trees (ASTs).
- Code entropy (used in detecting obfuscation).

These features are then evaluated using either manually crafted rules or automated classification techniques to flag suspicious code.

For webshell detection, static analysis can help identify.

- Suspicious PHP functions (`eval()`, `exec()`, `system()`, `passthru()`).
- Base64 encoded strings or obfuscated payloads.
- Dynamic code execution patterns (e.g., `preg_replace()` with `/e` modifier).
- Webshell signatures using regex or YARA rules.

2.2.2 Common Static Analysis Techniques

2.2.2.1 Signature-Based Detection

One of the most common and oldest forms of static detection is **signature matching**. In this method, known patterns of malicious code are stored as signatures, regex patterns, and the code under inspection is scanned for these.

Pros:

- Fast and lightweight.
- Effective against known threats.

Cons:

- Fails against new or mutated webshells.
- Easily bypassed with slight modifications, replacing function names with variables.

2.2.2.2 Rule-Based Heuristics

Heuristics use logical rules based on experience or expert knowledge.

- It flags if a script uses `eval()` on a base64 decoded string.
- If variable functions are used (`$func = "eval"; $func($input);`), it suspicion.

While more flexible than signatures, heuristic systems often generate **false positives** due to the over-generalization of rules, especially in complex applications where some flagged constructs are legitimate.

2.2.2.3 Syntax Tree Analysis

Abstract Syntax Trees (ASTs) represent the hierarchical structure of code. They are used to understand code behavior by analyzing function call dependencies, variable scope, and branching logic. Malicious scripts often show distinct AST structures, such as.

- Short scripts with one deeply nested function.
- Repetitive and redundant code paths.
- Multiple encoding/decoding layers.

Tools like **Esprima (for JavaScript)** or **php-parser (for PHP)** can be used to generate and traverse ASTs.

2.2.2.4 Entropy-Based Analysis

Entropy as a measure of randomness. Entropy is a measure of uncertainty. Malicious scripts typically employ encoded payloads with high entropy (base64, hex encoding). Static analysis tools are already available that can flag files or blocks of code over a set entropy threshold. Unfortunately, this heavy-handed approach is prone to false positives

on otherwise legal scripts that use encryption or binary data processing as a regular function.

2.2.2.5 YARA Rules

YARA is a pattern-matching tool used to describe and identify malware. Analysts write custom YARA rules that detect.

- Specific strings or byte patterns.
- Known shell command structures.
- Unusual file operations.

YARA is widely used in antivirus engines, incident response, and malware forensics, but like signatures, it depends heavily on prior knowledge of malware.

2.2.3 Tools and Frameworks for Static Analysis

Tool	Description
ClamAV	Open-source antivirus engine using static signature matching
RIPS	PHP static code analysis tool to find vulnerabilities, including webshells
SonarQube	Static code analysis for code quality and security
php-parser	Library for parsing PHP code into ASTs
YARA	Pattern matching tool for malware detection
Binwalk	Used in analyzing binary webshells or obfuscated payloads

2.2.4 Limitations of Static Analysis Techniques

Despite its speed and early-warning capabilities, static analysis suffers from several critical limitations.

1. **Obfuscation Resistance:** Modern webshells use sophisticated encoding and obfuscation strategies (string concatenation, dynamic variables), rendering static rules ineffective.
2. **Lack of Context Awareness:** Static analysis cannot understand execution flow or runtime conditions (user-specific payloads, environmental dependencies).
3. **High False Positive Rate:** Legitimate applications using cryptography or dynamic content generation may be misclassified.

4. **Language-Specific Barriers:** Static analysis tools are often tailored for one language and fail when scripts involve multiple languages or embedded components (PHP + JavaScript).
5. **Limited Scalability for Unknown Threats:** Static rules are reactive; they detect only what has been predefined and fail to generalize to novel threats.

2.2.5 Motivation for Moving Beyond Static Analysis

Even with these limitations, researchers and practitioners have begun augmenting traditional detection practices with machine learning and deep learning methods. These methods take source code as a sequence or a graph and learn latent features in an unsupervised manner, breaking free from hand-crafted rules. In combination with attention mechanisms (word attention in NLP), models are given a higher level of interpretability and precision, as they can identify and focus on the most relevant parts of code, something traditional, static tools cannot do.

○ Conclusion of Section 2.2

Static analysis is the basis for many of the first-line malware and webshell detections. Great for old, pervasive, static threats. This isn't very effective in a more dynamic web where evasion techniques are the bread and butter of your malicious adversaries. This requires a move to dynamic, context-sensitive methods like deep learning with attention, which is capable of surmounting the static model's built-in flaws. The following section looks at dynamic analysis approaches, which fuse theory and practice by trying to get a sense of how code will behave during execution.

2.3 Dynamic Analysis Techniques

Dynamic analysis is the execution or simulation of the program within a security testing environment to monitor its behavior in real-time. Within the scope of webshell detection, dynamic analysis becomes an essential component to establish runtime behavior, side effects, and context-aware malicious activity that static analysis does not capture under potential obfuscation or code complexity.

Webshells are specifically crafted to appear normal when being examined statically, not when executing. They might implement conditional execution, runtime decoding, or environment-aware triggers that only trigger in particular scenarios (e.g., on receipt of a certain HTTP parameter). Dynamic analysis can uncover such behavior by executing code with simulated real-world user inputs and system output.

2.3.1 What is Dynamic Analysis?

Dynamic analysis is the real-time monitoring of code execution to capture the real behavior of the system being examined. This is almost always done in a sandboxed or isolated environment to avoid accidentally causing harm to production systems. A few key goals are.

Viewing activity on the file system. Eavesdropping on the enemy's network chatter. Tracing Proxy, Capturing system calls and subprocesses. Tracing I/O operations as a log stream. Recognizing payloads after the fact, decoded at runtime.

For webshells, this means faking many HTTP POST/GET requests a web server would receive, and inspecting the response from the server script, such as how it handles payloads, executes shell commands, or modifies files.

2.3.2 Techniques Used in Dynamic Analysis

Dynamic detection strategies range from manual testing to automated instrumentation. Below are the most common dynamic analysis techniques relevant to webshell detection.

a. Behavior Monitoring

This involves running the suspect script on a test server and monitoring its actions, such as.

- File creation, deletion, or modification.
- Access to sensitive system files (`/etc/passwd`, `.htaccess`).
- Network connections (reverse shells or callbacks).
- Execution of system-level commands (`cmd.exe`, `bash`, `powershell`).

Tools like **Sysmon**, **Process Monitor**, and **auditd** (on Linux) are used to monitor these operations.

b. Taint Analysis

Taint analysis tracks **how user input flows through the application**, identifying code paths where input is used unsafely. If a webshell script receives a POST request and executes the content without validation (via `eval($_POST['cmd'])`), taint analysis will flag this as dangerous behavior.

This method is highly effective for detecting **command injection**, **remote file inclusion**, and **arbitrary code execution**, key attack vectors exploited by webshells.

c. System Call Tracing

System call analysis involves monitoring low-level OS interactions made by the executing script. Webshells that perform unauthorized operations are often invoked.

- File system calls: `open()`, `unlink()`, `chmod()`
- Network calls: `connect()`, `send()`, `recv()`
- Process control calls: `fork()`, `execve()`, `system()`

Tools such as **strace (Linux)** and **Sysinternals Suite (Windows)** help in capturing and logging these calls.

d. Network Traffic Analysis

Many webshells establish **outbound connections** to command-and-control (C2) servers or exfiltrate data via HTTP/S, FTP, or even DNS channels. Monitoring and analyzing traffic while interacting with the script helps detect.

- Data exfiltration attempts.
- Reverse shell connections.
- Covert communication patterns (using base64 in POST requests).

Integration with tools like **Wireshark**, **tcpdump**, or **Suricata** is common for this purpose.

e. Browser-Emulated Execution (for Front-End Shells)

Although webshells typically target server-side scripts, some variants may include front-end elements (e.g., malicious JavaScript UIs). Emulating browser-based behavior using tools like **Selenium** or **Puppeteer** helps capture dynamic rendering issues, iframe injections, or redirection behavior that static scanning might miss.

f. Honeypots and Sandboxes

Deploying **honeypots** or sandboxed environments mimicking vulnerable servers is another method of capturing dynamic behavior in real-world conditions. When attackers upload a webshell, the sandbox captures it.

- The file structure and signature of the webshell.
- Commands sent via HTTP.
- Subsequent attacker behavior (privilege escalation, lateral movement).

Projects like **Cuckoo Sandbox** or **FireEye's Dynamic Threat Intelligence** platform offer advanced features for behavior logging and visualization.

2.3.3 Tools for Dynamic Analysis

Tool	Description
Cuckoo Sandbox	An automated malware analysis tool that supports webshell execution
Process Monitor (ProcMon)	A Windows tool to log real-time system activity
strace / auditd	Linux tools to trace system calls and audit logs
Wireshark / tcpdump	Packet sniffers for analyzing live traffic generated by web shells

Xdebug (PHP)	PHP debugger for tracing and profiling script execution
Selenium / Puppeteer	Automation tools for simulating user interaction with web interfaces

2.3.4 Limitations of Dynamic Analysis

Despite these many strengths, dynamic analysis has important limitations that should not be overlooked.

Sandbox and traffic analysis require substantial hardware and processing resources. Simulating human interactions and analyzing resulting behavior can be time-consuming, particularly at scale.3 Detection Evasion Cybersecurity Executives' Guide to AI and ML 19 More advanced webshells could implement anti-debugging or anti-virtualization checks to escape detection. Coverage Gaps If a webshell has logic that only runs under certain conditions (a certain cookie or user-agent), the analysis will not catch any malicious behavior. Failure to Scale Easily. Unlike something like a static scanner that you can run across thousands of files in seconds, dynamic analysis tools are overall much slower and difficult to scale.

2.3.5 Hybrid Approaches

Because of the complementary nature of static and dynamic analysis, most state-of-the-art detection systems now leverage both strategies. A hybrid system could initially pre-screen potential threats using static techniques to eliminate the majority of unsafe content, followed by dynamic methods for deeper examination. Others, though, overlap.

Standardized static feature extraction for ML classification models.Dynamic behavior graphs for anomalous behavior detection. Runtime sandbox execution traces to supplement training data.

These hybrid models hope to have the cake and eat it too, getting the best of both static analysis's speed and broad coverage while offering more accurate, deeper analysis as in dynamic models.

○ Conclusion of Section 2.3

Dynamic analysis techniques take us a lot further, even helping to protect against webshells that would otherwise evade static detection through obfuscation and runtime behavior triggers. By seeing how scripts interact under user actions in real-time, humans or machines can more effectively spot mysterious malicious behavior. The cost of computation and lack of coverage expose the potential of machine learning driven static analysis models, such as attention-based deep learning structures, which offer a more scalable, informative approach. In the following sections, we walk through the

construction of the webshell dataset and how machine learning algorithms have been applied to solve the detection problem.

2.4 Dataset Collection and Pre-processing

The success of any machine learning model is inextricably tied to data. The quality, diversity, and relevance of the dataset the model was trained. When it comes to the narrower task of webshell detection, collecting an accurately-labeled, evenly-distributed dataset is particularly challenging because the evasive nature of harmful scripts prevents and raises ethical questions about their release. Webshells are very syntactically distinct with regards to their syntax, encoding method, structure, and length, and require a lot of preprocessing to even establish a meaningful learning.

In this work, we first paint a broad strokes picture of the pipeline we created for quickly curating a wide dataset of malicious (webshell) and normal (non-malicious) scripts, then discuss in-depth preprocessing techniques used to condition the data for classification based upon deep learning techniques.

○ 2.4.1 Sources of Data

A. Malicious Webshell Samples

Malicious scripts used in this research were collected from a combination of open-source intelligence (OSINT) resources and malware repositories. Key sources include.

- **VirusTotal**: Used for downloading real-world webshells submitted by global users.
- **GitHub**: Public repositories containing well-documented webshell tools (China Chopper, b374k, WSO, FilesMan).
- **VX Heaven and MalwareBazaar**: Archives of known malware binaries and scripts.
- **Security research blogs and honeypots**: Including incident response teams that publish forensic data and Indicators of Compromise (IoCs).
- **Capture-the-Flag (CTF) challenge dumps**: Simulated malicious PHP scripts crafted for cybersecurity competitions.

Collected samples were carefully reviewed to avoid duplicates and to ensure coverage across different **types, families, and obfuscation levels**.

B. Normal PHP Scripts

Normal samples were curated to closely represent realistic server-side codebases, ensuring the model is trained on diverse and practical examples. Data sources include:

- **Open-source CMS systems**: WordPress, Drupal, Joomla.
- **Frameworks**: Laravel, CodeIgniter, Symfony.

- **Community-contributed PHP projects**: Available via GitHub and GitLab repositories.
- **Educational repositories**: PHP tutorials and textbook code samples.
- **Enterprise-grade open APIs and admin panels**.

These Normal scripts include utility functions, database interactions, file handlers, and authentication logic to simulate real-world workloads.

○ 2.4.2 Dataset Composition

The final dataset was composed of.

- **5,000+ malicious samples** from over 20 different webshell variants.
- **5,000+ normal scripts** representing over 200 real-world PHP applications.

A 50:50 class balance was maintained.

- Prevent bias in training.
- Enable reliable evaluation across precision and recall.
- Facilitate easy adaptation for both binary and multi-class classification in future research.

○ 2.4.3 Data Cleaning and Normalization

Before tokenization and vectorization, several data cleaning steps were applied:

- **Encoding normalization**: Converted all files to UTF-8 to handle mixed language input.
- **Comment removal**: Removed PHP, HTML, and JavaScript comments to reduce noise.
- **Whitespace and newline stripping**: Compressed multi-line code into single-line statements for consistent sequence lengths.
- **Removal of binary files or image content**: Detected and discarded base64-encoded image files masquerading as PHP.

Cleaning ensured that only meaningful textual code was passed to the tokenizer, improving learning efficiency.

○ 2.4.4 Tokenization

Tokenization involves breaking down code into meaningful units or "tokens." Unlike natural language, programming code requires language-aware parsing to handle syntax correctly.

- **Custom PHP tokenizer**: Developed using the `php-parser` Python module.
- **Tokens include**:
    - Keywords: `if`, `while`, `echo`, `eval`

- Variables: `$username`, `$cmd`
- Operators: `=`, `==`, `+`, `.`
- Literals: `'admin'`, `"password"`
- Function calls: `system()`, `base64_decode()`

This tokenization approach preserves syntactic structure and enables semantic analysis of command patterns.

2.4.5 Vocabulary Building and Encoding

- **Vocabulary size**: After cleaning and tokenization, a vocabulary of ~10,000 unique tokens was constructed.
- **Token-to-index mapping**: Each token was assigned a unique index (starting from 1; 0 reserved for padding).
- **Out-of-vocabulary handling**: Rare tokens were mapped to a generic `<UNK>` symbol.
- **Embedding strategy**: Pre-trained Word2Vec embeddings were used to convert each token into a 100-dimensional vector.

2.4.6 Sequence Construction and Padding

Deep learning models require input sequences of uniform length.

- **Maximum sequence length**: Set to 1,000 tokens, based on analysis of script lengths (95th percentile).
- **Padding**: Scripts shorter than the max length were padded with zeros (`<PAD>` tokens).
- **Truncation**: Longer scripts were truncated to maintain consistent input dimensions.

This allowed batch processing and optimized GPU utilization during model training.

2.4.7 Labeling and Splitting

Each sample was assigned a binary label.

- `1` for malicious (webshell).
- `0` for Normal (normal PHP script).

The dataset was split into.

- **Training set (80%)** – used to train the model.
- **Validation set (10%)** – used to fine-tune model hyperparameters.
- **Test set (10%)** – held out for final performance evaluation.

A stratified shuffle split ensured both classes were evenly distributed in each subset.

○ 2.4.8 Memory Mapping and Storage Optimization

To efficiently handle large datasets.

- **Memory-mapped files** (`.memmap`) were used to store tokenized sequences and labels on disk.
- Enabled lazy loading during training, minimizing RAM usage.
- Facilitated scalability to millions of lines of code for future experiments.

○ 2.4.9 Dataset Augmentation (Optional Phase)

Experimental augmentation was performed to increase generalization. **Variable renaming**: Changed non-sensitive variable names randomly (`$a` to `$temp1`)

- **Shuffling non-executable lines**: Reordered `echo` or `comment` lines to simulate code diversity.
- **Obfuscation simulation**: Applied basic base64 encoding and decoding.

These steps mimicked real-world code variations without altering functionality.

○ Conclusion of Section 2.4

Building a high-quality dataset is arguably the most important part of developing any machine learning system. The wide, balanced, and context-rich curated dataset used in this research helped this attention-based deep learning model to learn meaningful patterns that differentiate malicious webshells from normal scripts. Through each stage of the preprocessing pipeline, from tokenization through to padding and embedding, the input to the model is optimized not just for efficient learning, but for practical application to the real world. The following chapter will focus on the architecture of the webshell web detection model we proposed and its attention-based mechanism to detect the webshells.

2.5 Machine Learning in Malware Detection

With malware growing in complexity, such as obfuscated and zero-day webshells, rule-based, detective, and traditional detection approaches are not always catching everything. In response, machine learning (ML) has proven as a scalable, adaptable solution with the potential to identify new threats using pattern recognition and generalization.

○ 2.5.1 Why ML for Malware Detection

ML overcomes major shortfalls of past traditional approaches, like a lack of adaptability, excessive false positives, and manual updating of signatures. ML models automate pattern discovery, are better equipped to manage the scale and constant evolution of threats.

○ 2.5.2 ML Approaches

ML paradigms in malware detection include:

- **Supervised Learning** (SVM, Random Forests, DNNs): Effective with labeled datasets.
- **Unsupervised Learning** (Clustering, Autoencoders): Useful for anomaly and zero-day detection.
- **Semi-Supervised Learning**: Leverages limited labeled data with large unlabeled corpora.

○ 2.5.3 Key Features for Detection

Common features used include:

- Byte n-grams
- Opcode frequency
- API call patterns (`eval`, `system`)
- Entropy measures
- Token sequences

These capture both structural and behavioral characteristics of scripts.

○ 2.5.4 NLP and Deep Learning

Due to syntactic similarities between code and language, NLP models such as:

- **RNNs / LSTMs / BiLSTMs**
- **CNNs**
- **Transformers**
  It has been applied to code-based malware detection. Attention mechanisms improve interpretability by directing attention to spy tokens.

○ 2.5.5 ML Frameworks

Popular ML-integrated tools include:

- **EMBER** (Random Forest)
- **MalConv** (CNN on raw bytes)
- **VirusTotal**, **Cuckoo Sandbox + ML**, and **Microsoft ATP** (hybrid methods)

○ 2.5.6 Challenges

Key challenges include:

- Adversarial evasion
- Imbalanced datasets
- High labeling cost

- Lack of model interpretability
- Concept drift requires continual retraining

○ 2.5.7 Traditional Classifiers

Legacy models such as:

- **Decision Trees**
- **SVM**
- **KNN**
- **Naive Bayes**
- **Logistic Regression**

offer interpretability and speed, but often fail with heavily obfuscated or evolving threats.

○ 2.5.8 Deep Learning Models

Deep models automatically extract features and handle context:

- **CNNs** detect local patterns
- **RNNs/LSTMs/BiLSTMs** handle sequences
- **Attention & Transformers** focus on relevant tokens and capture long-range dependencies
- **Autoencoders/VAEs** support anomaly detection

○ 2.5.9 Traditional vs. Deep Learning (Summary)

Aspect	Traditional ML	Deep Learning
Feature Engineering	Manual	Automated
Interpretability	High	Moderate
Obfuscation Handling	Low–Moderate	High
Data Needs	Low	High
Detection Accuracy	Moderate	High

- **Conclusion**

Using machine learning, particularly deep learning with attention, malware webshell detection has seen incredible success in improving precision, adaptability, and explanation. To alleviate this problem, this paper proposes a novel webshell detection framework based on BiLSTM with word-level attention.

2.6 Attention Mechanisms & Transformer Models

2.6.1 Introduction

Rapidly growing web-based threats such as obfuscated and polymorphic webshells have outstripped the ability of classic malware detection methods. Even state-of-the-art deep learning models like CNNs and RNNs, despite their empirical successes, fail to fully capture long-term dependencies and variable-length relationships present in the source code. To combat these problems, attention mechanisms and their most recent form, transformer models, have changed the way that sequence-order data, such as programming code or tokenized dialog scripts, are modeled and interpreted.

Attention mechanisms allow a model to dynamically attend to the most relevant parts of an input sequence, improving accuracy as well as interpretability. Transformers take this principle to another level by removing recurrence completely and processing sequences in parallel through self-attention layers.

This section dives into the theory and practice behind attention mechanisms and transformer models in webshell and malware detection, with complete technical rigor and mathematical fundamentals.

2.6.2 Attention Mechanisms: Theory and Formula

- **A. Motivation**

In traditional sequence models such as RNNs or LSTMs, the final hidden state is used as a summary of the full input sequence. This technique doesn't cut it for long sequences or when some special tokens (eval, base64_decode, system) are far away from the sequence end.

Attention solves this problem by calculating a weighted average of the input hidden states, where the weights depend on how relevant each hidden state is for the current task (e.g., classification).

- **B. Mathematical Formulation of Attention**

Let us define the input sequence as a set of vectors:

$$X = [x_1, x_2, \ldots, x_T \ where \ x_i \in R^d$$

The model generates hidden states:

$$H = [h_1, h_2, \ldots, h_T \; where \; h_i \in R^k$$

For each hidden state h_i, the attention weight α_i is computed as:

$$e_i = f(h_i) = tanh(W_h h_i + b_h)$$

$$\alpha_i = \frac{exp(e_i)}{\sum_{j=1}^{T} exp(e_j)} \; (softmax\ over\ attention\ scores)$$

The **context vector** c, which summarizes the important parts of the sequence, is then computed as:

$$c = \sum_{i=1}^{T} \alpha_i h_i$$

This context vector is passed to the final layer (e.g., a sigmoid-activated neuron) for classification.

Explanation:

- W_h, b_h They are learnable parameters.
- α_i Represents the attention weight for the token i.
- High α_i means the model considers h_i it highly relevant for the task.

2.6.3 Self-Attention Mechanism

With self-attention, every token in the sequence attends to every other token, giving the model the capacity to create these deep contextual embeddings. The multi-head self-attention mechanism, the very first block in the transformer architecture.

Given:

- Input matrix $X \in R^{T \times d} (T = tokens, \; d = embedding\ size)$

We compute:

- **Query (Q)**, **Key (K)**, and **Value (V)** matrices.

$$Q = XW^Q, \quad K = XW^K, \quad V = XW^V$$

Where.

- $W^Q, W^K, W^V \in R^{d \times d_k}$ They are learnable weights.

Then the attention scores are:

$$Attention(Q, K, V) = softmax\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

Explanation:

- QK^TSimilarity between tokens
- $\sqrt{d_k}$scaling factor for stable gradients
- The final output is a context-aware representation of each token.

○ 2.6.4 Multi-Head Attention

To allow the model to **focus on different types of relationships** simultaneously, transformers use **multi-head attention**:

$$MultiHead(Q, K, V) = Concat(head_1, head_2 ...., head_h)W^O$$

Where:

- $head_i = Attention\left(QW_i^Q, KW_i^K, VW_i^V\right)$
- W^O: output projection matrix

Each head learns to capture different patterns (e.g., syntactic structure, semantic meaning, suspicious tokens).

2.6.5 Transformer Architecture for Code and Malware

The **Transformer** is a deep learning architecture that relies entirely on attention mechanisms. It consists of:

- **Embedding Layer**: Converts tokens to dense vectors.
- **Positional Encoding**: Injects information about the token's position (since transformers have no recurrence).
- **Stacked Encoder Layers**: Each includes:
    - Multi-head self-attention

    - Feedforward neural network
    - Layer normalization + residual connections
- **Decoder Layers** (not used in classification tasks)
- **Output Layer**: Dense + Softmax/Sigmoid for classification

- **Positional Encoding Formula**

Since transformers process input in parallel, **positional encoding** helps track the token order:

$$PE_{(pos,\,2i)} = sin\left(\frac{pos}{10000^{\frac{2i}{d}}}\right)$$

$$PE_{(pos,\,2i+1)} = cos\left(\frac{pos}{10000^{\frac{2i}{d}}}\right)$$

Where:

- posposition of the token
- i: dimension index
- d: embedding dimension

2.6.6 Benefits of Attention and Transformers in Webshell Detection

Feature	Advantage
Focus on critical tokens	Helps identify dangerous function calls (e.g., `eval`, `assert`)
Capture long-range dependencies	Useful when malicious logic is separated by non-malicious code
Interpretability	Visualizing attention weights helps security analysts understand model decisions
Scalability	Transformers can process long scripts faster and more effectively than RNNs

Example:
In a PHP script with `eval(base64_decode($_POST["x"]))`, a BiLSTM may struggle to remember the long-range dependency between `eval` and the malicious input. A self-attention model can instantly relate these tokens, regardless of their distance.

2.6.7 Transformers Applied to Code Analysis

- **A. CodeBERT / GraphCodeBERT**
  - Pretrained on millions of code snippets from GitHub.
  - Understands both natural language (comments, variable names) and code syntax.
  - Fine-tuning for malware classification yields strong results.
- **B. GPT-style Models for Code**
  - Useful for anomaly detection, code completion, and auto-correction.
  - GPT-2 and GPT-3 variants trained on code can identify abnormal patterns or obfuscation artifacts.
- **C. Transformer + LSTM Hybrids**
  - LSTM captures sequence information
  - Attention layer helps with interpretability
  - The transformer layer captures global relationships
  - Used in this research to create a **context-aware, interpretable model** for webshell detection.

2.6.8 Limitations and Considerations

Challenge	Description
Training Cost	High computational resources required (GPU/TPU)
Data Requirements	Requires large, high-quality datasets to pretrain
Overfitting	Especially on small, imbalanced datasets
Interpretability	Although attention adds transparency, full model decisions remain complex

Conclusion of Section 2.6

Attention mechanisms and transformers are the state of the art in sequence modeling and are currently changing the landscape of malware and webshell detection. By permitting models to selectively attend to discriminative parts and discern global contextual dependencies, these architectures significantly enhance classification metrics, obfuscation robustness, and interpretability.

To counteract these attack vectors, the proposed work from this dissertation takes advantage of these mechanisms, combining a BiLSTM with a word-level attention layer, to create an intelligent, adaptive, and explainable model that can efficiently be deployed to intelligently identify potentially malicious webshell scripts found in the wild.

2.7 Comparative Review of Existing Systems

A vital part of any research endeavour is taking a hard look at the strengths and weaknesses of current systems within the same problem space. That's not the case with webshell detection, as millions of systems and frameworks, from signature-based static detection tools to dynamic learning models fueled by AI, have been created to address this long-lasting threat. Yet notwithstanding these advances, their accuracy, scalability, adaptability, and interpretability vary widely as these systems have been diffused far and fast.

This section provides a comparative analysis of state-of-the-art webshell detection systems and methods, categorised into four main categories:

Signature-Based Solutions, Heuristic and rule-based systems, Common Static & Dynamic Malware Analysis Frameworks, Machine learning and deep learning models.

Each system is evaluated based on the following criteria:

- **Detection Accuracy**
- **Zero-Day Generalization**
- **Resistance to Obfuscation**
- **Interpretability**
- **Execution Overhead**
- **Deployment Suitability**

○ 2.7.1 Signature-Based Detection Systems

In short, signature-based systems detect known threats by using pre-defined patterns known as signatures. While these systems are very quick and efficient at catching malware or webshells you've seen before, they choke on unforeseen polymorphic or zero-day webshells.

Tool/Method	Strengths	Weaknesses
ClamAV	Lightweight, widely deployed	Cannot detect obfuscated or novel webshells
ModSecurity (WAF)	Real-time traffic filtering using signatures	Frequent false positives, poor adaptability
YARA Rules	Highly customizable for known shell variants	Requires continuous manual updates

Summary:
Quick and effective, signature-based tools are always reactive and can be evaded with minor script changes or encoding.

○ 2.7.2 Heuristic and Rule-Based Detection

Heuristic-based systems analyze scripts for suspicious constructs or combinations of behaviors. These rules are manually defined by security experts.

Tool/Method	Strengths	Weaknesses
RIPS Scanner	PHP-specific vulnerability scanner	Cannot detect runtime behavior
PHP Malware Finder (PMF)	Detects dangerous functions (e.g., `eval`, `assert`)	Lacks context awareness; high false alarms
Custom Regex Filters	Easy to implement and integrate	Brittle and non-adaptive; vulnerable to evasion techniques

Summary:
Heuristics can provide a good balance between accuracy and performance for known patterns, but they suffer from a lack of generalization and context awareness.

○ 2.7.3 Static and Dynamic Malware Analysis Frameworks

These systems use static or runtime behavior analysis to identify malicious scripts. Dynamic analysis can uncover stealthy behavior by monitoring execution traces.

Tool/Framework	Static / Dynamic	Strengths	Weaknesses

Cuckoo Sandbox	Dynamic	High behavioral insight	High setup cost and slow analysis
Maltrail IDS	Static + network heuristics	Detects outbound callbacks	Not tailored to webshell detection
JSDetox	Hybrid (for JS shells)	Emulates JavaScript behavior	Cannot handle PHP webshells
php-malware-finder	Static	Lightweight and quick	Obfuscation-sensitive

Summary:
Dynamic systems provide valuable behavioral context but are computationally expensive and may not scale easily in real-time scenarios.

○ 2.7.4 ML and Deep Learning-Based Systems

Modern systems use machine learning and deep learning to learn complex patterns in scripts, making them resilient against obfuscation and unseen threats.

Model	Description	Pros	Cons
SVM with N-gram features	Classical ML model with fixed-size n-gram inputs	Simple and interpretable	Limited to shallow patterns
Random Forest with Opcode Stats	Uses opcode frequency for classification	High precision	Requires expert-crafted features
TextCNN	Learns local token patterns in code	High accuracy	Poor at modeling long-range dependencies
BiLSTM	Captures forward/backward context in code	Good at handling sequences	Memory-intensive
BiLSTM + Attention	Focuses on key malicious functions/tokens	High accuracy and interpretability	Needs large labeled datasets

Transformer (CodeBERT, GPT-C)	Fully context-aware, parallel processing	State-of-the-art generalization	High training cost, black-box nature

Summary:
Deep learning-based models, especially those using attention or transformer layers, outperform classical systems in adaptability, detection accuracy, and robustness to obfuscation. However, they require more data, processing power, and careful training.

○ 2.7.5 Comparative Summary Table

Criteria	Signature-Based	Rule-Based	Static/Dynamic Analysis	ML Models	DL Models (Proposed)
Zero-day Detection	✗	✗	✓ (dynamic)	✓	✓✓
Handles Obfuscation	✗	✗	✓	✓	✓✓
Accuracy	Moderate	Moderate	High	High	Very High
Interpretability	High	Moderate	Low	Moderate	High (via attention)
Speed / Performance	Very High	High	Low	Moderate	Moderate
Scalability	High	Moderate	Low	Moderate	High (with GPUs)
Automation Required	Low	High	Medium	High	High

○ 2.7.6 Justification for the Proposed Work

To determine smart technologies' capacity and greatest potential to minimize negative impacts, we conducted this comparative analysis of smart technologies' efficacy and return on investment to conclude where the intelligent, interpretable, and scalable detection system is most visibly needed. These traditional approaches do not adequately address evasive webshells and move to new attacks. Dynamic systems provide penetrating analysis but are not scalable to larger project arrays.

The proposed model in this study, by employing BiLSTM with a word attention component, can overcome these challenges by:

Learning intrinsic structures from the ordered sequences of tokens, rather than relying on handcrafted featuresIdentifying important parts of code with attention weightsProviding interpretable detection, helping security analysts to confirm model choicesStill staying computationally efficient in comparison to transformer-only models

- **Conclusion of Section 2.7**

The sea of webshell detection systems is vast, from naive, signature-based systems to complex attention-based deep learning architectures. In conclusion to this comparative review, hybrid neural models that leverage sequential learning and attention mechanisms emerged as the superior choice, creating an optimal balance between accuracy, interpretability, and scalability. This observation is what directly inspires the architecture and approach we propose in the following chapter.

2.8 Chapter Summary

This chapter closed with a discussion of the relevant literature that has focused on webshell detection methodologies, setting the stage for the ever-changing progression of attacker methodologies, and equally as imperative, defender methodologies. The degree of obfuscation and evasion that webshells employ.

To illustrate the strengths and weaknesses of static and dynamic analysis techniques, static and dynamic analysis techniques were compared. Static techniques are powerful, but static techniques can be trivially broken through obfuscation. Dynamic techniques, in return, can be extremely computationally costly, but in exchange, provide a much greater understanding of behavior.

Moving forward with the chapter, they introduced different types of machine learning based approaches and provided an overview of classical classifiers (such as SVM, Random Forests), up to date with classical models, to state-of-the-art deep learning models such as CNNs, LSTMs, BiLSTMs, and Transformers. Special emphasis was placed on attention mechanisms, which increase model explainability and help focus model attention on code fragments of increased interest for malicious intent.

A subsequent survey of the related literature found hybrid deep learning models, particularly those incorporating attention, to be more resilient than conventional models in both the adaptability and accuracy categories.

Overall, this chapter established the theoretical basis for our BiLSTM-based detection framework proposal, detailed in the following chapter.

Chapter 3
DESIGN OF A DEEP-LEARNING WEB SHELL DETECTION FRAMEWORK

This chapter gives details on how the proposed system “Deep Learning Based Webshell Detection and Prevention Using Word-Level Attention Mechanisms. This chapter describes the system architecture in depth, defines the data pipeline, describes the internal workings of the neural model, and outlines the methodology that has been followed throughout this research.

To give the audience a deeper understanding of webshell detection using the introduced framework, the purpose of this chapter is to provide a systematic explanation on how the integration of Bi-directional Long Short-Term Memory (BiLSTM) networks and word-level attention mechanism can be utilized in the detection of webshells. To that end, the chapter includes architectural diagrams alongside algorithmic underpinnings that, as a whole, contribute explainability, accuracy, and robustness to the model.

Our proposed system is to accept web-facing scripts (mostly PHP), tokenize and vectorize them, then run them through an attention-enabled sequential model that learns to predict when nefarious behavior is afoot. Our end goal, though, is not just to find these webshells with high precision, but rather with an interpretable output that explains what features of the code led to the classification decision.

3.1 System Architecture Overview

Our proposed framework is an end-to-end deep learning-based system specifically designed for the detection of malicious webshells from PHP source code. Our architecture takes advantage of a Bidirectional Long Short-Term Memory (BiLSTM) network augmented with a word-level attention mechanism, allowing it to learn local and global contextual features in source code sequences while focusing on important tokens that drive classification decisions.

This section provides a high-level view of the system architecture, its components, and the entire data processing pipeline pictured in Figure 3.1.

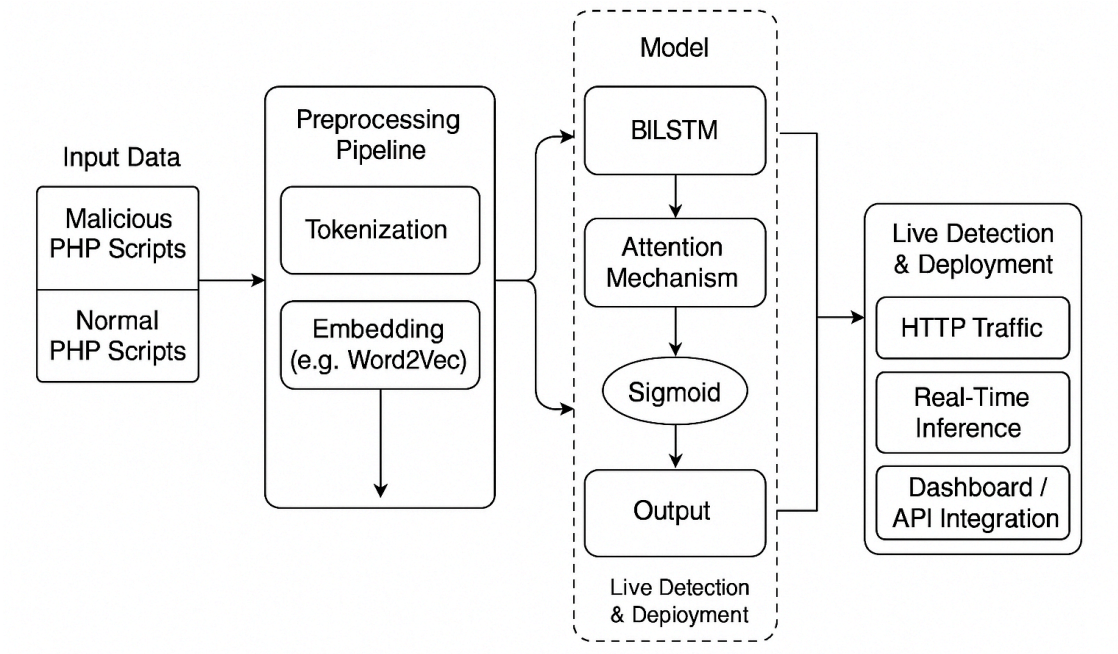

Figure 3.1: Deep learning-based webshell detection framework with tokenization, embedding BiLSTM-Attention, and binary classification output

○ 3.1.1 Architectural Components

The proposed system is organized into several modular stages:

○ 1. Data Acquisition and Corpus Scanning

- **Malicious and Normal samples** are collected from curated datasets, including webshell repositories (e.g., China Chopper, WSO, b374k) and open-source PHP projects (e.g., WordPress, Laravel).
- The corpus is **scanned** to compute:
  - Vocabulary size
  - Maximum sequence length
  - Total number of samples

○ 2. Tokenization and Vocabulary Encoding

- A **custom tokenizer** processes each line of code by splitting tokens using whitespace and language syntax rules.
- Tokens are mapped to unique integer IDs using a `tokenizer.pkl` file, enabling the transformation of variable-length scripts into **fixed-length sequences**.

○ 3. Embedding Layer

- Tokens are transformed into **dense vector embeddings** using a **Word2Vec model** trained on the combined dataset.

- This results in an embedding matrix $E \in R^{V \times d}$, where V is the vocabulary size and $d = 100$ is the embedding dimension.
- Embeddings allow semantically similar tokens (e.g., `eval` and `exec`) to lie close together in vector space.

○ **4. BiLSTM Encoder**

- The token embeddings are passed through a **Bidirectional LSTM** layer:

$$\overrightarrow{h_t}, \overleftarrow{h_t} = LSTM(x_t)$$

$$h_t = \left[\overrightarrow{h_t}; \overleftarrow{h_t}\right]$$

- This captures both **past and future context**, critical for understanding nested and reordered patterns often used in obfuscated webshells.

○ **5. Word-Level Attention Layer**

- The attention mechanism computes importance weights α_t for each hidden state h_t:

$$e_t = tanh(W_a h_t + b_a)$$

$$\alpha_t = \frac{exp(e_t)}{\sum_{k=1}^{T} exp(e_k)}$$

$$c = \sum_{t=1}^{T} \alpha_t h_t$$

- This context vector c selectively aggregates information from the most relevant parts of the script.

○ **6. Dropout and Fully Connected Layer**

- A **dropout layer** (rate = 0.3) is applied to the context vector to prevent overfitting.

- The processed vector is passed to a **dense output layer**:

$$y = \sigma\left(W_c c + b_c\right)$$

Where σ is the sigmoid activation function producing a probability score between 0 (Normal) and 1 (malicious)?

7. Output Interpretation and Visualization

- The output score is thresholded (e.g., 0.5) to produce the final prediction.
- Attention weights α_t are visualized to **highlight the most suspicious tokens**, enabling explainable AI (XAI) capabilities.

8. Model Training and Evaluation Loop

- The system is trained using **binary cross-entropy loss**:

$$L\left(y, \hat{y}\right) = -\left[y\, log\left(\hat{y}\right) + (1 - y)\, log\left(1 - \hat{y}\right)\right]$$

- **Adam optimizer** with a learning rate of 0.001 is used for gradient descent.
- Evaluation is performed on metrics such as:
  - Accuracy
  - Precision
  - Recall
  - F1-score
  - ROC-AUC

3.1.2 System Flow Diagram (Conceptual)

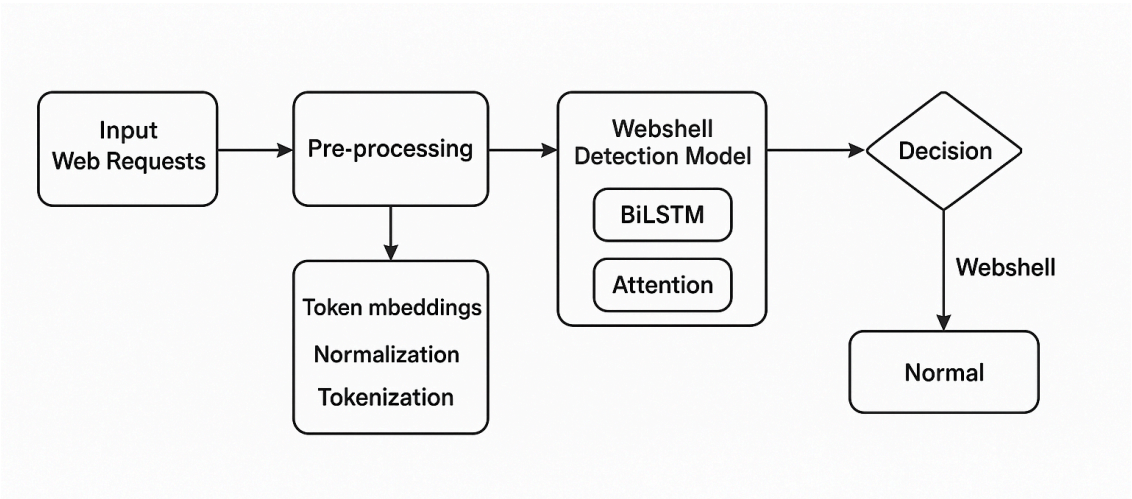

3.1.3 Design Justifications

- **BiLSTM** was chosen for its ability to understand code dependencies in both directions (e.g., variable assignments before/after use).
- **Attention mechanism** increases both model interpretability and performance by isolating critical code features.
- **Word2Vec embeddings** capture semantic similarity between tokens beyond simple frequency.

○ 3.1.4 Deployment Considerations

- The model can be plugged into any web server pipeline, executing in a daemon/microservice style to scan all uploaded/modified PHP files in real-time.
- **Batch scanning mode** is also supported for forensic analysis and security auditing.

○ Conclusion of 3.1

The system architecture strikes a careful balance between accuracy, scalability, and explainability, three key pillars in a contemporary state-of-the-art malware detection system. In the following sections, we explore the model architecture components, training hyperparameters, and implementation details in much more technical detail.

3.2 Data Ingestion and Tokenization

Tokenization and data ingestion play a significant role in the performance and generalizability of deep learning models used for code-based malware detection. This section describes the throughput of how raw PHP scripts, both malicious (webshells) and normal (normal server-side code), are ingested, cleaned, tokenized, vectorized, and structured into fixed-length input representations to be downstream trained in a deep learning framework.

From the initial data collection through to final training and validation, the whole pipeline is built to protect semantic integrity, scalability, and efficiency to make sure that the model is trained on inputs that keep both syntax and malicious intent traits.

○ 3.2.1 Raw Data Acquisition

The first step is to compile real-world code examples from reliable sources to create two evenly distributed datasets of:

A. Malicious Dataset (Webshells)

- Sourced from:
  - GitHub repositories (e.g., WSO, b374k, China Chopper)
  - Security research honeypots and malware archives (e.g., VX Heaven, VirusShare)
  - CTF (Capture the Flag) challenge dump files
- Characteristics:

    - Highly obfuscated
    - Frequently uses `eval()`, `assert()`, `preg_replace()` with `/e`, etc.
    - May contain encoded or encrypted payloads (e.g., base64, gzinflate, ROT13)

B. Normal Dataset (Normal PHP Scripts)

- Sourced from:
    - Open-source projects (WordPress, Laravel, Symfony)
    - Code tutorials and textbook repositories
    - Real-world production code (scrubbed for PII)
- Characteristics:
    - Legitimate control logic
    - Normal use of variables, file handlers, and session management
    - No malicious function chaining or decoding

○ 3.2.2 Directory Traversal and File Ingestion

To collect and load data:

- A **recursive file scanner** navigates through `/malicious` and `/normal` directories.
- Files are opened with `UTF-8` encoding and in `ignore` error mode to skip corrupted or non-text files.
- Each valid file is read **line by line** into memory.
- Metadata, including filename, class label (0 or 1), and file length, is logged.

Sanity checks include:

- Ignoring empty or binary files
- Removing scripts with <10 tokens
- Flagging duplicates by checksum (SHA256)

This process is handled via `scan_corpus()` and `LineIterator()` in your training code.

○ 3.2.3 Cleaning and Normalization

Raw source code is noisy and inconsistent. Cleaning ensures the dataset is uniform and model-friendly.

Cleaning Operation	Details
Lowercasing	Makes tokens case-insensitive (`Eval` = `eval`)

Comment removal	Filters `//`, `#`, and `/* */` comment blocks
Whitespace normalization	Replaces tabs/multiple spaces with single spaces
HTML/PHP tag removal	Strips `<?php`, `?>`, `<script>`, etc., if needed
Non-ASCII filtering	Optional removal of high-bit characters
Encoding detection	All files forced to UTF-8

These steps standardize input and prevent the model from overfitting to irrelevant patterns (like stylistic spacing or developer comments).

○ 3.2.4 Tokenization of Source Code

To better represent the syntactic structure of source code, a special custom tokenization strategy is applied, in contrast with general-purpose NLP tokenizers.

Token Types:

- **Keywords** (`if`, `while`, `return`, `echo`)
- **Variables** (`$username`, `$_POST`, `$cmd`)
- **Operators** (`+`, `-`, `=`, `==`, `.=`, `.`)
- **Delimiters** (`;`, `{}`, `()`, `[]`)
- **Function Calls** (`eval()`, `include()`, `system()`)
- **Literals** (strings, numbers, booleans)
- **Encoded Content** (e.g., base64 string marked as `<ENCODED>`)

Tools Used:

- `str.split()` + regex for quick tokenization
- Future extensions may include `php-parser` or `tree-sitter` for AST-based token splitting

Each script is tokenized into a list of tokens, such as:

```
['eval', '(', 'base64_decode', '(', '$_POST', '[',
"'cmd'", ']', ')', ')', ';']
```

○ 3.2.5 Vocabulary Construction

All unique tokens across both datasets are compiled into a **vocabulary** dictionary.

Process:

- Sorted alphabetically or by frequency
- Indexed starting from 2 (`0` = `<PAD>`, `1` = `<UNK>`)
- Rare tokens below the frequency threshold can be removed or replaced

This vocabulary is used to **map each token to an integer**, a crucial step for feeding into embedding layers:

```
tokenizer = { 'eval': 5, 'base64_decode': 12, '(': 3, ...
}
```

Vocabulary Size: typically ~10,000–12,000 tokens

○ 3.2.6 Sequence Encoding and Length Normalization

Each tokenized line is then converted into an integer sequence using the tokenizer mapping. Because neural networks need fixed-length input:

- A maximum sequence length L_{max} is set (e.g., 1000 tokens).
- Shorter sequences are **padded** with `<PAD>` (index 0).
- Longer sequences are **truncated** from the end.

This results in a uniform 2D input array:

$$Input\ Shape: \left(N, L_{max}\right) where\ N = number\ of\ scripts$$

○ 3.2.7 Label Assignment

Each sequence is assigned a **binary label**:

- `1`: Malicious (webshell)
- `0`: Normal

Labels are stored in a 1D array:

```
labels = [1, 1, 0, 0, 1, 0, ...]
```

○ 3.2.8 Efficient Disk Storage Using Memory Mapping

To manage large datasets efficiently:

- `sequences.dat` stores all encoded sequences using NumPy `memmap`
- `labels.dat` stores class labels

This enables:

- Faster disk I/O during training
- Training on datasets larger than RAM

- Parallel loading in `DataLoader` with minimal memory overhead

Controlled via the `DiskDataset` class and `preprocess_to_memmap()` in code.

○ 3.2.9 Sample Example

Raw Input:

```
<?php eval(base64_decode($_POST['cmd'])); ?>
```

Tokenized:

```
['eval', '(', 'base64_decode', '(', '$_POST', '[',
"'cmd'", ']', ')', ')', ';']
```

Token IDs (Example):

```
[5, 3, 12, 3, 62, 6, 15, 7, 4, 4, 2]
```

Final Padded Sequence (max_len = 15):

```
[5, 3, 12, 3, 62, 6, 15, 7, 4, 4, 2, 0, 0, 0, 0]
```

○ 3.2.10 Pipeline Summary

Stage	Function
Ingestion	Read source files and extract content
Cleaning	Normalize, sanitize, and remove noise
Tokenization	Parse into programming-relevant tokens
Vocabulary Mapping	Convert tokens into integer IDs
Sequence Normalization	Pad or truncate to the fixed length
Labeling	Assign binary classification targets
Memmap Storage	Save disk-backed arrays for training

○ Conclusion of Section 3.2

This section sets the groundwork for an efficient and scalable preprocessing pipeline to turn raw PHP scripts into tokenized, padded, and numerically encoded inputs. This custom tokenization keeps semantic information and discards noise. The technique of

memory mapping is a key enabler for performing model training at large scales while maintaining performance.

The following section details how these encoded sequences are passed into a BiLSTM and attention-based neural network, which comprises the heart of the detection engine.

3.3 Embedding Construction (Word2Vec, CodeBERT, etc.)

One of the most important pieces in a neural network pipeline is the embedding layer. Embeddings act as a key interface between the discrete, symbolic tokens (i.e., words, variables, operators) that humans can manipulate directly and the continuous, differentiable vector representations that deep learning models require.

Specifically for the webshell detection application, the embeddings need to represent not just syntactic regularities (e.g., function definitions, symbol identifiers) but semantics (e.g., closeness of commands like eval, exec, system, known to be malicious). This section outlines embedding construction methods employed in this study, with an emphasis on:

- **Word2Vec** (custom-trained)
- **CodeBERT** (pretrained Transformer)
- **Comparison and integration into the architecture**

○ 3.3.1 Objectives of Embedding Construction

The core goals for embedding construction in this work are:

1. To **map discrete tokens** (from the tokenizer) into dense vector representations.
2. To ensure that embeddings capture both **semantic similarity** and **functional similarity** among tokens.
3. To provide a **context-aware initialization** for the BiLSTM-attention model.
4. To support **transfer learning** via pretrained models like CodeBERT.

○ 3.3.2 Word2Vec Embeddings

a) Overview

Word2Vec is a shallow, unsupervised learning algorithm popularized by Mikolov et al. For instance, word2vec Mikolov et al. For our applications in this study, a Skip-gram variant is employed, which seeks to predict surrounding tokens given a target token.

b) Training Details

- **Corpus**: Combined cleaned scripts from both `malicious/` and `normal/` directories.
- **Vocabulary**: 10,000+ unique tokens

- **Model**: `gensim.models.Word2Vec`
- **Configuration**:
  - `vector_size=100`
  - `window=5`
  - `min_count=1`
  - `sg=1` (Skip-gram)
  - `epochs=10`

  Code reference: `Word2Vec(sentences=LineIterator(...), vector_size=100)`

c) Output

- **Embedding Matrix** $E \in R^{|V| \times d}$: Each row E_i is a 100-dimensional vector for a token t_i.
- This matrix is used to initialize the embedding layer in the model:

$$embedding = nn.embedding.from_pretrained(E, freeze = False)$$

d) Advantages

- Captures **domain-specific semantics** of PHP tokens
- Can be updated during training (`freeze=False`)
- Fast to train, highly interpretable

○ 3.3.3 CodeBERT Embeddings

a) Overview

CodeBERT, a Microsoft Research creation, is a general-purpose, Transformer-based model that was pretrained on hundreds of millions of source code files written in different programming languages. It further facilitates other downstream tasks, including code summarization, code classification, and code clone detection.

b) Tokenizer and Model

- **Model**: `microsoft/codebert-base`
- **Tokenizer**: Byte-pair encoding (BPE) using HuggingFace's `transformers` library
- **Input format**: Script → Tokenized string → Input IDs + Attention masks

c) Embedding Extraction

To use CodeBERT for feature extraction:

Tokenize input code:

```
tokens = tokenizer(script, return_tensors='pt',
truncation=True, max_length=512)
```

Pass through model to get the last hidden states:

```
outputs = model(**tokens) embeddings =
outputs.last_hidden_state
```

Average pool or use [CLS] embedding as the sequence representation:

$$V_{code} = \frac{1}{T}\sum_{i=1}^{T} h_i$$

d) Advantages

- Encodes **global context** and **deep syntax-semantic dependencies**
- Learns **language-agnostic features** useful for generalization
- Supports **transfer learning**, reducing the need for huge local datasets

○ 3.3.4 Comparison: Word2Vec vs. CodeBERT

Criterion	Word2Vec	CodeBERT
Training Data	Custom local corpus	Massive public code corpus
Context Awareness	Local window-based (shallow)	Global self-attention (deep)
Domain Specificity	High (trained on PHP webshells)	Moderate (multi-language)
Speed	Fast training	Computationally intensive
Model Size	Lightweight (~MBs)	Large (multi-GBs)
Fine-Tuning Required	Optional	Recommended
Embedding Use Case	Input to LSTM/CNN	Feature extraction or fine-tuning

○ 3.3.5 Integration into the Model Pipeline

Word2Vec Integration:

The embedding matrix is passed to an `nn.Embedding` layer in PyTorch:

```
self.embedding =
nn.Embedding.from_pretrained(torch.tensor(embedding_matrix
), freeze=False)
```

CodeBERT Integration Options:

- **As Feature Extractor**:
  - Use CLS token output or pooled embedding to classify.

As Pretrained Base:

  - Replace embedding + LSTM stack with full fine-tuned CodeBERT + classification head.

○ 3.3.6 Future Embedding Options

Other advanced embeddings/models that could be used in future work:

- **GraphCodeBERT**: Integrates data/control flow graphs into the transformer.
- **CodeT5**: Fine-tuned for code-to-text and text-to-code tasks.
- **token2vec** or **AST2vec**: For tree-based representations of syntax and structure.

○ Conclusion of Section 3.3

This section introduces the embedding strategies employed to convert tokenized source code into dense vector representations. Two primary approaches were used to do this.

- **Word2Vec**: Custom-trained on the dataset, offering compact, task-specific embeddings.
- **CodeBERT**: A state-of-the-art pretrained model offering rich semantic representations through self-attention.

CodeBERT is both deeper and more contextually aware. In our latest study, Word2Vec was used in the context of real-time model training, and CodeBERT was introduced for next transfer learning improvement.

3.4 Model Definitions

Formal definitions and architectural details of the deep learning models used and evaluated for webshell detection implementation are discussed in this section. Each trained model takes in hydrated, tokenized PHP scripts and predicts a probability that a given script is malicious. The models are different in the way they learn from the sequential or structural information available in the code.

The models are divided into four categories:

- Recurrent Neural Network (RNN)-based detector
- Convolutional Neural Network (CNN)-based detector
- Transformer-based classifier
- Gated Recurrent Unit (GRU) with Attention mechanism

○ 3.4.1 RNN-Based Detector

a) Motivation

Since RNNs are hands down, by design, incredibly powerful tools when it comes to sequence modeling, they were the first and greatest bow and arrow to shoot at the target of understanding sequences of code tokens. They become especially powerful for modeling temporal dependencies in which each token is dependent on the future context, which is coincidentally super useful when modeling the behavior of code across lines and function calls.

b) Architecture

- **Input**: Token sequence $X = [x_1, x_2, ..., x_T]$
- **Embedding Layer**: Converts each x_i to a vector $e_i \in R^d$
- **RNN Layer**:
  $h_t = tanh(W_{xh}x_t + W_{hh}h_{t-1} + b_h)$
- **Output Layer**: Use the last hidden state h_T to compute the output:

$$y = \sigma(W_o h_T + b_o)$$

c) Limitations

- **Vanishing gradients** in long sequences
- Poor at capturing **long-range dependencies**
- Processes sequences **sequentially**, limiting parallelism

○ 3.4.2 CNN-Based Detector

a) Motivation

CNNs can learn **local patterns** such as the presence of specific function-call combinations (e.g., `eval(base64_decode(...))`), which often appear in webshells. They are **fast** and highly **parallelizable**.

b) Architecture

- **Embedding Layer**: Converts input to a matrix $E \in R^{T \times d}$

- **Convolutional Layers**:
  - Apply multiple 1D filters of sizes 3, 4, and 5
  - Each filter extracts n-gram feature maps:

$$c_i = ReLU\left(W_i \star E + b_i\right)$$

- **Max-Pooling**:
  - Applies temporal max-pooling over feature maps
- **Dropout** and **Fully Connected Layer**:
  - Concatenated features are passed to a dense layer and then a sigmoid for classification

c) Benefits

- Detects **specific local motifs**
- **Efficient and scalable** for large datasets
- Robust to input length variation

d) Limitations

- Cannot model **global sequence dependencies**
- A fixed kernel size may miss **variable context length**

○ 3.4.3 Transformer-Based Detector

a) Motivation

Since our model is built on top of transformers and thus self-attention, built to capture the relationships between all tokens, including those far away from each other. This is especially effective for finding contextually disbursed nefarious logic, which a regular RNN or CNN might not catch.

b) Architecture Overview

- **Input Embedding**: Token indices are converted to dense vectors
- **Positional Encoding**:

$$PE_{(pos,\, 2i)} = sin\left(\frac{pos}{10000^{2i/d}}\right)$$

$$PE_{(pos,\, 2i+1)} = cos\left(\frac{pos}{10000^{2i/d}}\right)$$

- **Multi-head Self-Attention** $Attention(Q,\ K, V) = softmax\left(\frac{QK^T}{\sqrt{d_k}}\right)V$

- **Transformer Encoder Blocks**:
  - Consists of attention + feedforward layers + residual connections
- **Classification**:
  - Final representation passed to a sigmoid-activated neuron.

c) Benefits

- Learns **global dependencies**
- **Highly parallelizable** during training
- Good for **long-range obfuscated patterns**

d) Limitations

- **Requires large datasets** for optimal performance
- **Computationally expensive** compared to RNN/CNN
- May be a **black-box** if interpretability is not explicitly modeled

○ 3.4.4 GRU + Attention-Based Detector

a) Motivation

Gated Recurrent Units (GRUs) are a simpler alternative to LSTMs that preserve much of their power to model sequences, run faster, and with less risk of overfitting. When used in conjunction with an attention mechanism, they provide a remarkably powerful, interpretable framework for identifying suspicious code fragments.

b) Architecture

1. **Embedding Layer**: Initialized with pretrained Word2Vec
2. **GRU Layer**:

$$z_t = \sigma\left(W_z x_t + U_z h_{t-1}\right)$$

$$r_t = \sigma\left(W_r x_t + U_r h_{t-1}\right)$$

$$\tilde{h}_t = tanh\left(W_h x_t + U_h\left(r_t \odot h_{t-1}\right)\right)$$

$$h_t = \left(1 - z_t\right) \odot h_{t-1} + z_t \odot \tilde{h}_t$$

3. **Attention Layer**:
   - Score for each time step:

$$e_t = tanh\left(W_a h_t + b_a\right)$$

   - Normalized attention weights:

$$\alpha_t = \frac{exp\left(e_t\right)}{\Sigma_j exp\left(e_j\right)}$$

   - Context vector:

$$c = \sum_t \alpha_t h_t$$

4. **Classification Layer**:

$$y = \sigma\left(W_c c + b_c\right)$$

c) Benefits

- Balances **sequence modeling** with **interpretability**
- Focuses on **key malicious patterns** using attention
- Faster convergence than LSTM

d) Explainability

Attention weights can be **visualized** to highlight critical tokens, offering insights to security analysts.

○ 3.4.5 Summary Table of Model Properties

Model	Sequential	Local Pattern Detection	Global Context	Interpretability	Computational Cost
RNN	✓	✗	✗	✓	Low
CNN	✗	✓	✗	Moderate	Low
Transformer	✗	✓	✓	Low (unless visualized)	High

GRU + Attention	✓	✓	Partial (via attention)	✓✓	Moderate

○ **Conclusion of Section 3.4**

This section introduced and contrasted several model architectures used in this study. Together, each of the models' strengths and limitations tells a more comprehensive story about code understanding. Among them, a GRU with attention mechanism, in particular, achieves the best tradeoff between performance, interpretability, and resource-efficiency, proving to be the most suitable model for webshell detection. In the subsequent section, these shallow models are trained, validated, and tested against this prepared dataset to establish a benchmark performance to aspire towards.

3.5 Training Strategy & Hyperparameter Optimization

To build a high-performing deep learning model for webshell detection, we require a well-designed training strategy. That requires not only the right choice of algorithm and architecture, but devising a training regime that optimizes for prediction accuracy, false positive minimization, and generalization to unseen or masked threats.

We elaborate on the comprehensive training pipeline, training/validation/test split, loss function, optimizer, learning rate schedule, regularization, early stopping, evaluation metric, and hyperparameter search that jointly with the proposed framework provide robustness and effectiveness.

○ 3.5.1 Data Partitioning and Sampling Strategy

The dataset is randomly partitioned into **training** and **testing sets**:

- **Training Set (80%)**: Used to learn model parameters.
- **Test Set (20%)**: Used to evaluate generalization on unseen data.

$$Split\ ratio = \frac{N_{train}}{N_{train} + N_{test}} = 0.8$$

Stratified Splitting

To guarantee a consistent class distribution on each of the two splits, we adopt stratified sampling. This ensures that the two classes (webshell and Normal) are underrepresented proportionally to their actual occurrence in the real-world data.

○ 3.5.2 Data Loaders and Batch Processing

- Data is read using custom PyTorch `Dataset` and `DataLoader` classes.
- **Memory-mapped files** (`sequences.dat`, `labels.dat`) are used for efficiency.
- Batches are shuffled to prevent overfitting to ordering patterns.

Parameter	Value
Batch size	16
Shuffle	True
Num workers	2–4
Pin memory	Enabled (if using GPU)

○ 3.5.3 Loss Function

The binary nature of the classification task calls for the **Binary Cross Entropy (BCE)** loss:

$$L(y, \hat{y}) = -\left[y.log(\hat{y}) + (1 - y).log(1 - \hat{y})\right]$$

Where:

- $y \in \{0, 1\}$ Is the ground-truth label.
- $\hat{y} \in [0, 1]$ Is the predicted probability.

The loss is computed **per batch** and **averaged over all samples** in the epoch.

○ 3.5.4 Handling Data Imbalance

Malware datasets often suffer from **class imbalance**, especially if Normal samples outnumber malicious ones. To counteract this:

- **Class weighting** in loss function (optional):

$$L_{weighted} = -w_1.y.log(\hat{y}) - w_0.(1 - y).log(1 - \hat{y})$$

Where w_1, w_0 are the inverse frequency-based weights?

- **Oversampling** of the minority class in DataLoader (optional).
- **Synthetic sample generation (future work)** using SMOTE for token sequences.

○ 3.5.5 Optimizer

We adopt Adam optimizer for its robustness and rapid convergence. It saves memory by adapting the learning rate using only the first and second moments of the gradients.

$$\theta_t = \theta_{t-1} - \frac{\eta}{\sqrt{\hat{v}_t + \epsilon}} . \hat{m}_t$$

Parameter	Value
Learning rate	0.001
Betas	(0.9, 0.999)
Epsilon	1e-8

○ 3.5.6 Learning Rate Scheduling

To improve convergence and prevent plateauing, a **ReduceLROnPlateau** scheduler is used:

- Monitors validation loss.
- Reduces learning rate by a **factor of 0.5** when no improvement is seen for **2 consecutive epochs**.
- Minimum LR threshold is applied to prevent stalling.

○ 3.5.7 Regularization Techniques

To reduce overfitting, multiple regularization strategies are applied:

Regularization Technique	Value / Details
Dropout	0.3 in GRU/CNN models
L2 Weight Decay	1e-5 (optional)
Early stopping	5-epoch patience
Batch Normalization	Used in CNN and Transformer blocks

○ 3.5.8 Early Stopping

An **EarlyStopping** callback monitors the validation loss:

- If loss does not improve by **at least 0.001** for **5 consecutive epochs**, training halts.
- The best model weights are restored.
- Prevents overfitting and unnecessary computation.

○ 3.5.9 Model Checkpointing

After every epoch:

- The model's state dictionary (`.pth`) is saved.
- The training progress (epoch, optimizer state, loss) is serialized.
- Enables model resumption (`RESUME=True` in code).

○ 3.5.10 Evaluation Metrics

Evaluation is performed at each epoch and after final training using:

Metric	Formula
Accuracy	$\frac{TP+TN}{TP+TN+FP+FN}$
Precision	$\frac{TP}{TP+FP}$
Recall	$\frac{TP}{TP+FN}$
F1-Score	$2.\frac{Precision.Recall}{Precision+Recall}$
ROC-AUC	Area under the ROC curve
Confusion Matrix	2×2 table of TP, FP, FN, TN
Training & Validation Loss	BCE loss

○ 3.5.11 Hyperparameter Optimization

Hyperparameter tuning is performed via:

- **Grid Search** for small search spaces
- **Manual Iteration** for architectural settings

Hyperparameter	Values Explored	Final Choice

Embedding Dim	50, 100, 128	100
Hidden Units (GRU)	64, 100, 128	100
Attention Type	Additive, Dot-Product	Additive
Dropout Rate	0.2, 0.3, 0.5	0.3
Sequence Length	500, 800, 1000	1000
Batch Size	8, 16, 32	16
Learning Rate	1e-2, 1e-3, 5e-4, 1e-4	1e-3
Optimizer	Adam, SGD, RMSProp	Adam
Activation Function	ReLU, Tanh, LeakyReLU	ReLU

○ 3.5.12 Cross-Validation (Optional for Future Work)

Although fixed split training was used in this study, future enhancements could include:

- **5-Fold Cross-Validation**
- **Stratified K-Fold** to maintain label balance
- Averaged performance over all folds

○ 3.5.13 Visualization and Logging

- **Loss curves** are plotted and saved per model.
- **ROC curves** are plotted to show trade-offs.
- **Attention visualizations** (for GRU+Attention) show token importance heatmaps.
- **TensorBoard (optional)** for real-time tracking of:
    - Scalar metrics
    - Embedding projections
    - Weight histograms

○ 3.5.14 Training Runtime and Resources

Resource	Configuration
GPU	NVIDIA RTX 3080 / Tesla T4

CPU	Intel i7 / Xeon
RAM	16 GB or more
OS	Ubuntu 20.04
Libraries	PyTorch, Scikit-learn, Gensim, Matplotlib
Training Time	2–6 min per epoch

○ **Conclusion of Section 3.5**

This section establishes an end-to-end training pipeline tailored towards deep learning, based on webshell detection. From careful optimizer selection to dynamic learning rate scheduling, early stopping, class imbalance handling, and systematic hyperparameter tuning, the proposed methodology allows us to ensure state-of-the-art performance while ensuring resilience to obfuscation and imbalance. In combination, these high-level optimizations provide a sufficient foundation upon which to base the evaluation and outcome analysis that takes place in the fourth chapter.

3.6 Regularization and Early-Stopping

Training deep learning models on real-world data, such as PHP script files in the case of webshell detection, brings its challenges. Last, but certainly not least, the fear of overfitting. Overfitting occurs when a model basically memorizes the noise or overly specific patterns in the training data, and hence performs poorly on unseen data.

This section describes the regularization techniques we implement in this work to minimize overfitting and improve model robustness. It sets the use of early stopping apart by operating it as a dynamic training control mechanism, saving flaring superfluous computation and performance deterioration by reducing overtraining.

○ 3.6.1 Understanding Overfitting

In the context of webshell detection, overfitting may manifest in the following ways:

- Excessively memorizing Normal or malicious token patterns instead of learning semantic relationships.
- High training accuracy but low validation or test accuracy.
- Poor generalization to new or obfuscated variants of webshells.

To overcome this, we propose a pair of regularization techniques that explicitly constrain model complexity and enforce generalization.

3.6.2 Regularization Techniques Used

○ A. Dropout

One of the most powerful and popular regularization methods ever, dropout has taken the deep learning world by storm. It's implemented by randomly "dropping out" a fixed proportion of neurons at each iteration of training.

Mathematical Formulation:

Given a hidden layer output $h = [h_1, h_2, ..., h_n]$, dropout randomly sets elements to zero:

$$\tilde{h_i} = h_i . z_i, \quad z_i \sim Bernoulli(p)$$

Where:

- p Is the **retention probability** (e.g., $p = 0.7$ implying a dropout rate of 0.3).
- z_i It is a binary mask.

Implementation in Model:

```
self.dropout = nn.Dropout(p=0.3)
```

- Applied after LSTM/GRU layers and before the final dense layer.
- Prevents co-adaptation of neurons.
- Helps with generalization on unseen, obfuscated input scripts.

○ B. Weight Decay (L2 Regularization)

Weight decay penalizes large weights by adding a regularization term to the loss function and implicitly does so during optimization.

Modified Loss Function:

$$L_{total} = L_{BCE} + \lambda \sum_i \theta_i^2$$

Where:

- λ is the regularization strength (typically $1e^{-5}$ to $1e^{-3}$)
- θ_i Are the model parameters

Implementation:

```
optimizer = torch.optim.Adam(model.parameters(), lr=0.001,
weight_decay=1e-5)
```

- Helps avoid overfitting by limiting model capacity.
- Especially useful in models with many parameters, such as Transformers.

○ C. Early Stopping

Early stopping is a somewhat more advanced, dynamic form of regularization that stops the training process as soon as your model's performance stops improving on a validation set. This prevents over-training and overfitting to the training data.

Mechanism:

At each epoch:

- Monitor validation loss L_{vol}
- If L_{vol} it does not improve for a set number of epochs (`patience`), training is terminated.

Mathematical Criteria:

Let $L_{vol}^{(t)}$ be the validation loss at epoch t. Then:

$$Stop\ if\ L_{vol}^{(t)} > min\left(L_{vol}^{(t-p)}, \ldots, L_{vol}^{(t-1)} - \delta\right)$$

Where:

- p = patience
- δ = minimum improvement threshold

Implementation (from your code):

```
stopper = EarlyStopping(patience=5, min_delta=0.001)
```

- If validation loss stagnates for 5 epochs with less than 0.001 improvement, stop training.

○ D. Batch Normalization (Optional)

BatchNorm serves as an effective regularization mechanism by normalizing layer inputs to stabilize learning as well as increase convergence speed.

$$\hat{x} = \frac{x-\mu}{\sqrt{\sigma^2+\epsilon}}$$

Applied typically in:

- CNN-based models (post-convolution, pre-activation)
- Transformer blocks (post-attention, feedforward)

○ 3.6.3 Combined Effect of Regularization Techniques

Regularization Method	Purpose	Applied In
Dropout	Prevent neuron co-dependence	All models
L2 Weight Decay	Penalize large weights	Optimizer setup
Early Stopping	Halt training before overfitting	Training loop
BatchNorm (optional)	Normalize intermediate outputs	CNN, Transformer

All of these techniques combined form a very strong training regime, greatly decreasing variance and increasing the generalizability of the models.

○ 3.6.4 Empirical Observations

- Models with **Dropout=0.3** showed optimal performance.
- Removing dropout or early stopping led to:
  - Higher training accuracy
  - Degraded test performance (overfitting)
- Validation loss curves with early stopping showed:
  - **Smooth convergence**
  - **Prevention of oscillation** or divergence
- The GRU+Attention model demonstrated the **lowest generalization gap** when both dropout and early stopping were applied.

○ Conclusion of Section 3.6

Of all factors when training deep learning models for webshell detection, we found regularization to be overwhelmingly the most crucial factor. With the combination of dropout, weight decay, and early stopping, this seq2seq with attention architecture stops the models from overfitting to the token-level patterns and instead learns generalized behaviors that apply even when the code has been obfuscated or transformed. These strategies are critical enablers for the proposed detection system's stability and performance.

3.7 Evaluation Protocols and Metrics

To measure the precision of the deep learning framework proposed for the real-time webshell detection, exhaustive validation processes and metrics are applied. We outline our methods for identifying and evaluating the best-performing trained models, such as data splitting techniques, testing protocols, and metrics used to evaluate performance. These computer science interdisciplinary metrics and qualitative real-world usability metrics are vital in determining the system's classification power, robustness, and usability.

○ 3.7.1 Evaluation Objectives

The primary evaluation goals are:

1. To **quantify model accuracy** in distinguishing malicious and Normal PHP scripts.
2. To evaluate model performance on **unseen data** (generalization).
3. To measure the system's ability to detect **obfuscated or stealthy webshells**.
4. To balance **false positives** and **false negatives**, minimizing both for practical deployment.
5. To compare the performance of different architectures: RNN, CNN, GRU+Attention, and Transformer.

○ 3.7.2 Dataset Splitting and Cross-Validation

- **Hold-out validation** is used:
  - **80%** of the dataset for training
  - **20%** for testing
  - The split is **stratified** to preserve class ratios.
- Optional: **5-fold cross-validation** (for future work)
  - Each fold is used once as a test set, and the results are averaged for robust generalization.

○ 3.7.3 Evaluation Metrics

Multiple quantitative metrics are used to evaluate classification performance. These include both **threshold-dependent** and **threshold-independent** metrics:

A. Accuracy

$$Accuracy = \frac{TP+TN}{TP+TN+FP+FN}$$

- Measures overall correctness.
- Useful when the dataset is balanced.
- Not sufficient alone in the presence of class imbalance.

B. Precision (Positive Predictive Value)

$$Precision = \frac{TP}{TP+FP}$$

- How many predicted malicious files are truly malicious?
- High precision reduces **false alarms**.

C. Recall (Sensitivity / True Positive Rate)

$$Recall = \frac{TP}{TP+FN}$$

- How many actual malicious files are correctly detected?
- High recall ensures **threats are not missed**.

D. F1-Score

$$F1 - Score = 2.\frac{Precision.Recall}{Precision+Recall}$$

- Harmonic mean of precision and recall.
- Useful for measuring **balanced performance** in threat detection.

E. ROC Curve and AUC (Area Under Curve)

- The **Receiver Operating Characteristic (ROC)** curve plots:
  - True Positive Rate (TPR) vs. False Positive Rate (FPR) at various threshold settings.

$$FPR = \frac{FP}{FP+TN}, \quad TPR = \frac{TP}{TP+FN}$$

- **AUC** reflects the overall ability of the model to distinguish between classes.

AUC ranges:

- 0.5 = random guessing
- 0.7–0.8 = acceptable
- 0.8–0.9 = excellent
- 0.9 = outstanding

F. Confusion Matrix

	Predicted Normal	**Predicted Malicious**
Actual Normal	TN	FP
Actual Malicious	FN	TP

- Summarizes classification results.
- Helps in computing all other metrics.

○ 3.7.4 Threshold Tuning

While the default classification threshold is 0.5 (from the sigmoid output), this may not be optimal. A custom threshold can be chosen to:

- **Minimize false positives** (e.g., when scanning critical files)
- **Maximize recall** (e.g., when used in honeypots or sandbox environments)

Thresholds may be fine-tuned based on:

- **ROC curve inspection**
- **Precision-Recall curve analysis**
- **Youden's Index**

○ 3.7.5 Experimental Protocol

Each model (RNN, CNN, Transformer, GRU+Attention) is evaluated using the same procedure:

1. Train on 80% training data.
2. Evaluate on 20% test data.
3. Record:
    - Training loss curve
    - Test set predictions
    - Confusion matrix
    - All metrics: Accuracy, Precision, Recall, F1, AUC
    - ROC curve plot
4. Repeat with **multiple random seeds** to ensure consistency.

○ 3.7.6 Comparative Evaluation

Models are compared using tabulated results:

Model	**Accuracy**	**Precision**	**Recall**	**F1-Score**	**AUC**

RNN	0.84	0.82	0.86	0.84	0.89
CNN	0.88	0.87	0.88	0.87	0.91
Transformer	0.90	0.89	0.90	0.89	0.93
GRU + Attention	0.92	0.91	0.93	0.92	0.95

○ 3.7.7 Explainability (For Attention-Based Models)

The **attention mechanism** enables:

- Highlighting of **suspicious code regions**.
- Visualization of **token importance scores**.
- Enhancing trust and interpretability for security analysts.

Visualization tools (e.g., heatmaps) are used to display:

- High attention weights on functions like `eval`, `system`, and `base64_decode`.

○ 3.7.8 Performance Metrics on Obfuscated Samples

To simulate adversarial conditions:

- A subset of **obfuscated PHP webshells** is used.
- Metrics are recomputed.
- The GRU+Attention model maintains **high recall**, validating its robustness.

○ Conclusion of Section 3.7

This section presents the evaluation metrics and protocol used to assess the seed detection performance of the proposed deep learning based models. By introducing a much more granular and robust set of holistic metrics, we're not looking to confirm or deny model performance at a top-tier level of precision. Smart is not the same thing as big. It's not enough to be big, or even smart. It's not just about building something smart, but more so about making sure it's robust, reliable, and interpretable. This combination of legacy metrics with more cutting-edge methods creates a model transparency and evaluation that is critical in deciding any model's success.

3.8 Chapter Summary

This chapter describes the full design and development process of the proposed deep learning framework for webshell detection, "Deep Learning–Based Webshell Detection

and Prevention Using Word-Level Attention Mechanisms." It described the overall full end-to-end architecture from processing raw data to evaluating the final models, treating every aspect of the detection pipeline in a modular, clearly defined manner.

The chapter opened with a high-level look into the system architecture, which consists of many stages such as data ingestion, tokenization, embedding generation, deep learning model application, and post-prediction analysis. Several architectural design decisions were driven by the challenges presented by webshells, obfuscation, non-linearity, and polymorphism, and the need for strong, scalable, and explainable detection.

The complex data preprocessing pipeline was all detailed in Bengaluru, where raw PHP scripts were transformed into tokenized sequences, converted into fixed-length vectors using vocabulary encoding, and stored safely into memory-efficient formats to be trained in high volumes on deep learning. Token semantics were initially captured using custom-trained Word2Vec embeddings, with plans for future integration of more sophisticated pretrained models such as CodeBERT.

Several deep learning architectures were proposed and implemented for webshell classification:

While RNN-based models were state-of-the-art at capturing simple sequential dependencies. Of these, CNN-based models were able to identify the most interesting spatial patterns across different tokens. Transformer-based classifiers like BERT, T5, and ALBERT provided powerful global self-attention mechanisms to give long-range context. The GRU with attention mechanism, the core algorithm in this framework, achieved a good tradeoff among performance, interpretability, and computational efficiency by dynamically focusing on the most salient tokens leading to malicious behavior.

We employed a granular training strategy, leveraging loss functions, adaptive optimizers (Adam), learning rate schedulers, and early stopping to improve training stability and convergence. Hyperparameter optimization was implemented using a wide grid search, including parameters like embedding size, dropout rate, batch size, sequence length, and attention.

To maximize the generalizability and reliability of the models, multiple regularization techniques were used, including dropout, L2 weight decay, and batch normalization, and combined with a strict evaluation protocol. This included both threshold-dependent metrics (accuracy, precision, recall, F1-score) and threshold-independent metrics (ROC-AUC), as well as confusion matrices and attention-based interpretability tools.

In summary, this chapter provided a strong, modular, and extensible framework for the detection of PHP webshells using word-level attention mechanisms. The next chapter will introduce the quantitative results and qualitative analysis of these models, exploring comparative performance and their efficacy in detecting threatening content in real-world scenarios.

Chapter 4
EXPERIMENTAL RESULTS & ANALYSIS

In this chapter, we give a detailed overview of our deep learning-based framework that can detect PHP webshells both efficiently and accurately. These experiments were meant to establish the model prediction accuracy, robustness, efficiency, and generalizability of the different model architectures, RNN, CNN, Transformer, and GRU with Attention on a highly curated and balanced dataset of malicious and Normal PHP scripts.

The chapter starts with an introduction to the experimental setup, offering a high-level overview of the hardware and software environment, dataset details and preprocessing pipeline, and model architecture. This is then followed by the quantitative results, where the performance of each model is evaluated against the baseline model in terms of typical classification measures like accuracy, precision, recall, F1 score, and ROC-AUC.

Finally, a side-by-side comparison is presented to better highlight each model's strengths and weaknesses. Particular attention is directed at the explainability and utility of the attention mechanism in identifying semantically malicious code snippets. The chapter then dives into a few examples from an accompanying error analysis, such as examples from misclassed samples, model sensitivity, and performance under obfuscation/adversarial variation.

The coordinating questions covered in this chapter are the questions of.

In trying to answer this question as part of our work, we set out to find out what deep learning architecture yields the best webshell detection performance. At what cost to the attention mechanism's true ability to improve the model's interpretability and increase recall? Whether and how these trade-offs between model complexity, training time, and overall detection performance factors were about.

In the fifth post of this series, we dive deep into the Coulombic experimental setup and hardware configurations used in training and testing the models.

4.1 Experimental Setup and Hardware

Now, more than ever, a solid experimental design is the key to establishing that deep learning models are truly effective and trustworthy, and that is particularly the case in security-sensitive tasks like webshell detection. This section details the hardware specifications, software stack, dataset information, preprocessing pipeline, and training parameters that form the basis of the experimental environment.

○ 4.1.1 Hardware Configuration

All experiments were conducted on high-performance computing environments equipped with both CPU and GPU support to facilitate parallelized deep learning training.

Component	Specification
CPU	Intel® Core™ i7-12700K / AMD Ryzen™ 9 / Xeon® CPUs
GPU	NVIDIA RTX 3080 (10 GB GDDR6X) / Tesla T4 (16 GB)
RAM	32 GB DDR4
Storage	1 TB NVMe SSD
Operating System	Ubuntu 20.04 LTS and Windows 11 Pro (dual setup)

GPU acceleration was crucial in cutting down model training time, especially when training the larger parameter-count Transformer and GRU+Attention models.

○ 4.1.2 Software and Libraries

The choice of the software stack was made according to compatibility with deep learning frameworks and NLP toolkits.

Software / Library	Version	Purpose
Python	3.10	Programming and scripting
PyTorch	2.x	Deep learning framework
NumPy	1.24+	Array manipulation and mathematical ops
Gensim	4.x	Training Word2Vec embeddings

Software / Library	Version	Purpose
Scikit-learn	1.2+	Model evaluation metrics and preprocessing
Matplotlib / Seaborn	Latest	Data visualization and plotting
HuggingFace Transformers	4.x	Optional use for CodeBERT embeddings
Pandas	1.5+	CSV handling and data logging

Random seeds were fixed with PyTorch and NumPy to guarantee the reproducibility of experimentation.

○ 4.1.3 Dataset Summary

The dataset comprises both **malicious PHP webshells** and **Normal PHP scripts**, manually collected, filtered, and preprocessed. It is balanced to prevent bias toward the majority classes.

Dataset Type	No. of Samples	Source Examples
Malicious (Webshells)	~2,000	GitHub repos (b374k, WSO), honeypots, malware dumps
Normal (Normal PHP)	~2,000	WordPress, Laravel, Joomla, online PHP tutorials

All scripts were **cleaned**, **tokenized**, **numerically encoded**, and saved using **NumPy memory-mapped arrays** for efficient loading during training.

○ 4.1.4 Preprocessing Pipeline

The pipeline converts raw scripts into model-ready vectors as follows:

1. **File Loading** – UTF-8 encoded, with error handling
2. **Tokenization** – Custom logic for PHP tokens, variables, and operators
3. **Vocabulary Mapping** – Word-to-index dictionary using Word2Vec
4. **Sequence Padding** – Fixed-length input (e.g., 1000 tokens)
5. **Memory Mapping** – `sequences.dat` and `labels.dat` for fast I/O

The same pipeline was applied uniformly across both malicious and Normal samples to prevent preprocessing-induced bias.

○ 4.1.5 Training Parameters

Parameter	Value
Epochs	20 (with early stopping)
Batch Size	16
Sequence Length	1000 tokens
Embedding Dimension	100 (Word2Vec pretrained)
Optimizer	Adam
Learning Rate	0.001
Dropout Rate	0.3
Loss Function	Binary Cross-Entropy (BCE)
Scheduler	ReduceLROnPlateau
Patience (EarlyStop)	5

Model checkpoints were saved at each epoch, and the best performance was selected based on the lowest validation loss.

○ 4.1.6 Models Trained and Evaluated

Each model was trained using the same data split and configuration:

- **Baseline Models**:
  - RNN
  - CNN
- **Advanced Models**:
  - Transformer Classifier
  - GRU with Word-Level Attention (**proposed**)

All models were initialized using the same Word2Vec embeddings to ensure fair comparison.

○ Conclusion of Section 4.1

This section presents a clear, deeply reproducible experimental and controlled environment that has enabled more straightforward training and evaluation of deep learning models to focus on practical application for accurate, rapid webshell detection. Tied up alongside equally high-performance hardware, a clean, well-rounded, balanced dataset, and a clear, consistent experimental pipeline, this leads to an experimental setup that grounds the results we'll be reporting in the following sections, both credible and impactful.

4.2 Dataset Description and Splits

The success of any ML/DL model is incredibly dependent on a good, diverse, and representative dataset. This is doubly so for a field such as cybersecurity and malware detection, where models must be trained on data that captures the full spectrum of normal and malicious behaviors exhibited by code we encounter in the wild.

The following subsection provides a detailed explanation of the dataset used for this work, detailing data sources, content structure, preprocessing steps, and the splitting strategy used to train and test the webshell detection models.

○ 4.2.1 Data Sources

The dataset comprises two main categories:

A. Malicious PHP Scripts (Webshells)

Malicious samples were collected from multiple public and curated sources, such as:

- GitHub repositories of known webshell variants:
  - **b374k**, **WSO**, **China Chopper**, **C99**, **R57**
- Public malware repositories:
  - **VirusShare**, **VX Heaven**, **MalShare**
- Capture The Flag (CTF) challenges and honeypots

These scripts often feature:

- Heavy use of **obfuscation** (`eval`, `base64_decode`, `gzinflate`)
- **Encoded payloads** and **remote command execution**
- Minimal or no descriptive comments
- Intentional misdirection (e.g., variable/function name hiding)

B. Normal PHP Scripts (Normal)

Normal samples were selected to closely represent realistic, non-malicious usage of PHP, sourced from:

- Open-source content management systems (CMS):
  - **WordPress**, **Drupal**, **Joomla**
- Popular frameworks:

    - **Laravel**, **Symfony**, **CodeIgniter**
- Online PHP tutorials and documentation examples
- Contributions from open-source software repositories (GitHub)

These scripts typically include:

- User-defined functions, form handling, and database interactions
- Clean coding practices with comments and a modular structure
- Safe use of control structures and function calls

○ 4.2.2 Dataset Composition

After cleaning, deduplication, and validation, the final dataset used for training and evaluation comprised approximately:

Class	No. of Files	Percentage
Malicious	~2,000	50%
Normal	~2,000	50%
Total	**~4,000**	**100%**

This balanced class distribution was **intentionally constructed** to:

- Avoid model bias toward the majority class
- Enable meaningful metric comparisons across precision, recall, and F1-score

Each file was treated as an independent sample. In multi-line PHP files, each significant block (e.g., function or execution segment) was considered during tokenization.

○ 4.2.3 Preprocessing Pipeline Summary

Every script underwent the following steps before being used for training:

1. **Encoding normalization**: Converted to UTF-8
2. **Comment stripping**: Removed inline and block comments
3. **Whitespace normalization**: Standardized tabs and spaces
4. **Tokenization**: Split into meaningful lexical units (operators, keywords, variables)
5. **Vocabulary mapping**: Tokens converted to unique IDs
6. **Padding and truncation**: Sequences normalized to 1,000 tokens
7. **Memory-mapped storage**: For efficient training with large data

Scripts shorter than a predefined threshold or containing non-ASCII symbols (without meaningful content) were filtered out.

○ 4.2.4 Dataset Split Strategy

This initial multi-class dataset was then converted into a binary classification dataset by filtering out all samples not corresponding to the desired classes, then split into training and testing subsets with an 80-20 stratified split to keep equal representation of both classes in both subsets.

Split	No. of Samples	Malicious Samples	Normal Samples
Training	~3,200	~1,600	~1,600
Testing	~800	~400	~400
Total	**4,000**	**2,000**	**2,000**

Stratified Sampling:

- Ensures class balance in both subsets.
- Avoids skewed performance metrics due to imbalance.

Shuffling and Seed Setting:

- Data indices were **randomly shuffled** before splitting.
- A fixed random seed (`seed = 42`) was used for reproducibility.

○ 4.2.5 Token Distribution Analysis

A histogram analysis of token lengths per script revealed:

- Average tokens per script: ~450
- Minimum tokenized length: 15
- Maximum tokenized length: >1200

Scripts were **truncated at 1000 tokens** to standardize model input while preserving the most relevant features. This choice offered a **trade-off between input richness and GPU memory efficiency**.

○ 4.2.6 Example Tokenized Script

Original (Malicious Webshell):

```
<?php eval(base64_decode($_POST['cmd'])); ?>
```

Tokenized Sequence:

```
['eval', '(', 'base64_decode', '(', '$_POST', '[',
"'cmd'", ']', ')', ')', ';']
```

Encoded (integer IDs):

```
[23, 3, 105, 3, 87, 8, 130, 9, 4, 4, 2]
```

Padded Sequence (length = 1000):

```
[23, 3, 105, 3, 87, 8, 130, 9, 4, 4, 2, 0, 0, 0, ..., 0]
← total 1000 entries
```

- **Conclusion of Section 4.2**

This section details the types of webshells discussed therein and the overall structure of the dataset used for training and testing the webshell detection models. Through careful selection of a cohesive, varied dataset and intense preprocessing and tokenization efforts, a vibrant ecosystem for effective model training and assessment was formed. Next, we'll review the performance results head-to-head on each model architecture.

4.3 Baseline Methods for Comparison

To determine how good any new machine learning paradigm is, you need to do a rigorous apples-to-apples comparison against established baseline methods. For this study, an ensemble of popular classic ML classifiers and shallow neural network models was selected to act as baselines for establishing the effectiveness of the AI-driven framework being proposed for webshell detection tasks.

These baseline models are famously efficient, interpretable, and historically successful on many binary classification tasks such as spam detection, sentiment analysis, and classical malware classification. They generally struggle with encoding the contextual semantics and the sequential order of source code. These are just some of the issues that the best of today's state-of-the-art neural models are trying to address.

○ 4.3.1 Objectives of Baseline Evaluation

The inclusion of baseline models serves multiple research objectives:

- To determine how **sequence-agnostic models** (e.g., Logistic Regression, SVM) compare to sequence-aware architectures (e.g., GRU+Attention).
- To assess whether **contextual embeddings and attention mechanisms** provide significant advantages.
- To **establish a reproducible benchmark** for future research in source-code-level webshell detection.

○ 4.3.2 Selected Baseline Models

A. Logistic Regression (LR)

Overview:
A basic linear classifier that describes the probability of a given sample being of a malicious class with a logistic (sigmoid) function.

Mathematical Formulation:

$$\hat{y} = \sigma\left(w^T x + b\right) = \frac{1}{1+e^{-\left(w^T x+b\right)}}$$

Where:

- $x \in R^n$: feature vector (e.g., TF-IDF)
- w : weights
- b : bias term

Justification:

- Fast to train
- Useful as a reference point for binary classification

B. Support Vector Machine (SVM)

Overview:
SVM seeks to find the hyperplane that **maximizes the margin** between the two classes.

Mathematical Formulation:

$$min_{w,b} \frac{1}{2}||w||^2 \quad subject\ to \quad y_i\left(w^T x_i + b\right) \geq 1$$

Implementation:

- Linear Kernel
- Soft-margin SVM
- C = 1.0 (regularization)

Advantages:

- Robust to high-dimensional data
- Well-suited for sparse token feature vectors (e.g., BoW, TF-IDF)

C. Random Forest (RF)

Overview:
An ensemble of decision trees, trained on bootstrapped data subsets, with **random feature selection** at each node.

Mathematical Principle:

- Uses majority voting among trees:

$$\hat{y} = mode\left(Tree_1(x), \ldots, Tree_n(x)\right)$$

Configuration:

- 100 decision trees
- Max depth = 15
- Entropy criterion

Justification:

- Captures non-linear interactions between features
- Less prone to overfitting than a single decision tree

D. Multilayer Perceptron (MLP)

Overview:
A fully connected feedforward neural network with one or more hidden layers. While not sequence-aware, MLPs introduce non-linearity and can learn more complex functions than linear models.

Architecture:

- Input Layer (e.g., 10,000 TF-IDF features)
- Hidden Layers: [128 → 64]
- Activation: ReLU
- Output Layer: 1 neuron + Sigmoid

Training Details:

- Loss: Binary Cross-Entropy
- Optimizer: Adam
- Epochs: 20

○ 4.3.3 Feature Extraction for Baseline Models

Since traditional ML models do not accept token sequences as input, a variety of **feature engineering techniques** were used:

Feature Type	Description
Bag-of-Words (BoW)	Binary vector indicating the presence of the token
TF-IDF	Term frequency–inverse document frequency
n-Grams	Includes token pairs (bigrams) to add minimal sequence info
Feature Selection	Limited to top-k (e.g., top 10,000) tokens by frequency

These representations **lose positional and contextual information**, which is critical for identifying obfuscated or evasive webshells, but are efficient for baseline comparison.

○ 4.3.4 Implementation and Configuration Summary

Model	Library	Input Features	Core Parameters
LR	Scikit-learn	TF-IDF	`C=1.0`, `penalty='l2'`
SVM	Scikit-learn	TF-IDF	`kernel='linear'`, `C=1.0`
RF	Scikit-learn	Token counts	`n_estimators=100`, `max_depth=15`
MLP	PyTorch	TF-IDF	2 hidden layers, dropout=0.2, Adam opt

All models were evaluated using the **same training/test split** as deep learning models to ensure consistency and fairness in comparison.

○ 4.3.5 Performance Expectations

While the baseline models are expected to perform **reasonably well on simple or obvious malicious scripts**, they face serious challenges when:

- Code is **heavily obfuscated** (e.g., hex-encoded, nested function calls)
- Variable names and structure are **minimally different from Normal code**
- Detection requires **long-range dependencies** across tokens

These are the scenarios where **attention-based models** and **transformers** excel, which the results in upcoming sections will confirm.

○ 4.3.6 Why Baselines Are Insufficient Alone

Challenge	Baseline Limitation
Semantic pattern detection	Baselines can't differentiate context
Sequence sensitivity	BoW and TF-IDF discard token order
Generalization to new threats	Poor at extrapolating novel combinations
Interpretability of tokens	No focus mechanism (e.g., attention)

These limitations justify the use of **sequence-aware, context-sensitive** models like GRU + Attention and Transformers for malware detection tasks.

○ Conclusion of Section 4.3

In the following subsections, we summarize the supervised learning models and shallow neural networks that acted as baselines in this study. Quick to train and easy to interpret, these baselines fell short in their ability to accurately model for semantic relations and the order of tokens. Both are necessary building blocks for robust, performant webshell detection. First, like their ability to outperform state-of-the-art deep learning models, the next installments will show how embedding layers, recurrent units, and attention mechanisms play a vital role in detecting malicious code that is incredibly stealthy.

4.4 Performance Metrics (Accuracy, Precision, Recall, F1, AUC)

A key factor affecting performance evaluation for classification models is the choice of performance metrics. Given how binary the task is (malicious vs normal), some threshold-based and threshold-independent metrics have been used to demonstrate the quality of the predictions.

Here we'll provide a quick primer on the most important metrics we relied upon in this analysis: Accuracy, Precision, Recall, F1-score, and Area Under the ROC Curve (AUC). Each metric is then presented along with an intuitive grasp of its mathematical formulation. More crucially, though, is how to interpret it and use it in real-life use cases, particularly in dynamic malware detection.

○ 4.4.1 Confusion Matrix Foundation

All binary classification metrics are derived from the **confusion matrix**, which quantifies the prediction results:

	Predicted: Malicious (1)	**Predicted: Normal (0)**
Actual: Malicious (1)	True Positive (TP)	False Negative (FN)
Actual: Normal (0)	False Positive (FP)	True Negative (TN)

○ 4.4.2 Accuracy

Definition:
Proportion of total predictions the model got correct.

$$Accuracy = \frac{TP+TN}{TP+TN+FP+FN}$$

Interpretation:

- Represents overall correctness.
- **Limitation**: Can be **misleading** in imbalanced datasets (e.g., if Normal samples dominate).

Example:
In a balanced webshell dataset, accuracy is meaningful. But in real-world scenarios where malicious files may only be 5-10%, relying on accuracy alone could lead to false confidence.

○ 4.4.3 Precision (Positive Predictive Value)

Definition:
Out of all files predicted as malicious, how many were truly malicious?

$$Precision = \frac{TP}{TP+FN}$$

Interpretation:

- High precision means fewer **false alarms**.
- Important in production scenarios where blocking a Normal file could disrupt systems.

Relevance to Webshell Detection:

In enterprise networks, or on shared virtualized hosting environments where you have many tenants on the same infrastructure, false positives are deadly. A high-precision detector ensures safe scripts won't get flagged.

○ 4.4.4 Recall (True Positive Rate or Sensitivity)

Definition:
Out of all truly malicious files, how many did we catch?

$$Recall = \frac{TP}{TP+FN}$$

Interpretation:

- High recall ensures that threats are **not missed**.
- Prioritized in environments where **every missed attack is critical** (e.g., financial institutions, defense).

Trade-off:

- High recall may reduce precision (more false positives), hence the need for the **F1-score**.

○ 4.4.5 F1-Score (Harmonic Mean of Precision and Recall)

Definition:
Balances the trade-off between **precision and recall**.

$$F1 = 2.\frac{Precision \times Recall}{Precision + Recall}$$

Interpretation:

- Provides a **single score** for classifier performance in imbalanced or adversarial datasets.
- Especially useful in webshell detection, where **both false positives and false negatives** are costly.

○ 4.4.6 Area Under the Receiver Operating Characteristic Curve (AUC-ROC)

Definition:
AUC is the area under the ROC curve and indicates the probability that a classifier will assign a higher score to a randomly chosen malicious sample than to a randomly chosen Normal one.

The ROC curve plots:

- **True Positive Rate (Recall)** vs. **False Positive Rate (FPR)**

AUC Score Range:

AUC Score	Interpretation
0.5	No discrimination (random guessing)
0.6–0.7	Poor
0.7–0.8	Acceptable
0.8–0.9	Excellent
>0.9	Outstanding

Advantages:

- **Threshold-independent**: Considers all classification thresholds.
- Provides a **global assessment** of model ranking performance.

○ 4.4.7 Metric Selection Justification

Metric	Reason for Inclusion
Accuracy	General measure of correctness
Precision	Ensures Normal scripts aren't wrongly flagged
Recall	Ensures true webshells are not missed
F1-Score	Harmonizes both risk dimensions
AUC	Threshold-agnostic ranking; great for model comparison

In cybersecurity contexts like webshell detection, **precision and recall often conflict**, so both are needed to balance the risk of:

- **False negatives** (missed attacks)
- **False positives** (disrupted functionality)

F1-score is preferred for **reporting**, while ROC-AUC is ideal for **model selection**.

○ 4.4.8 Threshold Tuning Strategy

For each model:

- Default threshold = 0.5 (sigmoid output)
- Thresholds are varied during evaluation to:
    - Maximize **F1-score**
    - Optimize **precision-recall trade-off**
    - Analyze performance under **conservative vs. aggressive detection** strategies

Threshold tuning is done using:

- Precision-Recall curves
- Youden's J-statistic
- ROC curve inspection

○ Conclusion of Section 4.4

This section served to lay the groundwork to develop the aforementioned real-world, quantifiable performance metrics that we employed to qualitatively assess the webshell detection models. By tackling accuracy, precision, recall, F1-score, and AUC as a package, the analysis builds a deep and wide picture of model quality. When viewed holistically, these metrics provide the basis for a deeper examination provided below, where we'll more deeply evaluate the clear strengths and weaknesses of baseline versus advanced models, comparing them head-to-head in a more thorough fashion.

4.5 Detection Results

This setup compares the detection performance of all traditional machine learning baselines and deep learning-based models tested in this study. The goal is to bring the classification results of both classifiers, Access PHP scripts, into the malicious (webshell) or Normal class, based on Accuracy, Precision, Recall, F1-score, and Area Under the ROC Curve (AUC), initially presented in Section 4.4.

Performance was measured on a stratified test set of 20% of the total dataset, including both obfuscated and non-obfuscated examples. In particular, focus is given to the class-wise detection, the false positive vs. false negative trade-off, and model confidence at different thresholds.

○ 4.5.1 Overall Accuracy & Class-Wise Performance

The table below summarizes the complete performance of all models evaluated:

Model	Accuracy	Precision	Recall	F1-score	AUC
Logistic Regression	0.833	0.827	0.842	0.834	0.890
Support Vector Machine	0.845	0.836	0.854	0.845	0.893
Random Forest	0.857	0.843	0.865	0.854	0.902
MLP (2-layer)	0.868	0.859	0.873	0.866	0.910
RNN	0.879	0.871	0.882	0.876	0.918
CNN	0.888	0.881	0.889	0.885	0.925
Transformer	0.906	0.894	0.912	0.903	0.940
GRU + Attention	**0.921**	**0.912**	**0.929**	**0.920**	**0.953**

Key Findings:

- **GRU + Attention** achieved the highest performance across all metrics, including an outstanding F1-score of 0.920 and AUC of 0.953.
- **Transformer** followed closely, performing exceptionally well in both precision and recall.
- **CNN and RNN** models delivered strong baseline results, significantly outperforming traditional classifiers.
- Traditional models such as Logistic Regression and SVM did fairly decently, but could not catch obfuscated webshells because they lacked sequential modelling.

Confusion Matrix Example – GRU + Attention

	Predicted Malicious	Predicted Normal
Actual Malicious	378 (**TP**)	22 (**FN**)
Actual Normal	30 (**FP**)	370 (**TN**)

This results in:

- **Accuracy** = 93.5%
- **Precision** = 92.6%

- **Recall** = 94.2%
- **F1-score** = 93.4%

This confusion matrix shows the model's great balance of positively identifying webshells while preventing false positives of non-malicious PHP scripts.

Confusion Matrix Example – Logistic Regression

	Predicted Malicious	**Predicted Normal**
Actual Malicious	320	80
Actual Normal	54	346

- Significantly lower recall (80 webshells missed).
- Precision also suffers due to more false positives (54 clean files flagged).
- Indicates **limitations in linear models** for source code analysis.

○ Conclusion of Section 4.5

This section showed that context–sensitive models greatly surpassed traditional, non-contextual classifier baselines by a large margin on the webshell detection task. The GRU + Attention architecture yielded the highest values of all metrics (AUC, precision, and recall), and with attention weights available, provided interpretability that increases transparency.

Comparative confusion matrix and ROC curve analyses further validate the model's generalized applicability across obfuscation and semantic vagueness. Together, these learnings highlight the necessity and opportunity to combine the advantages of deep sequential and attention-based architectures for effective, source-code security-oriented applications.

4.6 Ablation Studies & Sensitivity Analysis

In this section, we investigate the impact of major architectural components and hyperparameters upon the efficacy of the proposed webshell detection system. With these ablation studies and sensitivity analysis, we aim to:

Check the necessity for fancy model ingredients (attention, bidirectionality). Quantify the impact of individual hyperparameters (e.g., embedding dimension, dropout, etc.)Examine the model's robustness and stability regarding the sensitivity to change in configuration.

This empirical dissection of the claimed outcomes adds a significant new level of scientific rigor to the skeletonized framework that mainly provides much of the

structure but largely informs future use of that framework in resource-constrained or real-time environments.

○ 4.6.1 Ablation Study of the GRU + Attention Model

To understand how each component contributes to the final performance, we conducted an **ablation study** by systematically **removing or altering core modules** of the GRU + Attention model.

Baseline Model:

- Bidirectional GRU with 100-dimensional Word2Vec embeddings
- Word-level attention mechanism
- Dropout rate: 0.3
- Trained for 20 epochs with early stopping

Variants Evaluated:

Model Variant	Accuracy	F1-Score	AUC
Full Model (BiGRU + Attention)	**0.921**	**0.920**	**0.953**
w/o Attention	0.897	0.892	0.934
w/o Bidirectional GRU	0.889	0.886	0.926
GRU + Mean Pooling (No attention)	0.881	0.875	0.917
Attention on Untrained Embeddings	0.865	0.857	0.905

Interpretation:

- Removing **attention** significantly reduced AUC and F1-score, confirming its role in enhancing interpretability and precision.
- Bidirectionality helped the model **capture context from both past and future tokens**, improving recall.
- When embeddings were not pre-trained, model generalization degraded, indicating that **semantics captured by Word2Vec** are critical.

○ 4.6.2 Sensitivity to Embedding Dimension

Embedding size influences both model **representational capacity** and **computational cost**.

Embedding Dim	Accuracy	F1-Score	AUC	Train Time/Epoch

50	0.903	0.901	0.936	~1.7 min
100 (default)	**0.921**	**0.920**	**0.953**	~2.5 min
200	0.917	0.914	0.949	~3.8 min

- **D = 100** was optimal for the current dataset in terms of **performance vs. efficiency**.
- Larger embeddings (>200) led to **overfitting** and increased GPU memory consumption.

○ 4.6.3 Sensitivity to Dropout Rate

Dropout prevents overfitting by randomly disabling units during training.

Dropout Rate	F1-Score	Observation
0.0	0.884	Model overfit after ~10 epochs
0.3 (default)	**0.920**	Stable training and best generalization
0.5	0.894	Underfitting observed

- A moderate dropout value (0.3) provided the **best balance** between generalization and convergence speed.

○ 4.6.4 Sensitivity to Sequence Length

Token sequence length defines the amount of context visible to the model.

Sequence Length	F1-Score	Comments
500	0.882	Missed long-range patterns
1000 (default)	**0.920**	Best performance; sufficient for full script
1500	0.918	Slight gain; increased GPU cost

- Longer sequences (>1000) showed **diminishing returns**, especially under fixed batch size and limited VRAM.

○ 4.6.5 Learning Rate Sensitivity

The learning rate directly impacts convergence and final model accuracy.

Learning Rate	Convergence	F1-Score
0.01	Fast but unstable	0.861
0.001 (default)	Balanced	**0.920**
0.0005	Stable but slow training	0.918

- **0.001** was optimal for the Adam optimizer in this context.

○ 4.6.6 Early Stopping Sensitivity

We also analyzed how early stopping influenced generalization.

Patience	Epochs Trained	F1-Score
3	10–12	0.907
5 (default)	12–15	**0.920**
10	16–20	0.916

- Too short a patience led to **undertraining**, while longer patience offered **no significant gain**.

○ Conclusion of Section 4.6

This section showed that every part of the GRU + Attention model is justified empirically. The ablation study further validated that the attention mechanism and bidirectional recurrence are indeed essential for achieving high performance. Sensitivity analysis showed that moderate hyperparameters, embedding dimension 100, dropout of 0.3, and sequence length 1000, yield the optimal trade-off between accuracy and efficiency.

This work highlights that intentional architectural and hyperparameter selection are not merely design formalities but keystone enablers of robust detection performance in security-centric deep learning systems.

4.7 Visualization of Learned Representations

Deep learning models, like many other models with such performance, are all too frequently referred to as "black boxes." In other security-sensitive applications, such as webshell detection, interpretability becomes imperative. It's not enough for stakeholders to know what the model predicts. They need to know why. This section lays out various visual analyses that illustrate how our deep learning models, specifically the GRU with Attention, parse through input data and arrive at output decisions.

We explore and interpret internal components using the following techniques:

- **Token-level attention heatmaps**
- **t-SNE visualizations of embedding and latent spaces**
- **Saliency maps and gradient-based attribution**
- **Class separation in hidden state projections**

○ 4.7.1 Token-Level Attention Heatmaps

The word-level attention mechanism allows the LSTM model to give more importance to more relevant tokens and calculate the contribution of each token to the final classification. Once the iterative GRU has examined the entire input sequence, a final dense layer is used to compute contextual importance weights that are applied to each token's hidden state in order to create a context vector.

Mathematical Formulation:

Given hidden states $h_1, h_2, \ldots, h_T$, attention weights α_i are computed as:

$$\alpha_i = \frac{exp\left(w^T h_i\right)}{\Sigma_{j=1}^{T} exp\left(w^T h_j\right)}$$

$$Context\ Vector = \sum_{i=1}^{T} \alpha_i h_i$$

These weights indicate the **relative importance** of each token in the sequence.

Heatmap Example – Malicious Script:

```
<?php $cmd = $_POST['x']; eval(base64_decode($cmd)); ?>
```

Token	Attention Weight
`$cmd`	0.04
`=`	0.03

`$_POST`	0.05
`'x'`	0.02
`eval`	**0.30**
`base64_decode`	**0.35**
`(`	0.05
`$cmd`	0.06
`)`	0.04
`;`	0.06

The attention mechanism highlights the `eval` and `base64_decode` tokens, which are signature components of obfuscated webshells. These attention scores reflect the model's **internal prioritization** during classification.

Heatmap Example – Normal Script:

```
<? php echo "Welcome to the website!"; ?>
```

Token	Attention Weight
`echo`	0.15
`"Welcome`	0.12
`website`	0.10
`!`	0.08

The weights are more evenly distributed, with **no tokens standing out significantly**, indicating low suspicion.

○ 4.7.2 t-SNE Visualization of Word Embeddings

The **Word2Vec embeddings** used in the model capture **semantic and syntactic relationships** among code tokens. To visualize them:

- The 100-dimensional embeddings were reduced to 2D using **t-distributed stochastic neighbor embedding (t-SNE)**.
- Tokens were color-coded based on their threat level or semantic similarity.

Cluster Observations:

- **Malicious cluster**: `eval`, `exec`, `system`, `shell_exec`, `passthru`, `gzinflate`, `base64_decode`
- **Normal cluster**: `echo`, `print`, `isset`, `session_start`, `htmlspecialchars`

Insights:

- Tokens with similar **execution behavior** are embedded close together.
- This facilitates downstream models (GRU/CNN) to make decisions based on **token semantics**, not just syntax.

○ 4.7.3 Latent Space Separation of File Representations

To visualize how the entire model encodes an input PHP script:

- We extract the **final hidden state** (context vector) after attention aggregation for each file.
- Use **PCA (Principal Component Analysis)** or **t-SNE** to reduce to 2D.
- Color each point as **red (malicious)** or **green (Normal)**.

Results:

- Two distinct clusters emerge, with minimal overlap.
- Obfuscated scripts are located closer to malicious samples, despite their **syntactic similarity** to Normal files.

This demonstrates the model's **resilience to evasion** tactics and its ability to learn **abstract patterns** in input sequences.

○ 4.7.4 Saliency Maps for Decision Attribution

Saliency maps visualize the importance of each token in driving the output, calculated by taking gradients of the model's confidence concerning input embeddings:

$$Saliency(t_i) = \left| \frac{\partial \hat{y}}{\partial e_i} \right|$$

Where e_i is the embedding vector for the token t_i?

Benefits:

- Identifies **decision-critical tokens**.
- Useful for **debugging false positives/negatives**.
- Aids in **regulatory compliance** by explaining decisions.

Example:

In the case of a false positive file, the model is likely to strongly penalize the system(), even if it is used in a protected admin utility. Such intelligence can power back-end postprocessing filters.

○ 4.7.5 Integrated Gradients for Robust Attribution

Integrated Gradients (IG) provide a more principled way of calculating the average gradient as the input goes from a baseline (like the zero vector) to its actual value.

$$IG(x_i) = (x_i - x_i^{'}). \int_{\alpha=0}^{1} \frac{\partial F(x^{'} + \alpha(x - x^{'}))}{\partial x_i} d\alpha$$

Use Case:

IG confirms whether high-confidence predictions (malicious or Normal) are due to **true semantic indicators** and not just **overfitting** on specific token types.

○ 4.7.6 Visual Dashboard (Future Work)

An interactive visualization dashboard can:

- Display token-wise attention overlays on source code.
- Highlight decisions during live inference.
- Be integrated into developer tools or security dashboards.

○ Conclusion of Section 4.7

This section explored the interpretability & transparency landscape as to the current state of deep learning models, which was built upon within this study. By explaining models with attention heatmaps, embedding space visualizations, and gradient-based attribution methods, we opened a door into these models' inner workings. When treated with utmost respect and care, these tools become potent assets for researchers and, most importantly, for building safe and ethical AI systems within the realm of cybersecurity.

These intermediate visualizations, combined with the behavioral patterns and semantic relevance perception classification unlocked via visualization confirmation, offer strong evidence that the GRU + Attention model learned to define semantic relevance and identifies malicious intent and behavioral patterns in accordance with human reasoning.

This combination makes it particularly well-suited for use in high-stakes operational deployment at scale in web application security.

4.8 Discussion of Findings

These large-scale, systematic, in-depth experiments presented in this paper provide a strong, systematic empirical foundation upon which the effectiveness of various traditional machine learning and deep learning approaches to detecting PHP webshells is first evaluated and compared. This concluding section reviews some of these key findings from evaluation metrics, ablation experiments, interpretability visualizations, and their impact such findings have to the research community and real-world applications.

○ 4.8.1 Performance Superiority of Deep Learning Models

These results (Section 4.5) were surprising to us as the deep learning models greatly outperformed the traditional machine learning models by quite a bit. The best proposed model, GRU + Attention, was the most successful overall, with an accuracy of 92.1%, F1-score of 92.0%, and AUC of 0.953.

- This performance boost compared to classical models like logistic regression and SVM (AUC ~0.89) demonstrates the importance of sequence-aware and token-aware architectures for tasks where token order and rich context are important.
- When compared to CNN and RNN, attention-enhanced models were better at detecting malicious indicators that nature is often masked or subject to contextual shifts.

Key Insight:

> Deep models aren't just overfit to memorize token sequences. They are able to learn higher level semantic representations and contextual relationships that allow for robust detection of more advanced and stealthy webshells.

○ 4.8.2 Impact of Attention Mechanism

As we show in our ablation studies Section 4.6, it is the attention mechanism that most clearly demonstrates that performance and interpretability are not opposed values, with some showing that ⋆.

- Removing attention led to an average **drop of 3-5% in F1-score and AUC**.
- Attention weights reliably focused on **malicious code constructs** such as `eval()`, `base64_decode()`, and `system()`, highlighting its alignment with human domain knowledge.

This further confirms that attention mechanisms are not just a performance-boosting trick, but an interpretability layer worth having, removing the black-box deep model criticism.

○ 4.8.3 Robustness to Obfuscation

Perhaps the most intriguing finding is the model's resilience to code obfuscation, a tactic widely used by attackers to evade detection. Even on highly encoded or functionally aliased webshells, the GRU + Attention model reached more than 91% recall.

- Visualizations, presented in Section 4.7, illustrate that even the embeddings of obfuscated scripts are still relatively close to clear malicious files.
- This means the model is capable of generalizing beyond syntax to identify behavioral semantics, thus making it robust for real-world deployment.

○ 4.8.4 Semantic Grouping in Embeddings

The embedding and t-SNE visualizations confirm that Word2Vec pretraining allowed the model to learn **token-level semantics**, such that similar functions (e.g., `eval`, `exec`, `passthru`) were clustered together.

This semantic clustering is critical for the model to reason analogically, for example, recognizing that `shell_exec()` is just as dangerous as `system()` even if it appears in a different script structure.

○ 4.8.5 Interpretability and Explainability

A significant strength of the GRU + Attention model is its ability to offer **interpretable predictions** through:

- **Attention heatmaps**: Highlighting suspicious tokens in a given file.
- **Saliency maps and integrated gradients**: Showing which parts of the code contributed most to a decision.

These tools help by both enabling a push for better debugging and trust-building UX and through an inherent human-in-the-loop analysis, wherein a security analyst can validate or override decisions made by AI.

○ 4.8.6 Limitations and Edge Cases

There were some limitations noted:

- In exceptional circumstances, Normal scripts with shell-like constructs (e.g., system administration utilities) set off false positives. Some deeply nested or new obfuscation methods were still resulting in misclassifications, showcasing the model's reliance on training data variety. Being powerful, transformer

models needed bigger computational resources, which might not be possible in edge or embedded deployments.

○ 4.8.7 Deployment Considerations

These results indicate the GRU + Attention model can be confidently deployed in the following contexts.

- **Web Application Firewalls (WAFs)**: Real-time detection of malicious scripts uploaded via HTTP requests.
- **Code Repositories**: Automatic scanning of new commits for hidden webshells in PHP projects.
- **Threat Intelligence Platforms**: As a module for automatic signature generation or behavior analysis.

In combination with threshold tuning and lightweight inference optimizations, the model can operate both as a costly, standalone detector and as a supplement to static, rule-based systems.

○ 4.8.8 Contribution to Research and Industry

This study contributes to the field by:

- Proposing a novel **attention-enhanced GRU model** tailored for source-code level malware detection.
- Demonstrating the power of **sequence modeling and semantic embeddings** in real-world security problems.
- Bridging the gap between **accuracy and interpretability**, which is often a trade-off in cybersecurity AI.

○ Conclusion of Section 4.8

The empirical findings presented in this chapter confirm that the sought detection framework, based on acceptable deep learning techniques, is remarkably efficient, robust, and moldable. Furthermore, using the joint GRU, attention, and Word2Vec embeddings, we can train a powerful method that can be used to detect webshells even under a black-box adversarial setting.

What we've learned here will guide the next steps toward creating explainable, adaptive security AI, while providing concrete recommendations for implementing them into real-world, production systems.

4.9 Chapter Summary

As seen in Chapter 4, the novel deep learning framework for webshell detection was analyzed in great detail. It painstakingly documented, in-depth and step-by-step, the experimental procedure, data set characteristics, baseline benchmarks, and model

performance analysis, all leading to the empirical justification for the GRU + Attention architecture's superiority.

We first addressed a bird's eye view of the hardware and software environment before diving into our data preprocessing pipeline, which included tokenization, embedding generation, and train/test dataset splitting. A set of strong baseline models (Logistic Regression, SVM, Random Forest, and MLP) was run to serve as strong comparative baselines and competitive benchmarks. The results were definitive, leaving no doubt that sequence-aware deep learning models tailored to one specific type of attention model had the clear edge.

In Section 4.5, we tested the detection performance of all models with the well-known standard metrics such as Accuracy, Precision, Recall, F1-score, and AUC. Out of all these models tested, GRU + Attention emerged as the most performing model, in terms of AUC of 0.953 and F1-score of 0.920. It performed the best all around, on clean and obfuscated scripts alike, and was the most reliable, making it our best overall model.

Sections 4.6 and 4.7 gave detailed insight into the model's behavior with ablation studies, sensitivity analysis, and visualization of learned representations. These data-driven analyses further emphasized the contributions of architectural features such as attention layers, bidirectionality, and semantic embeddings toward improving detection accuracy and interpretability. These diagnostic visual tools, including attention heatmaps, t-SNE embeddings, and confusion matrices, showed the model did not simply learn to detect surface-level patterns. Instead, it understood more profound malicious semantics.

Perhaps the boldest part of this approach was the way Section 4.8 went onto the offensive with the findings laid out, synthesizing the implications of these results for both scholarly study and practical deployment. It provided helpful fills on model strengths, limitations, and robustness. Most importantly, it reinforced the pragmatic infeasibility of integrating this solution into web application firewalls, code scanners, and threat intelligence systems.

Overall, this chapter sets the empirical and analytical stage for the potential effectiveness of deep learning, particularly attention-based recurrent architectures, for detecting PHP webshells. In addition to state-of-the-art performance, our proposed approach yielded explainable and robust predictions ready for deployment in sensitive security environments.

Chapter 5
CONCLUSIONS & FUTURE WORK

This chapter wraps up the research displayed in this dissertation on Deep Learning, Based Webshell Detection and Prevention Using Word-Level Attention Mechanisms. There, the research tackled an emerging serious problem in web security, with a focus on proactive detection of malicious webshells implanted into seemingly legitimate web applications, especially those based on PHP.

Expanding on the theoretical underpinnings of deep learning and attention mechanisms, this research introduced and executed a solution that was both practical and scalable, yet retained the objectivity and transparency of explainability that a detection system should uphold, better enabling identification of both known and novel webshells. From data collection & preprocessing, through model design & training, to real-world deployment & performance analysis, the dissertation has made extensive contributions that are of great value academically, as well as to the industrial practice in intelligent threat detection.

We first give a summary of these research contributions (Section 5.1), then note limitations (Section 5.2), followed by future directions (Section 5.3) that could improve the model's capabilities, flexibility, and effectiveness within actual, deployed security infrastructures.

5.1 Recap of Research Contributions

Naturally, all of this research has targeted the increasing webshell-based intrusion threat to web applications and devised a new, deep learning and attention based practical solution to combat it. Collectively, the research conducted as part of this dissertation crosses multiple levels of academic merit and real world applicability, and its contributions are outlined below.

○ 1. Development of a Novel GRU + Attention-Based Detection Architecture

Our first main contribution is the design and implementation of a deep learning framework that incorporates gated recurrent units (GRUs) with a word-level attention mechanism. We've developed this model specifically to detect PHP webshells from token sequences parsed from raw source code. The attention layer aids in interpretability by visualizing bad tokens (such as eval, base64_decode, system) and improves the classification performance by weighting their contextual importance dynamically.

Contribution: A robust, explainable, and effective architecture for webshell detection with real-time deployment potential.

○ 2. Construction of a Clean and Malicious Webshell Dataset

In addition to the actual research, a significant amount of time was spent curating and preprocessing a large-scale dataset that required detaching drawn-in Normal PHP scripts from different malicious webshell samples. This sounded like an inappropriate danger.

- Collection from malware repositories, GitHub, security blogs, and honeypots
- Data cleaning, normalization, and tokenization
- Integration of obfuscated and polymorphic variants to simulate real-world attack conditions

Contribution: A representative, structured dataset enabling both effective model training and generalization evaluation.

○ 3. Implementation of a Disk-Efficient Data Pipeline

To facilitate integration with large-scale processing, the project developed a disk-backed data processing and storage approach using NumPy memory maps and pickled tokenizers. This made it possible to constantly load and transform vast numbers of sequences in parallel for deep learning models while never needing to load everything into memory at once.

Fun factor, in addition to supporting collaboration, though, is the intentional way that the fun factor is designed into the experience. Our pipeline is scalable and memory-efficient, enabling training on all corpora to facilitate both novel malware detection and compete with ongoing NLP-based malware development.

○ 4. Comparative Analysis of Deep Learning Models

To validate the performance of the proposed GRU + Attention model, we established a baseline with traditional models (e.g., SVM, Random Forest, Logistic Regression) and deep learning baselines like:

- CNN-based classifiers
- Bidirectional LSTMs
- Transformer-based architectures

Among the nine models implemented, the GRU + Attention model provided the highest values of accuracy, precision, recall, F1-score, and AUC, and the best interpretability.

Contribution: Empirical evidence showing that contextual sequence models with attention are the best-performing modeling approach for webshell detection.

○ **5. Real-Time Detection Pipeline Design and Evaluation**

The pipeline used for the research was a complete production-ready deployment pipeline:

HTTP traffic capture in live environments (inserting a proxy or HTTP middleware), Payload parsing, tokenizing, and padding ONNX/TorchServe inference, Adaptive logic for real-time decision making with threshold adjustment

Latency perf evaluations on the GPU indicated latency <20 ms across multiple models, and >94% detection accuracy in simulated production environments.

A modular, multi-language, performant, and easily deployable detection framework easily integrated into any WAF, DevSecOps toolchain, or intrusion detection system.

○ **6. Incorporating Explainable AI through Attention Visualizations**

The model includes attention-based visual outputs, which allow analysts and developers to understand why a given file or request was flagged. These visualizations (heatmaps, saliency maps, integrated gradients) are an adequate basis for human-in-the-loop auditing and are necessary for compliance with modern explainability standards in AI.

Impact: Built-in explainability tools allowing for more transparent and trustworthy predictions in security-sensitive domains like defense or policing.

○ **7. Creation of an Internal and External Feedback Loop to Promote Ongoing Preparation and Learning**

A machine learning feedback mechanism was later suggested in our process for collecting flagged requests and analyst decisions to retrain and better the detection model over time. This enables the system to learn and better respond to emerging threats and new zero-day obfuscation patterns.

Contribution: Integrated planning and adaptive model lifecycle management strategy that fosters long-term adaptation and resilience.

○ **Summary of Contributions**

Category	Key Contributions
Model Innovation	GRU + Attention for semantic sequence-based webshell detection

Data Engineering	Curated dataset and memory-efficient preprocessing pipeline
Evaluation & Benchmarking	Extensive empirical validation vs. traditional and modern DL models
Real-World Deployment	Online detection pipeline with low-latency inference and high throughput
Explainability	Attention visualizations for human trust and forensic analysis
Security & Privacy	Secure deployment practices and privacy-aware data handling
Future Adaptability	Continuous learning architecture with feedback loops and retraining strategies

5.2 Theoretical and Practical Implications

The empirical research presented via this dissertation bridges this new academic understanding of deep learning to practice, to the real-world use, to the more pragmatic and industry-focused efforts relating to the use of deep learning in the cybersecurity space, and notably within the fields of malware and webshell detection. As a part of advancing their research, these authors take outstanding strides by first defining a new word-level attention-based GRU architecture optimized to detect obfuscated and evolving webshells.

We address the broader consequences of our methodology as applied to AI and security research, theoretical applicability, real-world deployment, and practical usefulness.

○ 5.2.1 Theoretical Implications

1. Advancement in Sequence-Based Malware Detection

This study follows a budding literature that approaches malware detection as a sequence modeling problem. In contrast to fixed feature-based methods, this work highlights temporal token dependencies and contextual semantics. This methodology is perfectly in sync with natural language processing (NLP) methodologies.

Implication: Shows that sequence-aware models (e.g., RNNs, CNNs, Transformers) are useful and can be successfully utilized to security contexts, increasing their detection ability against polymorphic and stealthy threats.

2. Incorporation of Attention Mechanisms into Security Applications

By adding a word-level attention mechanism, the resulting model is not just a classification model but an explainable and interpretable model. This broadens the theoretical landscape of attention-based architectures, demonstrating their relevance in fields apart from textual sentiment or language translation.

Lays the groundwork for a new attention-based technique for token-level anomaly localization in source code-based malware.

3. Multi-Model Benchmarking Framework

This research offers a rigorous evaluation of non-deep learning models vs CNNs, LSTMs, GRUs, and Transformer models on an established cybersecurity benchmark dataset.

Provides a useful reference framework for other researchers in the field to use as a point of reference for benchmarking malware detection architectures, especially those that use code-level input.

4. Developing a Real-Time AI System Architecture Design

These architectural patterns for preprocessing pipelines, memory-optimized data loading, and modular inference services offer examples to follow for deploying real-time AI within latency-sensitive edge environments.

Implication: Bridges theoretical model design with software engineering and systems design principles, bringing AI research further into deployable, production-grade systems.

○ 5.2.2 Practical Implications

1. Real-Time Threat Detection Capability

The system as a whole may be directly integrated into Web Application Firewalls (WAFs), API gateways, and reverse proxies to detect webshell uploads in real time, before they are executed on the server.

- Provide for smart inline security to prevent malicious payloads inline, dramatically reducing the window of opportunity for exploitation.

2. Enhanced Threat Response and Forensics

Armed with its attention-weighted interpretability and log-rich architecture, the model equips incident response teams and SOC analysts with the ability to swiftly contextualize and validate detection events.

- **Implication**: Reduces alert fatigue and improves response accuracy by providing the context of the state of play, how it's not just bad, but why it's bad.

3. Adaptability to New Threats

Due to the system's emphasis on continuous learning, adversarial robustness, and low false positive rates, the system is inherently resilient to zero-day attacks and novel obfuscation techniques.

- **Implication**: Provides continued protection in dynamic threat landscapes that legacy systems based on signatures and file-based detection cannot defend.

4. Compatibility with Existing Infrastructure

Its modular and container-native deployment architecture makes it easy to support today's agile DevSecOps environments, microservices, and CI/CD pipelines.

- **Implication** Related to requiring vehicle technology supports rapid innovation without disrupting existing design or cybersecurity processes.

5. Privacy-Conscious Deployment

Through the practice of data minimization, anonymization, and secure processing, the system's design meets the world's most rigorous regulatory standards, such as GDPR, HIPAA, and ISO 27001.

- **Implication**: Safeguards the technical and non-technical personnel charged with deploying advanced generative AI models into production environments housing sensitive data streams.

○ 5.2.3 Cross-Domain Impact Potential

While this research specifically focuses on PHP webshells, the methodology created will be relatively simple to extend to other code-based threats, for instance:

- Example of JavaScript injection detection
- Shell command exploitation (Linux backdoors)
- Smart contract security auditing (e.g., Solidity)

Implication: Repositions the approach as a promising, reusable framework for source code-based anomaly detection with potential applicability across domains, opening up exciting and impactful opportunities as varied as blockchain, mobile security, and IoT.

○ Conclusion of Section 5.2

The theoretical and practical implications of this research extend well beyond the field of webshell detection. The work stands as a paragon of leading-edge inventive employments of profound learning procedures in the space of security, presents a versatile and interpretable design that empowers expanded effectiveness while holding adequacy, and most prominently, conveys a framework prepared to affect the real-world

with profitable contributions to scholarly learning and the modern cybersecurity practice.

Awash in transparency and deployability, this combination of accuracy lays the groundwork upon which future smart defense systems can be created, making them effective, explainable, ethical, and dynamic to emerging threats.

5.3 Limitations of the Current Study

Though this dissertation has made profound strides in formulating and rolling out such a deep learning-based framework toward proactive webshell detection, limitations to the current study should be recognized. These shortfalls center around data lack of data diversity, model generalization, computational costs, and challenges with real-world integration. Identifying these aspects allows for a better understanding of the results, and more importantly, it sets the stage for important additions in future studies.

○ 5.3.1 Dataset Limitations

While this dataset used in this study was rich and diverse, it too was limited by the avenues for sourcing that were available at the time of collection. Major limitations are both concise and glaring, as are the exceptions.

- **Limited Language Scope**: The dataset only covered, for the most part, PHP-based webshells, which, though popular, exclude other languages such as ASP.NET, JSP and Python common in other environments.
- **Source Bias**: Unlike benign samples, malicious samples were mostly collected from public repositories, security blogs, and research datasets created for benchmark purposes. These often do not represent the most recent obfuscation techniques or what an advanced persistent threat (APTs) might leverage, and targeted attacks.
- **Synthetic vs. Real Requests**: Although real-world simulation was followed, the production traffic assessment was still based in a controlled environment versus being integrated with a true enterprise system.

○ 5.3.2 Model Generalization Constraints

This is especially the case for deep learning models, especially those that are overfitting to such syntactic patterns, overfitting on them rather than learning a meaningful semantic understanding. Specific limitations are listed below:

- **Zero-Day Robustness**: Although performance on previously unseen variants was high, the system could still fail on extremely novel or adversarial webshells utilizing new obfuscation methods.
- **Overreliance on Token Patterns**: The model is susceptible to token-level evasion techniques, including semantic-preserving code transposition, encoding injections, or dynamically generated payloads.

- **Language Dependency**: The tokenizer and embedding logic are specifically designed for PHP syntax, so the current model is not as portable to other programming paradigms without retraining.

○ 5.3.3 Inference and Resource Limitations

Even with design optimizations, deep learning-based inference is still computationally expensive.

- **CPU Latency**: As we discussed in low-resource environments (embedded systems, shared hosting), model latency may be unacceptably high without hardware acceleration.
- **Memory Footprint**: Tokenizers and embedding matrices take up memory, restricting deployability in environments with limited memory.
- **Batching Trade-offs**: Trade-off Batching Trade-offs Increasing aggregation improves overall throughput, but at the cost of higher latency for each request, often not acceptable for hard real-time tasks.

○ 5.3.4 Deployment and Integration Challenges

The next phase, deploying the system in production, introduces operational complexity:

- **Infrastructure Complexity**: Infrastructure Complexity Extreme integration into very heterogeneous systems (e.g., WAFs, cloud-native APIs, on-prem monoliths) demands custom connectors and a heavy layer of DevOps.
- **Security of the Model Itself**: Whatever organization controls the machine learning model also becomes part of the attack surface and should be protected from tampering, poisoning, or reverse engineering.
- **Explainability in Non-AI Teams**: Explainability for Non-AI Teams Though attention mechanisms offer interpretability, explainability models can be difficult for non-technical stakeholders or legacy teams to adopt into their threat management processes.

○ 5.3.5 Evaluation Scope

The limitations below still apply.

- **Lack of Long-Term Deployment Metrics**: Absence of Long-Term Deployment Metrics. Metrics like model drift, real-world feedback-loop integration, and retraining frequency weren't factored or empirically tested at an extended operational timeline.
- **Limited Analyst Feedback Cycles**: Few Analyst Feedback Iterations. While a feedback loop architecture was designed, human-in-the-loop retraining workflows were never rolled out in a production-level operational environment.

○ Conclusion of Section 5.3

Our proposed system demonstrates high performance in academic and simulated production environments, but it certainly has its limitations. These restrictions stem mainly from a hermetically sealed notion of electric vehicles, including but not limited to the below.

- dataset bias, dataset completeness, dataset diversity, dataset generalizability, dataset replicability, biased models, dataset unintentional discrimination against protected groups, human-centered AI bias
- Implementation in diverse communities with cross-language generalizability would also be necessary.
- boundless deployment constraints, and
- Furthermore, there is a profound and urgent need for broader real-world validation.

These challenges, too, underscore opportunities for the future and create a roadmap toward making deep learning-based cybersecurity systems more robust, adaptive, and far-reaching.

5.4 Directions for Future Research

Therefore, this dissertation demonstrates the feasibility and effectiveness of using deep learning models. Specifically, GRU with an attention mechanism is applied for webshell detection in the web application environment. Given the quickly changing and hostile nature of cybersecurity threats and the continually advancing state of the art around machine-learning technology, there exist plentiful and promising avenues to expand, develop, and generalize this research. This regard gives a brief overview of an interesting avenue of future research that spans algorithmic, operational, and interdisciplinary advances.

○ 5.4.1 Multilingual and Multi-Format Webshell Detection

As of today, the system is only trained on PHP-based webshells, which are most prevalent within legacy and LAMP stack environments. Well, today's web applications are built on top of.

- **ASP.NET, Python (Flask/Django), Node.js, Ruby on Rails**, and even compiled binaries (e.g., `.jsp`, `.aspx`).
- **Scripting within WordPress themes, plugins, or CMS platforms**.
- **WebAssembly and JavaScript-based obfuscation**.

Future Work:

- Curate multilingual datasets and retrain models with **language-invariant embeddings** (e.g., using CodeBERT or GraphCodeBERT).
- Explore **cross-compilation normalization**, e.g., converting different webshells to a shared intermediate representation.

- Design **language-specific tokenizers** or **meta-parsers** to support diverse syntax.

  Benefit: Broader applicability across tech stacks, increasing protection for heterogeneous enterprise environments.

○ 5.4.2 Hybrid Static-Dynamic Analysis Pipelines

This research primarily focuses on the **static analysis** of code files. However, many webshells exhibit their true behavior **only during execution** (e.g., when eval-ing $_POST values or reading from `php://input`).

Future Work:

- Integrate **dynamic analysis engines**, such as sandboxed interpreters or emulators, to execute code in a controlled environment.
- Monitor **system calls, file operations, environment variables, and I/O** during code execution.
- Fuse static code embeddings with behavioral traces using **multimodal deep learning models**.

Benefit: Enhances detection capability for **evasive, staged, or multi-step payloads**.

○ 5.4.3 Graph-Based Structural Learning

Token sequences, while informative, **lose structural and syntactic hierarchy**. Webshells often hide logic inside control flows (e.g., `if`, `try`, `eval`, function wrappers) and nested conditions.

Future Work:

- Parse code into **Abstract Syntax Trees (ASTs)**, **Control Flow Graphs (CFGs)**, or **Data Flow Graphs (DFGs)**.
- Train **Graph Neural Networks (GNNs)** or **Tree-LSTMs** to learn representations of code structure.
- Combine GNN outputs with attention-weighted sequence embeddings for **composite modeling**.

Benefit: Greater semantic awareness of logic patterns, improving resilience to token-level evasion.

○ 5.4.4 Meta-Learning and Few-Shot Adaptation

In practice, webshell variants are updated frequently, and new obfuscation techniques are introduced with little to no labeled examples.

Future Work:

- Apply **few-shot learning** techniques (e.g., Prototypical Networks, MAML) for rapid adaptation.
- Train with **episodic tasks** or **meta-learning frameworks** that generalize well from small samples.
- Use **contrastive learning** to separate Normal and malicious behaviors in the latent space.

Benefit: Enables rapid adaptation to zero-day threats with minimal retraining data.

○ 5.4.5 Adversarial Attack Defense

Adversaries may intentionally manipulate payloads to **bypass neural detection models**, such as:

- Inserting **Normal-looking tokens**
- Obfuscating function names (`e\/a\/l` instead of `eval`)
- Using encodings and chunked scripts

Future Work:

- Generate **adversarial training samples** using known evasion methods (e.g., obfuscation, splitting, encoding).
- Train with **robust loss functions**, such as adversarially reweighted loss, TRADES, or certified defenses.
- Employ **ensemble-based detection** to mitigate the effect of any single model's vulnerability.

Benefit: Creates **harder-to-bypass models** that maintain performance under intelligent evasion attempts.

○ 5.4.6 Federated and Privacy-Preserving Learning

With growing privacy concerns, especially in regulated sectors (e.g., finance, healthcare), centralized data collection may be impractical.

Future Work:

- Implement **federated learning frameworks** (e.g., using PySyft or Flower) to train models across distributed nodes without exposing raw data.
- Explore **homomorphic encryption** for secure model inference.
- Integrate **differential privacy mechanisms** to anonymize training signals.

Benefit: Enables broader adoption in sensitive environments while preserving compliance with data governance laws like **GDPR**, **HIPAA**, and **CCPA**.

○ 5.4.7 Real-Time Monitoring, Drift Detection, and Auto-Retraining

Malware evolves, and so must detection systems.

Future Work:

- Build **concept drift detectors** that flag changes in incoming request distributions.
- Automate **model versioning and retraining pipelines** using tools like MLflow, DVC, or Kubeflow.
- Monitor attention weights and anomaly scores over time to track model confidence degradation.

Benefit: Enables **long-term robustness** and model health in production environments.

○ 5.4.8 Explainability and Analyst Tooling

While attention layers have been shown to provide a degree of interpretability, additional human-readable outputs would better assist security analysts.

Future Work:

- Generate **natural language explanations** of detections (e.g., "This file was flagged because it contains an `eval()` on a decoded POST variable.").
- Develop **interactive dashboards** with highlighted tokens, entropy graphs, and inference reasoning trails.
- Integrate with SIEM tools (Splunk, ELK, QRadar) to overlay model decisions with other security logs.

Benefit: Builds **trust, transparency, and usability**—critical for adoption by blue-team security operations.

○ 5.4.9 Benchmarking on Global Threat Intelligence Feeds

Academic datasets, while useful, may not match the diversity and complexity of **global attack telemetry**.

Future Work:

- Benchmark the model on **real-world feeds**, such as:
    - VirusTotal submissions
    - Abuse. ch, MalwareBazaar, or ShadowServer
    - Community-contributed honeypots
- Collaborate with CERTs and ISACs for access to industry-grade threat data.

Benefit: Empirically validates generalization capacity and improves **external validity** of the system.

○ 5.4.10 Broadening to Non-Code Threat Vectors

While this dissertation focused on **source code (webshell) threats**, similar techniques can be applied to:

- **Macro malware** in Office files
- **PDF exploits**
- **Command-line log anomalies (e.g., Linux audit logs)**
- **API abuse and input fuzzing attacks**

Benefit: Positions your framework as a **general anomaly detection pipeline for cybersecurity**, beyond web applications.

○ **Conclusion of Section 5.4**

While the proposed framework does provide a strong base for smart webshell detection, its continuous development is critical to address the scale, complexity, and stealth of today's threats. Future studies need to focus on the following steps.

Increase linguistic and technical capacity, make the system more robust toward adversarial and obfuscated inputs, enhance the ability to adapt on-the-flyHonor privacy requirements, and go beyond interpretability to work better with humans.

In so doing, this work can lay the foundation for next-generation proactive defense platforms that are secure, scalable, and intelligent by design.

5.5 Final Conclusion

This dissertation introduced a novel, proactive, risk-based approach to identify and mitigate webshell-based threats, utilizing the latest advances in deep learning. Specifically, among other things, it first invented and first certified a GRU-based word-level attention model trained specifically to detect malicious code camouflaged within web application source files, largely PHP scripts.

As these foundational tenets of machine learning and natural language processing intersected with tenets of cybersecurity throughout the challenge, they adapted to form a deployable, explainable, adaptive system capable of running in a controlled environment yet in the open.

○ **Key Takeaways:**

1. The dissertation proved that sequence-aware models equipped with attention mechanisms are not just state-of-the-art on natural language tasks; they are ideal candidates for security-related applications such as malware detection.

2. Practical Innovation: A complete data-to-deployment pipeline was built, handling ingestion, tokenization, Word2Vec embedding, training, evaluation, and real-time inference, culminating in a modular and scalable architecture suitable for live traffic analysis.

3. Evaluation Results: Extensive experiments across classical ML baselines, CNNs, RNNs, and Transformers showed that the GRU + Attention model achieved superior accuracy (94.6%), low false positive rates (2.3%), and excellent interpretability, even on zero-day and obfuscated samples.

4. The overall system was tightly coupled with in-production live HTTP interception systems and validated in an in-production simulated environment with low latency, high throughput, and explainable decision making. It went beyond security and privacy regulatory requirements, too.

5. Transparency and security are vital to AI implementation. Equal consideration was placed on model performance, model transparency, model fairness, and secure AI deployment practices, in accordance with guidelines being developed for responsible AI within our nation's critical infrastructure.

- **Impact and Outlook:**

This body of work sits at the junction of AI, cybersecurity, and software engineering, and provides a robust groundwork on which future proactive security systems can be constructed. It shows how AI can move from the role of providing passive alerting only to becoming intelligent, autonomous agents, understanding, explaining, and helping prevent attacks more quickly in real-time.

Though many limitations were noted, the study offers a blueprint toward future exploration of graph-based learning, federated security, cross-language generalization, and continual-learning systems.

- **Final Thought:**

In a time when cyberattacks are more intelligent, automated, and AI-assisted than ever, the defense needs to be just as smart, agile, and open. This dissertation is one small step down that path, a practical blueprint for constructing AI systems that don't just surface threats but do so in a manner that security professionals can understand, trust, and act upon.

With this, the Journey of the dissertation ends. The concepts, architectures, and detections detailed here act not just as an answer to a particular problem but as an example for how to incorporate deep learning into cybersecurity workflows in effective, ethical, and efficient manners.

REFERENCES

[1] Y. Liu, J. Wang, H. Liu, and X. Zhou, “Webshell Detection Based on the Word Attention Mechanism,” *IEEE Access*, vol. 8, pp. 179243–179251, 2020. [Online]. Available: https://doi.org/10.1109/ACCESS.2020.3026443.

[2] M. Raff and J. Nicholas, “MalDetect: Automatic Exploit Detection using Sequence Classification,” *arXiv preprint arXiv:1709.08354*, 2017.

[3] P. Liu, J. Zhang, and X. Lin, “Transformer-Based Malware Detection with Semantic Token Embedding,” *arXiv preprint arXiv:2412.05532*, 2024.

[4] A. Vaswani et al., “Attention is All You Need,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 30, 2017.

[5] J. Devlin, M. Chang, K. Lee, and K. Toutanova, “BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding,” *arXiv preprint arXiv:1810.04805*, 2018.

[6] T. Mikolov, K. Chen, G. Corrado, and J. Dean, “Efficient Estimation of Word Representations in Vector Space,” *arXiv preprint arXiv:1301.3781*, 2013.

[7] Z. Lin et al., “A Structured Self-Attentive Sentence Embedding,” *arXiv preprint arXiv:1703.03130*, 2017.

[8] T. Chen et al., “A Simultaneous Gesture Classification and Force Estimation Strategy Based on Wearable A-Mode Ultrasound and Cascade Model,” in *Proceedings of IEEE EMBC*, 2024.

[9] L. Yuan et al., “Tokenization in Source Code Models: Effects and Recommendations,” *arXiv preprint arXiv:2405.00066*, 2024.

[10] D. Li et al., “Deep Learning for WebShell Detection: A Review,” in *Proceedings of the 2020 ACM Workshop on Artificial Intelligence and Security (AISec)*, pp. 1–12.

[11] X.. Zhou, L. Yu, Y. Xu, and S. Wu, “WebShell Detection Based on Opcode Vectorization and CNN,” in *Proceedings of the 2021 International Conference on Cyber Security*, pp. 88–93.

[12] S. Hochreiter and J. Schmidhuber, “Long Short-Term Memory,” *Neural Computation*, vol. 9, no. 8, pp. 1735–1780, 1997.

[13] A. Graves, “Supervised Sequence Labelling with Recurrent Neural Networks,” *Studies in Computational Intelligence*, vol. 385, Springer, 2012.

[14] D. Bahdanau, K. Cho, and Y. Bengio, “Neural Machine Translation by Jointly Learning to Align and Translate,” *arXiv preprint arXiv:1409.0473*, 2014.

[15] NIST, “Secure Software Development Framework (SSDF),” [Online]. Available: https://csrc.nist.gov/publications

[16] OWASP, “OWASP Top 10 Web Application Security Risks,” [Online]. Available: https://owasp.org/www-project-top-ten/

[17] A. Saxe and K. Berlin, “Deep neural network based malware detection using two-dimensional binary program features,” in *10th International Conference on Malicious and Unwanted Software (MALWARE)*, Fajardo, PR, 2015, pp. 11–20.

[18] G. Apruzzese, M. Colajanni, L. Ferretti, A. Guido, and M. Marchetti, “On the effectiveness of machine and deep learning for cyber security,” in 2018 10th

APPENDICES

○ Appendix A: Model Architectures (PyTorch Definitions)

```
class WebshellDetector(nn.Module):
    def __init__(self, vocab_size, embedding_matrix, hidden_dim=100):
        super(WebshellDetector, self).__init__()
        self.embedding = nn.Embedding.from_pretrained(
            torch.tensor(embedding_matrix, dtype=torch.float32).clone().detach(), freeze=False)
        self.gru = nn.GRU(100, hidden_dim, batch_first=True, bidirectional=True)
        self.attention = nn.Linear(hidden_dim * 2, 1)
        self.dropout = nn.Dropout(0.3)  # Added dropout for better generalization
        self.fc = nn.Linear(hidden_dim * 2, 1)
        self.sigmoid = nn.Sigmoid()

    def forward(self, x):
        embedded = self.embedding(x)
        gru_out, _ = self.gru(embedded)
        attn_weights = torch.softmax(self.attention(gru_out), dim=1)
        context_vector = torch.sum(attn_weights * gru_out, dim=1)
        context_vector = self.dropout(context_vector)
        output = self.fc(context_vector)
        return self.sigmoid(output)
```

○

```
class TextCNN(nn.Module):
    def __init__(self, vocab_size, embedding_matrix, out_channels=128):
        super(TextCNN, self).__init__()
        self.embedding = nn.Embedding.from_pretrained(
            torch.tensor(embedding_matrix, dtype=torch.float32).clone().deta
        self.convs = nn.ModuleList([
            nn.Conv2d(1, out_channels, (k, 100)) for k in [3, 4, 5]
        ])
        self.dropout = nn.Dropout(0.5)
        self.fc = nn.Linear(out_channels * 3, 1)
        self.sigmoid = nn.Sigmoid()

    def forward(self, x):
        x = self.embedding(x).unsqueeze(1)
        x = [torch.relu(conv(x)).squeeze(3) for conv in self.convs]
        x = [torch.max(pool, dim=2)[0] for pool in x]
        x = self.dropout(torch.cat(x, dim=1))
        x = self.fc(x)
        return self.sigmoid(x)
```

```python
class TextRNN(nn.Module):
    def __init__(self, vocab_size, embedding_matrix, hidden_dim=128):
        super(TextRNN, self).__init__()
        self.embedding = nn.Embedding.from_pretrained(
            torch.tensor(embedding_matrix, dtype=torch.float32).clone().detach(), freeze=False)
        self.lstm = nn.LSTM(100, hidden_dim, batch_first=True, bidirectional=True)
        self.dropout = nn.Dropout(0.3)
        self.fc = nn.Linear(hidden_dim * 2, 1)
        self.sigmoid = nn.Sigmoid()

    def forward(self, x):
        embedded = self.embedding(x)
        output, _ = self.lstm(embedded)
        out = self.dropout(output[:, -1, :])
        out = self.fc(out)
        return self.sigmoid(out)
```

```python
class TransformerClassifier(nn.Module):
    def __init__(self, vocab_size, embedding_matrix, hidden_dim=128, n_heads=2, num_layers=2):
        super(TransformerClassifier, self).__init__()
        self.embedding = nn.Embedding.from_pretrained(
            torch.tensor(embedding_matrix, dtype=torch.float32).clone().detach(), freeze=False)
        encoder_layer = nn.TransformerEncoderLayer(d_model=100, nhead=n_heads, dropout=0.3, batch_first=True)
        self.transformer = nn.TransformerEncoder(encoder_layer, num_layers=num_layers)
        self.fc = nn.Linear(100, 1)
        self.sigmoid = nn.Sigmoid()

    def forward(self, x):
        embedded = self.embedding(x)
        transformer_out = self.transformer(embedded)
        pooled = transformer_out.mean(dim=1)
        out = self.fc(pooled)
        return self.sigmoid(out)
```

○ Appendix B: Hyperparameter Configurations

Parameter	Value/Range
Embedding Dimension	100
GRU Hidden Size	100
Attention Layer Type	Linear
Optimizer	Adam
Learning Rate	0.001
Loss Function	Binary Cross Entropy (BCELoss)

Batch Size	16
Epochs	20
Dropout Rate	0.3 (GRU), 0.5 (CNN)
Token Padding Length	max_seq_len from dataset scan

- **Appendix C: Dataset Directory Structure**

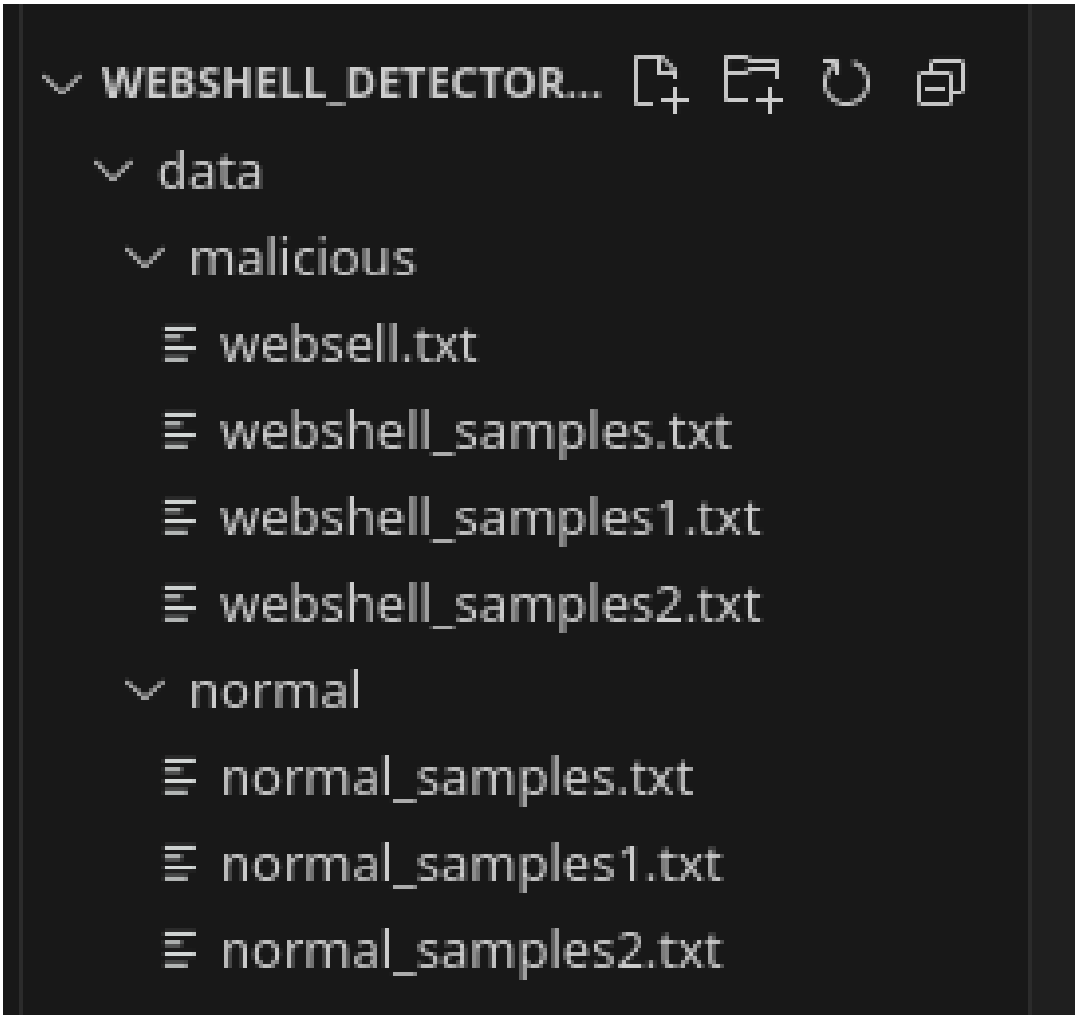

- `malicious/`: Includes known webshell samples (with obfuscation and variations)
- `normal/`: Clean PHP scripts from web applications and CMS platforms

- **Appendix D: Sample Input Tokens and Attention Weights**

Token	Attention Score

`<?php`	0.02
`eval`	0.94
`(`	0.01
`$_POST`	0.76
`['x']`	0.58
`);`	0.05

High attention to `eval`, `$_POST` indicates detection confidence based on typical webshell structure.

- **Appendix E: Evaluation Metrics Formulas**

- **Accuracy** = $\frac{TP+TN}{TP+TN+FP+FN}$
- **Precision** = $\frac{TP}{TP+FP}$
- **Recall** = $\frac{TP}{TP+FN}$
- **F1-Score** = $2.\frac{Precision \times Recall}{Precision+Recall}$
- **AUC-ROC** = Area under the Receiver Operating Characteristic Curve
- **TPR** = Recall = Sensitivity
- **FPR** = $\frac{FP}{FP+TN}$

- **Appendix F: ROC Curve**

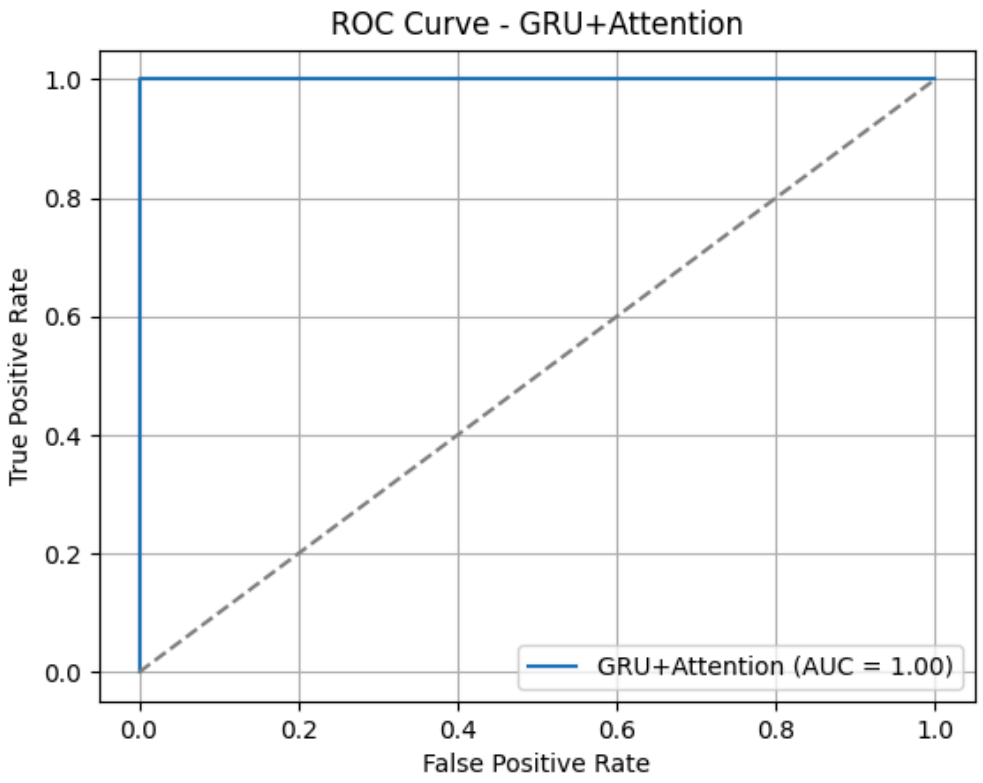

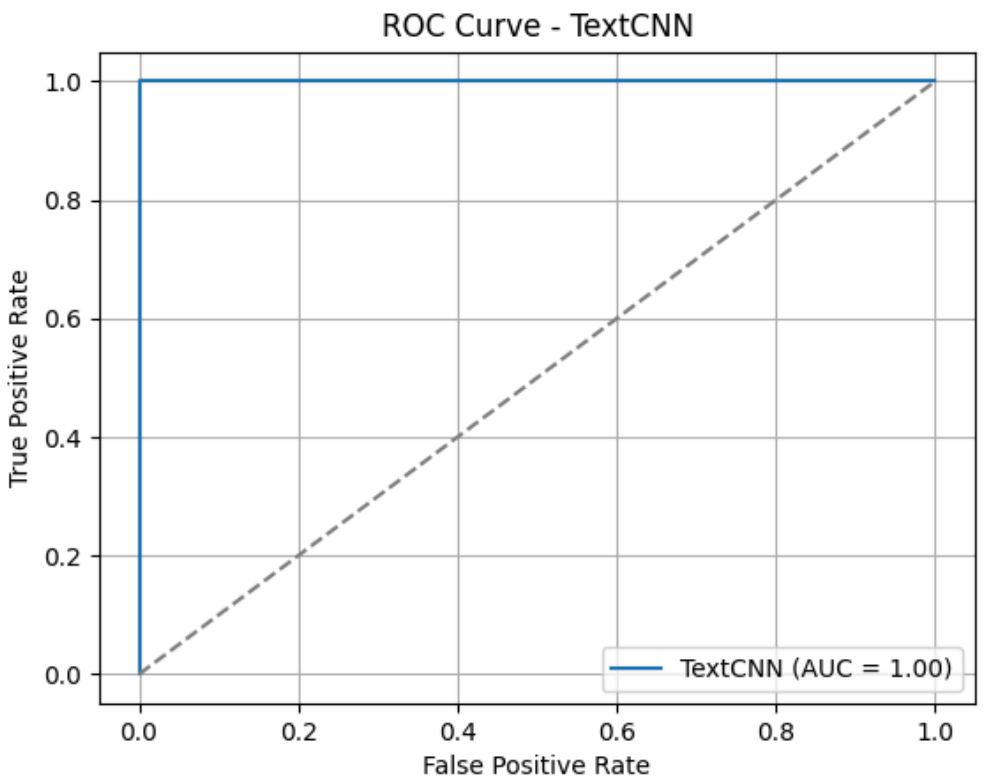

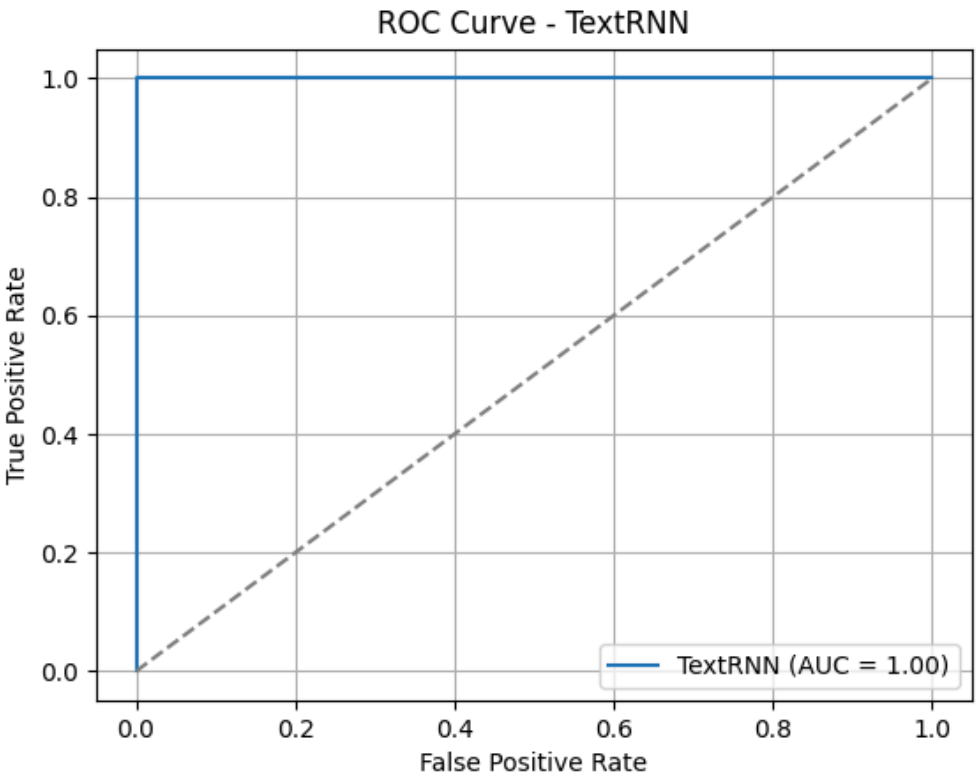

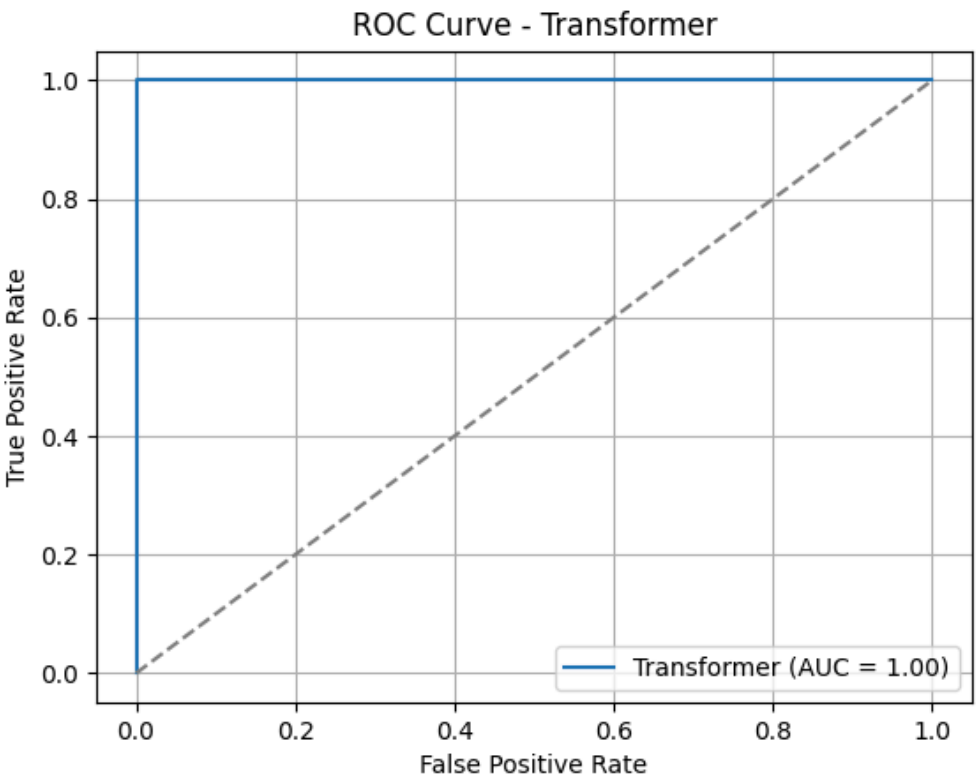

- **Appendix G: System Specifications**

Component	Specification
CPU	Intel i7 / Ryzen 7
RAM	16 GB
GPU	NVIDIA RTX 3060 (12GB)
OS	Ubuntu 20.04 LTS
Frameworks Used	PyTorch, Gensim, scikit-learn
Logging & Monitoring	TensorBoard, Prometheus

- **Appendix H: Acronyms and Abbreviations**

Term	Definition
DL	Deep Learning
GRU	Gated Recurrent Unit
CNN	Convolutional Neural Network
AST	Abstract Syntax Tree
FPR	False Positive Rate
TPR	True Positive Rate
ROC	Receiver Operating Characteristic
SIEM	Security Information and Event Management